


Training DGMs

October 25, 2022



We suppose that our training data come from an **original distribution** p_{data} (which is **not explicitly known**).

Instead, we have access to a set \mathcal{D}_{train} of n **samples** \mathbf{x}^i from p_{data} , which we suppose independent and identically distributed.

On the other side, we consider a family of **parameterized distributions** p_{θ} , parameterized by θ (for example, the parameters we saw on the auto-regressive models).

We want to determine an “optimal” θ^* such that p_{θ^*} is **as close as possible to** p_{data} .

Note that it is **impossible to get** p_{data} **exactly**: We do not have access to it but only to an approximation through the n samples.

Example



Each image can be seen as a vector \mathbf{x}^i of $28 \times 28 = 784$ binary variables.

The whole space of binary images is huge: $2^{784} \approx 10^{236}$. Of course, the **set of plausible figure images is much smaller but still very large**. So, even 1 million sample images would be little compared to the true support of p_{data} .

0.1 Loss functions for generative models

An important point to state the problem correctly is to evaluate **how far** a candidate p_θ is from p_{data} .

The **Kullback-Leibler (KL) divergence** between two distributions p and q is

$$D_{KL}(p||q) \triangleq \sum_{\mathbf{x}} p(\mathbf{x}) \log \frac{p(\mathbf{x})}{q(\mathbf{x})}.$$

- It is **not symmetric**: $D_{KL}(p||q) \neq D(q||p)$.
- It satisfies, for all p, q

$$D_{KL}(p||q) \geq 0,$$

with equality if and only if $p = q$.

To verify this statement, we use the fact that the KL expression can be interpreted as an **expectation** and the fact that $-\log$ is a **convex** function:

$$-\log \left(E_{\mathbf{x} \sim p} \left[\frac{q(\mathbf{x})}{p(\mathbf{x})} \right] \right) \leq E_{\mathbf{x} \sim p} \left[-\log \left(\frac{q(\mathbf{x})}{p(\mathbf{x})} \right) \right].$$

Now, the term on the left is:

$$-\log \left(E_{\mathbf{x} \sim p} \left[\frac{q(\mathbf{x})}{p(\mathbf{x})} \right] \right) = -\log \left(\sum_{\mathbf{x}} p(\mathbf{x}) \left[\frac{q(\mathbf{x})}{p(\mathbf{x})} \right] \right) = -\log \left(\sum_{\mathbf{x}} q(\mathbf{x}) \right) = -\log 1 = 0.$$

Hence:

$$E_{\mathbf{x} \sim p} \left[-\log \left(\frac{q(\mathbf{x})}{p(\mathbf{x})} \right) \right] \geq 0.$$

To come back to our original problem:

$$D_{KL}(p_{data} || p_{\theta}) = E_{\mathbf{x} \sim p_{data}} \left[-\log \left(\frac{p_{\theta}(\mathbf{x})}{p_{data}(\mathbf{x})} \right) \right] \quad (1)$$

$$= -E_{\mathbf{x} \sim p_{data}} [\log (p_{\theta}(\mathbf{x}))] + E_{\mathbf{x} \sim p_{data}} [\log (p_{data}(\mathbf{x}))]. \quad (2)$$

Now, observe that the second term does **not depend** on θ . Hence,

$$\theta^* = \arg \min_{\theta} -E_{\mathbf{x} \sim p_{data}} [\log (p_{\theta}(\mathbf{x}))] = \arg \max_{\theta} E_{\mathbf{x} \sim p_{data}} [\log (p_{\theta}(\mathbf{x}))].$$

This expression on the right is the **log likelihood**:

- high values when samples from p_{data} are **evaluated with high values on p_{θ}** (likely samples from p_{data} should be likely on p_{θ}),
- very low (negative) values if $p_{\theta}(\mathbf{x})$ is small.

Note that with this formulation:

- we **know** how to make p_{θ} as close as possible to p_{train} ,
- we **do not know** how close we will be at the end: there will still be an unknown part $E_{\mathbf{x} \sim p_{data}} [\log (p_{data}(\mathbf{x}))]$ that we are not taking into account.


How to compute $E_{\mathbf{x} \sim p_{data}} [\log (p_{\theta}(\mathbf{x}))]$?

Since we only have samples $\mathbf{x} \in \mathcal{D}_{train}$ from p_{data} , we will approximate the true expectation by an **empirical expectation**

$$E_{\mathbf{x} \sim p_{data}} [\log (p_{\theta}(\mathbf{x}))] \approx E_{\mathbf{x} \in \mathcal{D}_{train}} [\log (p_{\theta}(\mathbf{x}))].$$

Hence, training will be done by **maximizing the log-likelihood**:

$$\theta^* = \arg \max_{\theta} \frac{1}{|\mathcal{D}_{train}|} \sum_{\mathbf{x} \in \mathcal{D}_{train}} \log (p_{\theta}(\mathbf{x})) = \arg \max_{\theta} \sum_{\mathbf{x} \in \mathcal{D}_{train}} \log (p_{\theta}(\mathbf{x})).$$



Note that this is equivalent to **maximizing the likelihood** of the independent data within \mathcal{D}_{train} :

$$\theta^* = \arg \max_{\theta} \prod_{\mathbf{x} \in \mathcal{D}_{train}} p_{\theta}(\mathbf{x}).$$

We will call:

$$L(\theta, \mathcal{D}_{train}) \triangleq \prod_{\mathbf{x} \in \mathcal{D}_{train}} p_{\theta}(\mathbf{x}).$$

The approximation above is fundamentally a **Monte-Carlo estimation process**!

Idea: **approximate an expected value** (here, $E_{\mathbf{x} \sim p_{data}} [\log(p_{\theta}(\mathbf{x}))]$) through an **empirical expectation on samples** from the same distribution.

$$l = E_{\mathbf{x} \sim p_{data}} [\log(p_{\theta}(\mathbf{x}))] \approx \hat{l} = \frac{1}{|\mathcal{D}_{train}|} \sum_{\mathbf{x} \in \mathcal{D}_{train}} \log(p_{\theta}(\mathbf{x})) = \frac{1}{|\mathcal{D}_{train}|} \log L(\theta, \mathcal{D}_{train}).$$

Note that:

- l is a **deterministic** value (the value of a sum or integral).
- \hat{l} is a **random variable** (it is a weighted sum of r.v.).

We have some useful properties for this estimate \hat{l} :

- **Unbiased** estimate: $E_{p_{data}}[\hat{l}] = l$.
- **Convergence**: $\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n \log(p_{\theta}(\mathbf{x}^i)) = l$.

- **Variance** proportional to the inverse of the number of samples $\frac{1}{n}$:

$$\text{var} \left[\hat{l} \right] = \frac{1}{n} E_{\mathbf{x} \sim p_{data}} \left[(\log(p_{\theta}(\mathbf{x})) - l)^2 \right].$$

Example

A **biased coin** with head (H) and tail (T).

$$\mathcal{D}_{train} = \{H, H, H, T, H, T, H, H\},$$

p_{θ} would be a **Bernoulli distribution** with parameter θ (probability of having an H). How to choose θ according to what we have seen?

We will have:

$$p_{\theta}(\mathbf{x} = H) = \theta, \tag{3}$$

$$p_{\theta}(\mathbf{x} = T) = 1 - \theta, \tag{4}$$

and we can evaluate the likelihood of \mathcal{D}_{train} .

$$L(\theta, \mathcal{D}_{train}) = \theta^6(1 - \theta)^2,$$

and

$$\log L(\theta, \mathcal{D}_{train}) = 6 \log \theta + 2 \log(1 - \theta).$$

More generally, log-likelihood function

$$\log L(\theta, \mathcal{D}_{train}) = \#heads \log \theta + \#tails \log(1 - \theta).$$

Differentiating the expression above:

$$\frac{\partial \log L}{\partial \theta} = \frac{\#heads(1 - \theta) - \#tails\theta}{\theta(1 - \theta)} = \frac{\#heads - \theta(\#heads + \#tails)}{\theta(1 - \theta)},$$

that cancels at $\theta^* = \frac{\#heads}{\#heads + \#tails}$. Note that

$$\frac{\partial^2 \log L}{\partial \theta^2} = -\frac{\#heads}{\theta^2} - \frac{\#tails}{(1 - \theta)^2} < 0,$$

hence we have a maximum. So, in the example $\theta^* = \frac{3}{4}$.

Now, in general, things are not that easy!

$$\theta^* = \arg \max_{\theta} \log L(\theta, \mathcal{D}_{train}).$$

with

$$\log L(\theta, \mathcal{D}_{train}) = \sum_{\mathbf{x} \in \mathcal{D}_{train}} \log(p_{\theta}(\mathbf{x})),$$

but in most cases:

- it does not have a **closed form solution**;
- it is not even convex (so **its optimization is not trivial at all**).

The most common option is to use **gradient ascent methods**:

- Initialize $\theta^{(0)}$ at random.
- Compute $\nabla_{\theta} \log L(\theta, \mathcal{D}_{train})$ (e.g., by **back-propagation**).
- Apply gradient ascent:

$$\theta^{(k+1)} = \theta^{(k)} + \alpha_k \nabla_{\theta} \log L(\theta, \mathcal{D}_{train}).$$

*

In general, gives **decent results**.

What if \mathcal{D}_{train} is very large? $\log L(\theta, \mathcal{D}_{train})$ may be very expensive to compute:

$$\log L(\theta, \mathcal{D}_{train}) = n \sum_{\mathbf{x} \in \mathcal{D}_{train}} \frac{1}{n} \log(p_{\theta}(\mathbf{x})) = n E_{\mathbf{x} \in \mathcal{D}_{train}} [\log(p_{\theta}(\mathbf{x}))].$$

Hence we can use again Monte-Carlo:

- with one sample (**stochastic gradient**):

$$\log L(\theta, \mathcal{D}_{train}) \approx n \log(p_{\theta}(\mathbf{x})) \text{ for } \mathbf{x} \sim \mathcal{D}_{train},$$

- with a subset of \mathcal{D}_{train} (**mini-batch stochastic gradient**)

$$\log L(\theta, \mathcal{D}_{train}) \approx \sum_{\mathbf{x} \sim \mathcal{D}_{train}} \log(p_{\theta}(\mathbf{x})).$$

If the space for the parameterized p.d.f p_{θ} is limited, then you may have a lot of training data from p_{train} , but the **limitations of the model itself** p_{θ} make that you cannot fit well p_{data} : **bias**.

On the opposite, if your p_{θ} are very complex and expressive, then to fit them well you need an **infinite amount of data** in \mathcal{D}_{train} , which you **do not have**. Hence, given the **finiteness** of your training data, the fitting problem is an **ill-posed problem**: you may have lots of θ that fit well \mathcal{D}_{train} : **variance**.

In choosing one class of model (i.e. the design of p_{θ}), you may have to balance the two aspects:

- Distribution space too small: your model will **underfit** (strong bias).
- Distribution space too large: your model will **overfit** (strong variance).

To avoid overfitting:

- propose **simpler models**! (smaller networks, parameter sharing between parts of the networks);
- **regularization**:

$$\theta^* = \arg \max_{\theta} \log L(\theta, \mathcal{D}_{train}) + R(\theta);$$

- monitor the overfitting with a **validation dataset**.

0.2 Application to fitting the parameters of an auto-regressive models

In the case of auto-regressive models,

$$p(\mathbf{x}; \theta) = \prod_{i=1}^m p(\mathbf{x}_i | \mathbf{x}_{1:i-1}; \theta_i).$$

Then

$$L(\theta, \mathcal{D}_{train}) = \prod_{\mathbf{x} \sim \mathcal{D}_{train}} (p_{\theta}(\mathbf{x})) \quad (5)$$

$$= \prod_{j=1}^n \prod_{i=1}^m p(\mathbf{x}_i^j | \mathbf{x}_{1:i-1}^j; \theta_i). \quad (6)$$

and

$$\log L(\theta, \mathcal{D}_{train}) = \sum_{j=1}^n \sum_{i=1}^m \log p(\mathbf{x}_i^j | \mathbf{x}_{1:i-1}^j; \theta_i).$$

$$\nabla_{\theta_i} \log L(\theta, \mathcal{D}_{train}) = \sum_{j=1}^n \nabla_{\theta_i} \log p(\mathbf{x}_i^j | \mathbf{x}_{1:i-1}^j; \theta_i)$$

Depending on the implementation (shared parameters vs. non-shared parameters) you may estimate $\nabla_{\theta_i} \log L(\theta, \mathcal{D}_{train})$ independently.

Last, as we saw above, the stochastic gradient is used when the amount of training data is very large

$$\nabla_{\theta_i} \log L(\theta, \mathcal{D}_{train}) \approx n E_{\mathbf{x} \sim \mathcal{D}_{train}} [\nabla_{\theta_i} \log p(\mathbf{x}_i | \mathbf{x}_{1:i-1}; \theta_i)].$$

0.3 Conditional generative models

In some situations, we are interested in generating some \mathbf{x} given the values of another variable \mathbf{y} .

For example:

- generation of images “like the ones of ImageNet”, conditionally to the class;
- generation of images given a caption.

We could do it by modeling:

$$p(\mathbf{x}, \mathbf{y}; \theta),$$

but in fact it won’t be necessary since the value of \mathbf{y} will be fixed.

Instead:

$$p(\mathbf{x}|\mathbf{y}; \theta).$$

Hence we maximize:

$$L(\theta, \mathcal{D}_{train}) = \prod_{\{\mathbf{x}, \mathbf{y}\} \sim \mathcal{D}_{train}} \log(p_{\theta}(\mathbf{x}|\mathbf{y})). \quad (7)$$

and we can follow exactly the same process as above, except that all our components (giving the parameters of the conditional distributions) will now be depending on \mathbf{y} (i.e., \mathbf{y} is an “input” of the neural networks).

Conditional Image Generation with PixelCNN Decoders, A. Van Den Oord, N. Kalchbrenner, O. Vinyals, L. Espeholt, A. Graves, K. Kavukcuoglu. Proceedings of the NIPS’16: Proceedings of the 30th International Conference on Neural Information Processing Systems.

Conditional PixelCNN:

- Overall, same principle as PixelCNN.
- Modified version of the masked CNN
- Written as a function of \mathbf{y} .

Each layer has the form

$$\mathbf{z} = \tanh(\mathbf{W}_{k,f} \otimes \mathbf{x} + \mathbf{V}_{k,f}\mathbf{y}) \odot \sigma(\mathbf{W}_{k,g} \otimes \mathbf{x} + \mathbf{V}_{k,g}\mathbf{y}).$$

In the case of classes, \mathbf{y} can be encoded as a one-hot vector.



Lhasa Apso (dog)

[From original paper]



[From original paper]