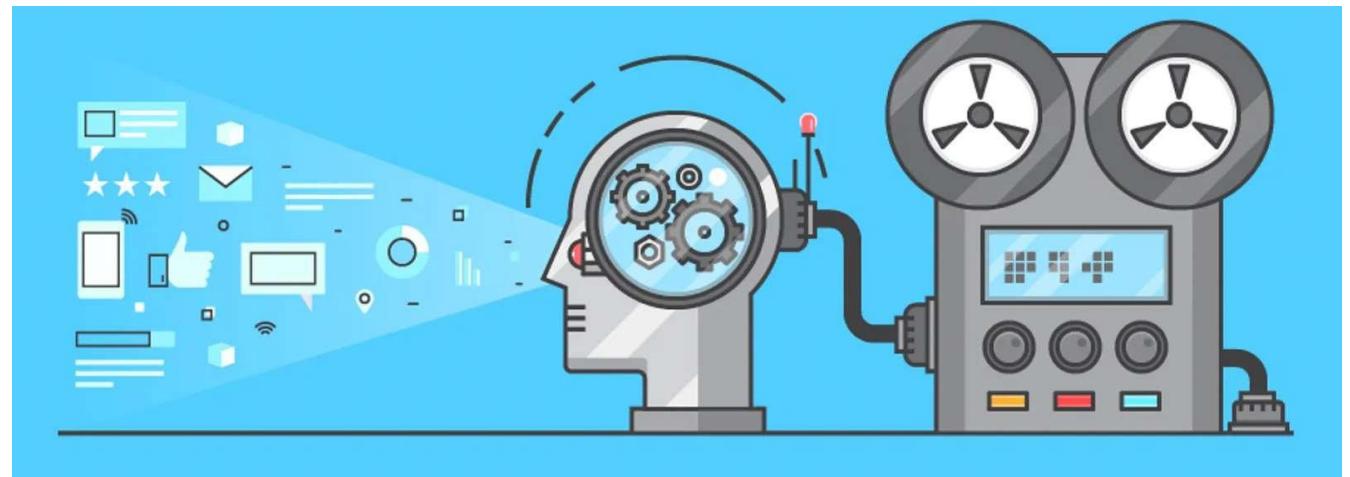


Deep Learning and Neural Networks

Topic 5:
Convolutional
Neural
Networks



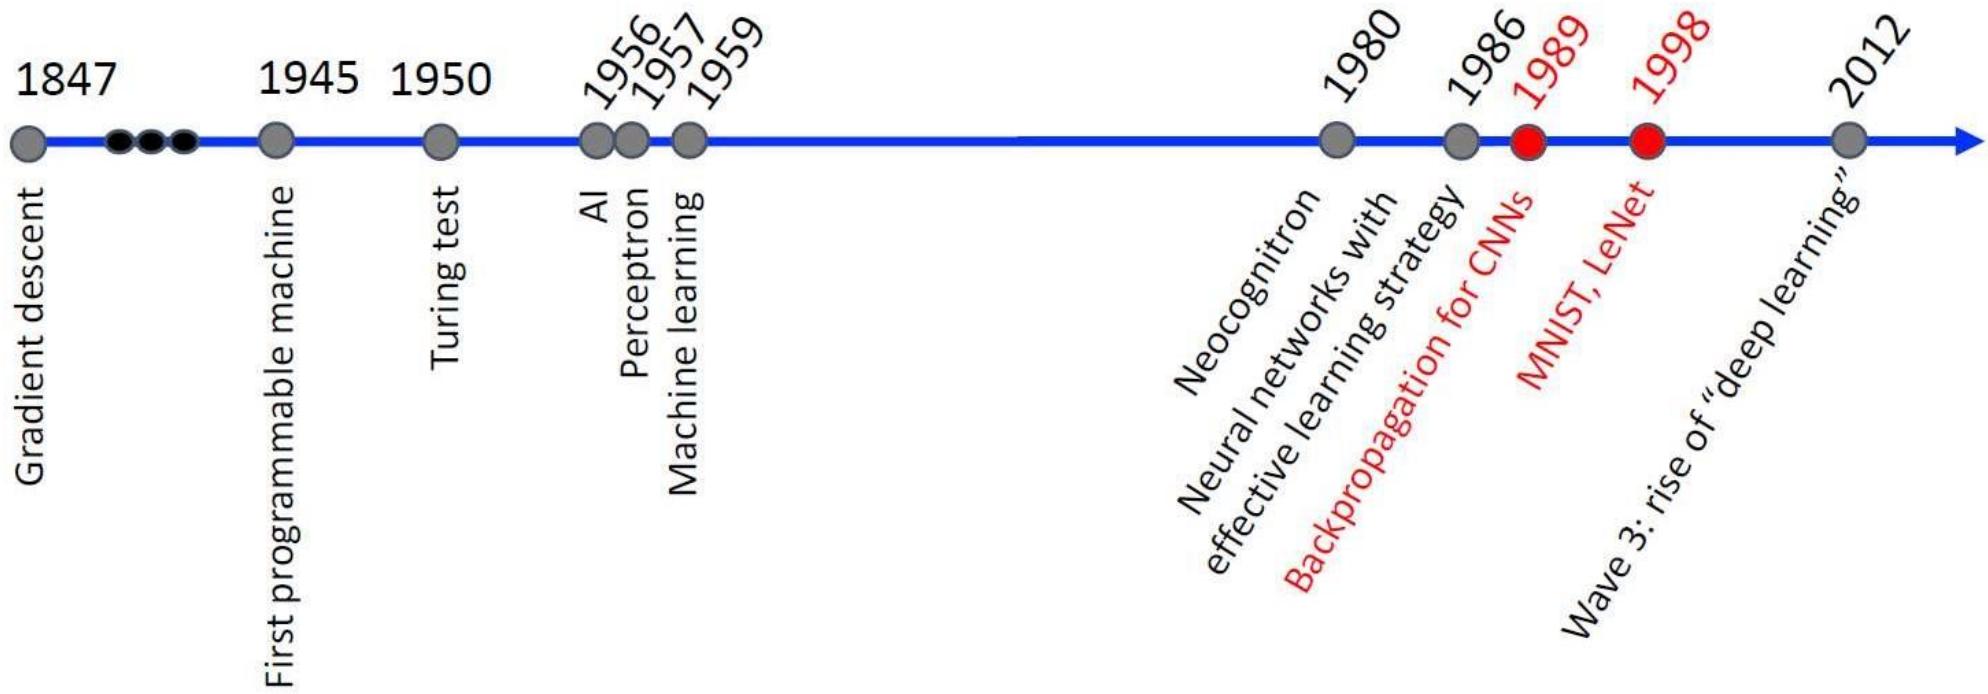
Ricardo Abel Espinosa Loera, McS

Researcher in DL & Computer Vision

BACKGROUND

Neural Networks for images

Historical Context: Inspiration



BACKGROUND

Neural Networks for images



Which are the different tasks in visual object recognition?



Image from [Fei Fei Li, 2011]

classification

Does this image contain a car? [yes/no]

detection

Does this image contains a car? [yes? where]

detection+

What objects does this image contain?
[where]

segmentation

What is the exact location? [object contour]

semantic interpretation

Semantic and geometric attributes
estimation

Which are the different tasks in visual object recognition?



classification

Does this image contain a car? [yes/no]

detection

Does this image contains a car? [yes? where]

detection+

What objects does this image contain?
[where]

segmentation

What is the exact location? [object contour]

semantic interpretation

Semantic and geometric attributes
estimation

Which are the different tasks in visual object recognition?



classification

Does this image contain a car? [yes/no]

detection

Does this image contain a car? [yes? where]

detection+

What objects does this image contain?
[where]

segmentation

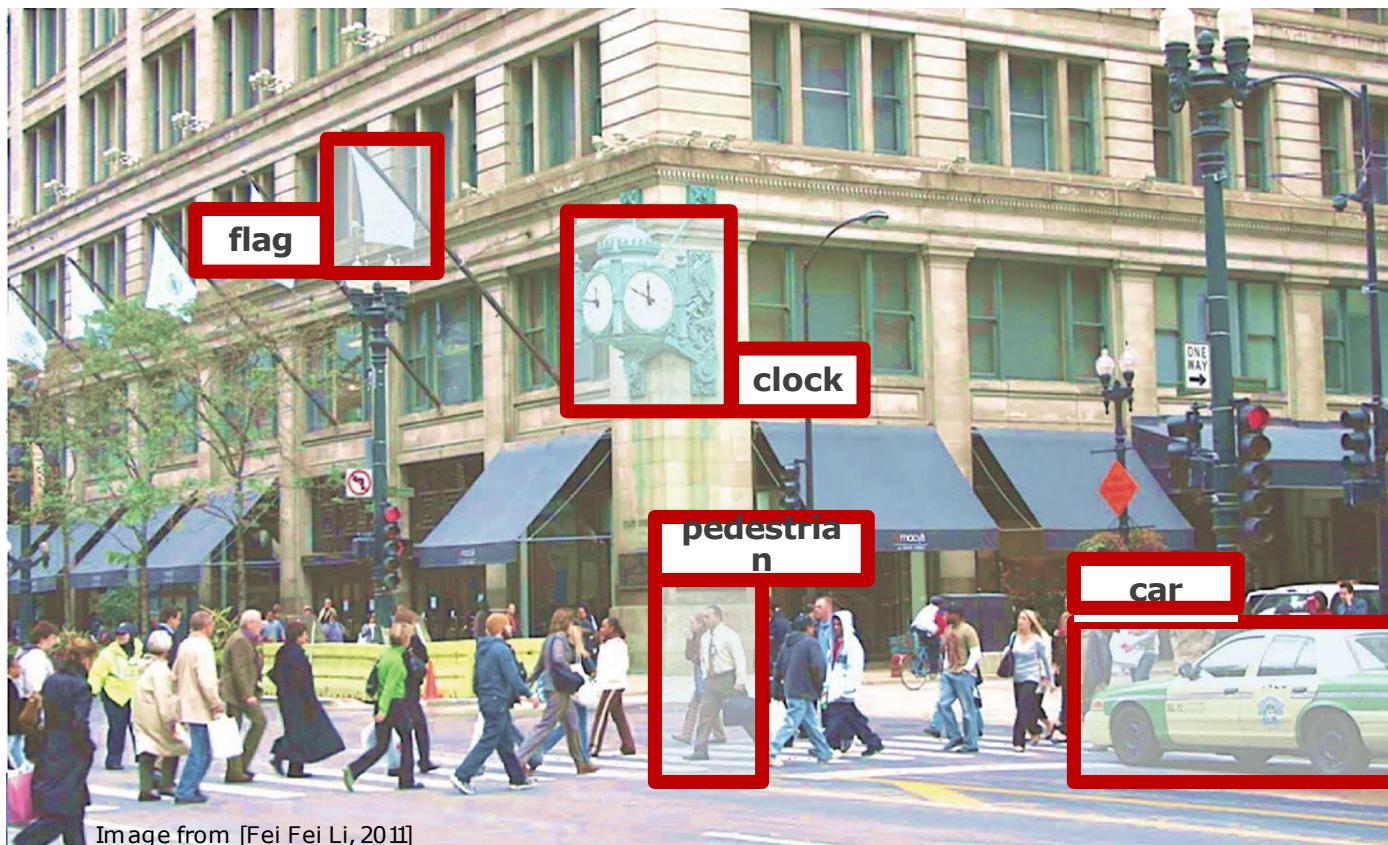
What is the exact location? [object contour]

semantic interpretation

Semantic and geometric attributes
estimation

Image from [Fei Fei Li, 2011]

Which are the different tasks in visual object recognition?



classification

Does this image contain a car? [yes/no]

detection

Does this image contains a car? [yes? where]

detection+

What objects does this image contain?
[where]

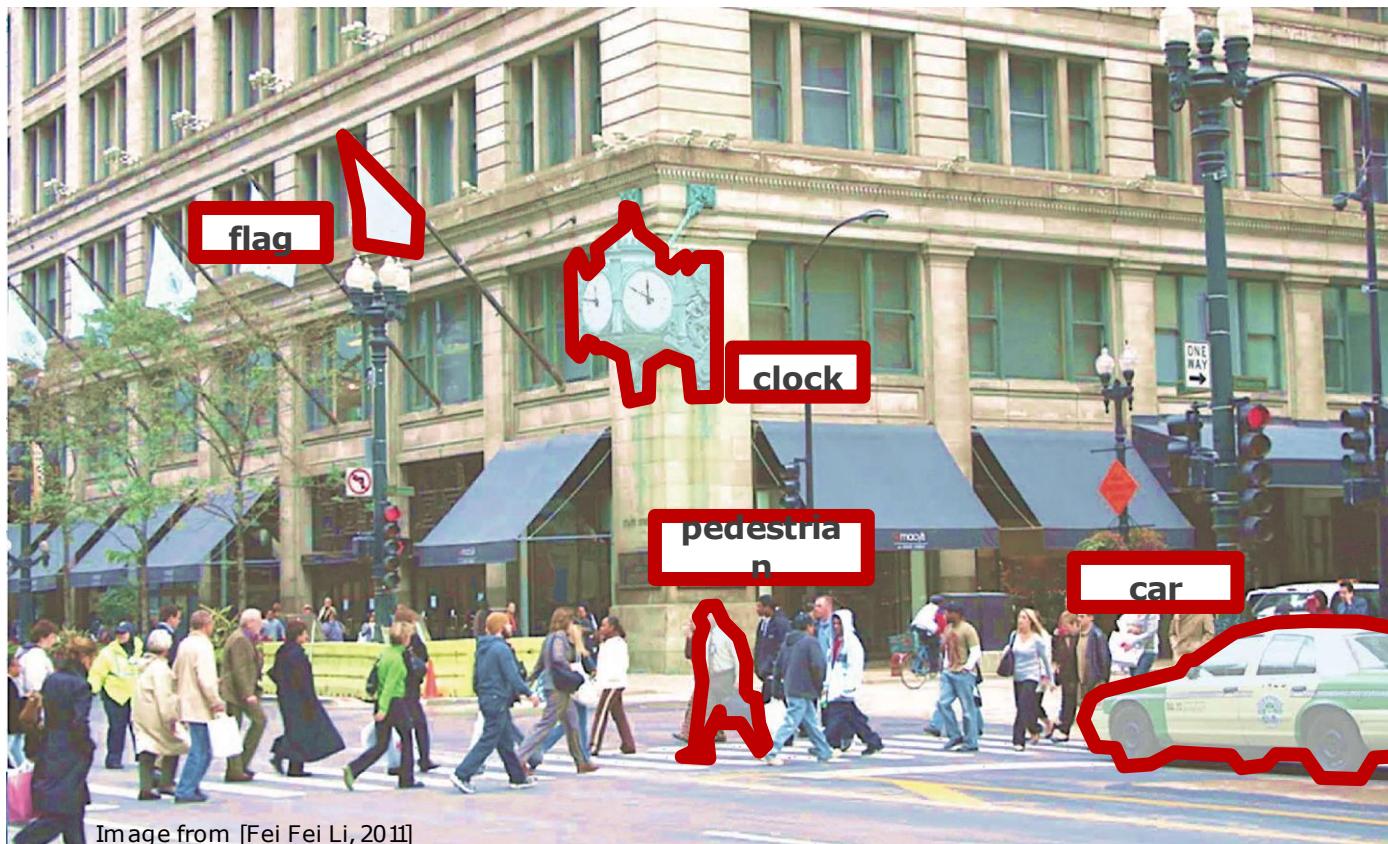
segmentation

What is the exact location? [object contour]

semantic interpretation

Semantic and geometric attributes
estimation

Which are the different tasks in visual object recognition?



classification

Does this image contain a car? [yes/no]

detection

Does this image contains a car? [yes? where]

detection+

What objects does this image contain?
[where]

segmentation

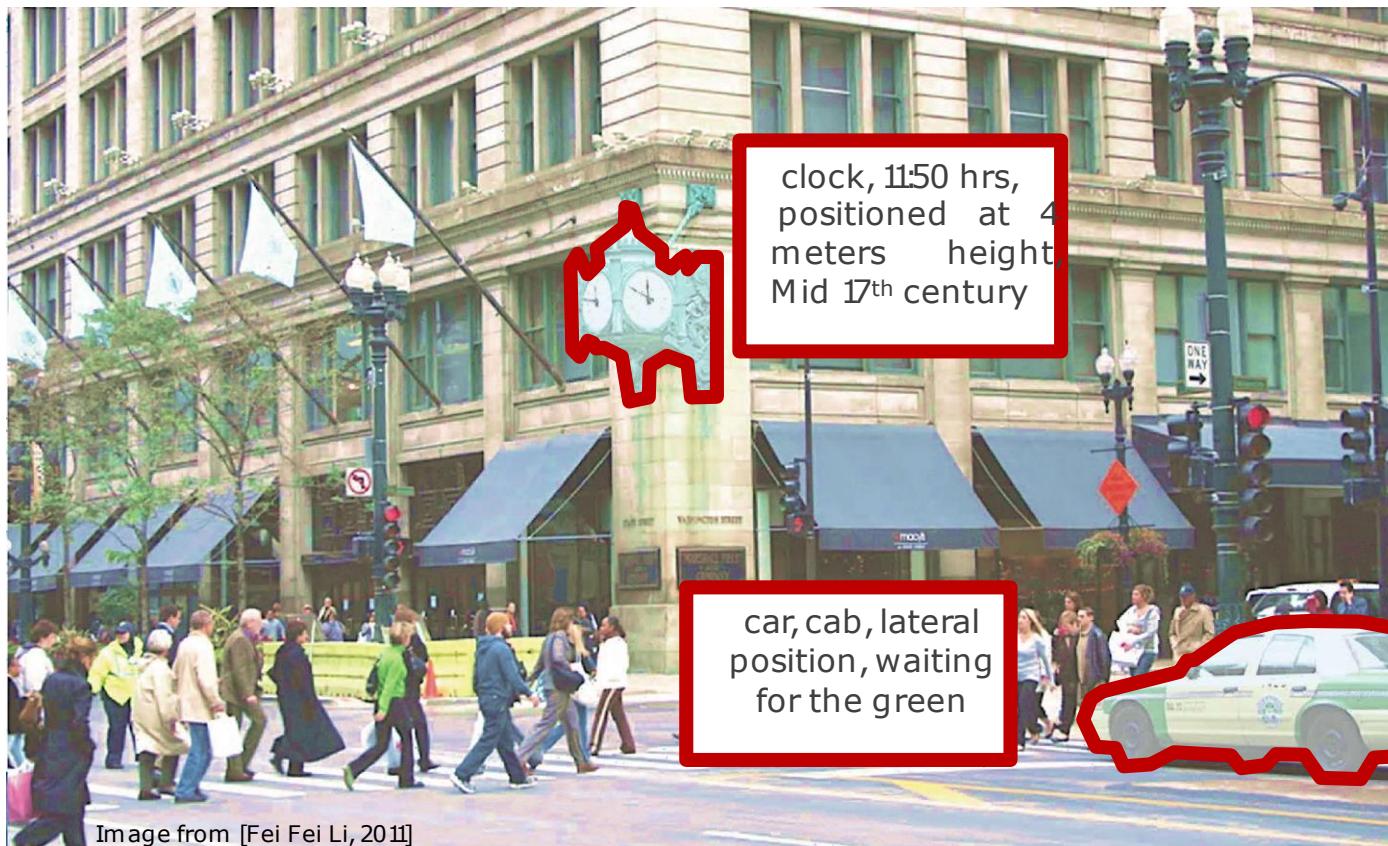
What is the exact location? [object contour]

semantic interpretation

Semantic and geometric attributes
estimation

Image from [Fei Fei Li, 2011]

Which are the different tasks in visual object recognition?



classification

Does this image contain a car? [yes/no]

detection

Does this image contains a car? [yes? where]

detection+

What objects does this image contain?
[where]

segmentation

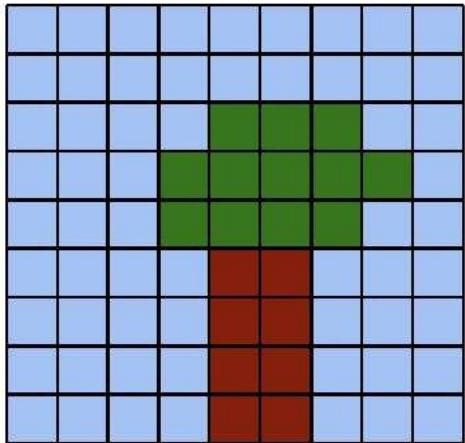
What is the exact location? [object contour]

semantic interpretation

Semantic and geometric attributes
estimation

BACKGROUND

Neural Networks for images



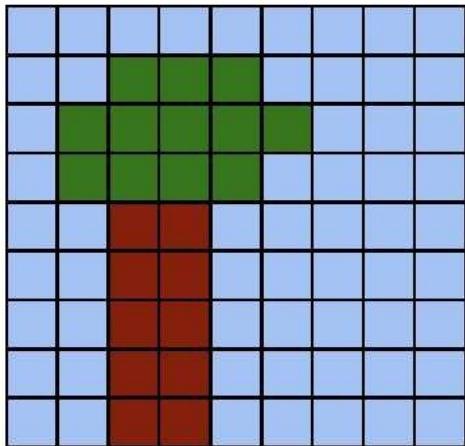
A digital image is a 2D **grid of pixels**.

A neural network expects a **vector of numbers** as input.



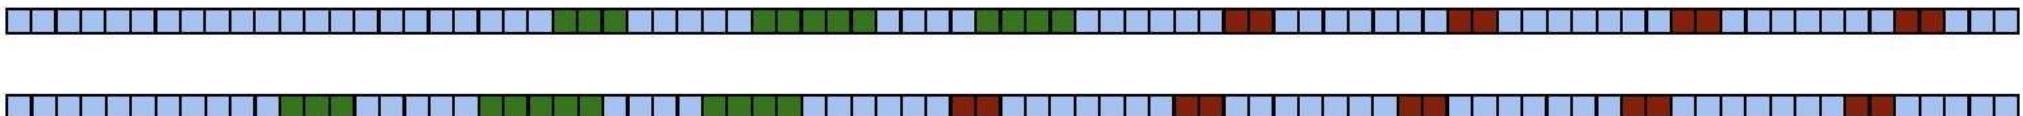
BACKGROUND

Neural Networks for images



A digital image is a **2D grid of pixels**.

A neural network expects a **vector of numbers** as input.



BACKGROUND

Locality and traslation invariance



Locality: nearby pixels are more strongly correlated

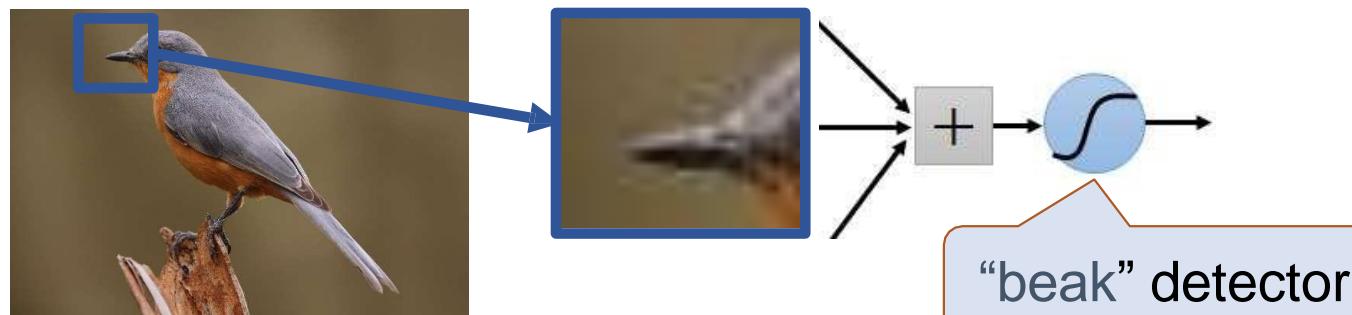
Translation invariance: meaningful patterns can occur anywhere in the image

BACKGROUND

Locality and traslation invariance

- Some patterns are much smaller than the whole image

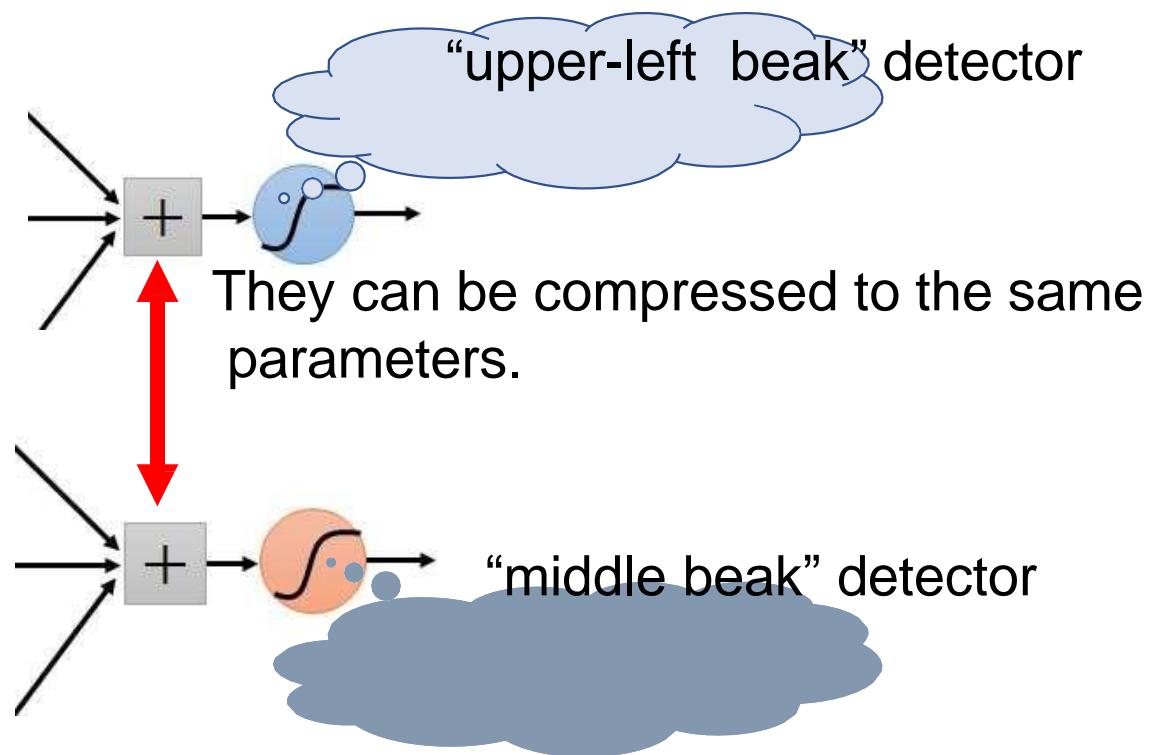
Can represent a small region with fewer parameters



BACKGROUND

Locality and traslation invariance

What about training a lot of such “small” detectors and each detector must “move around”.



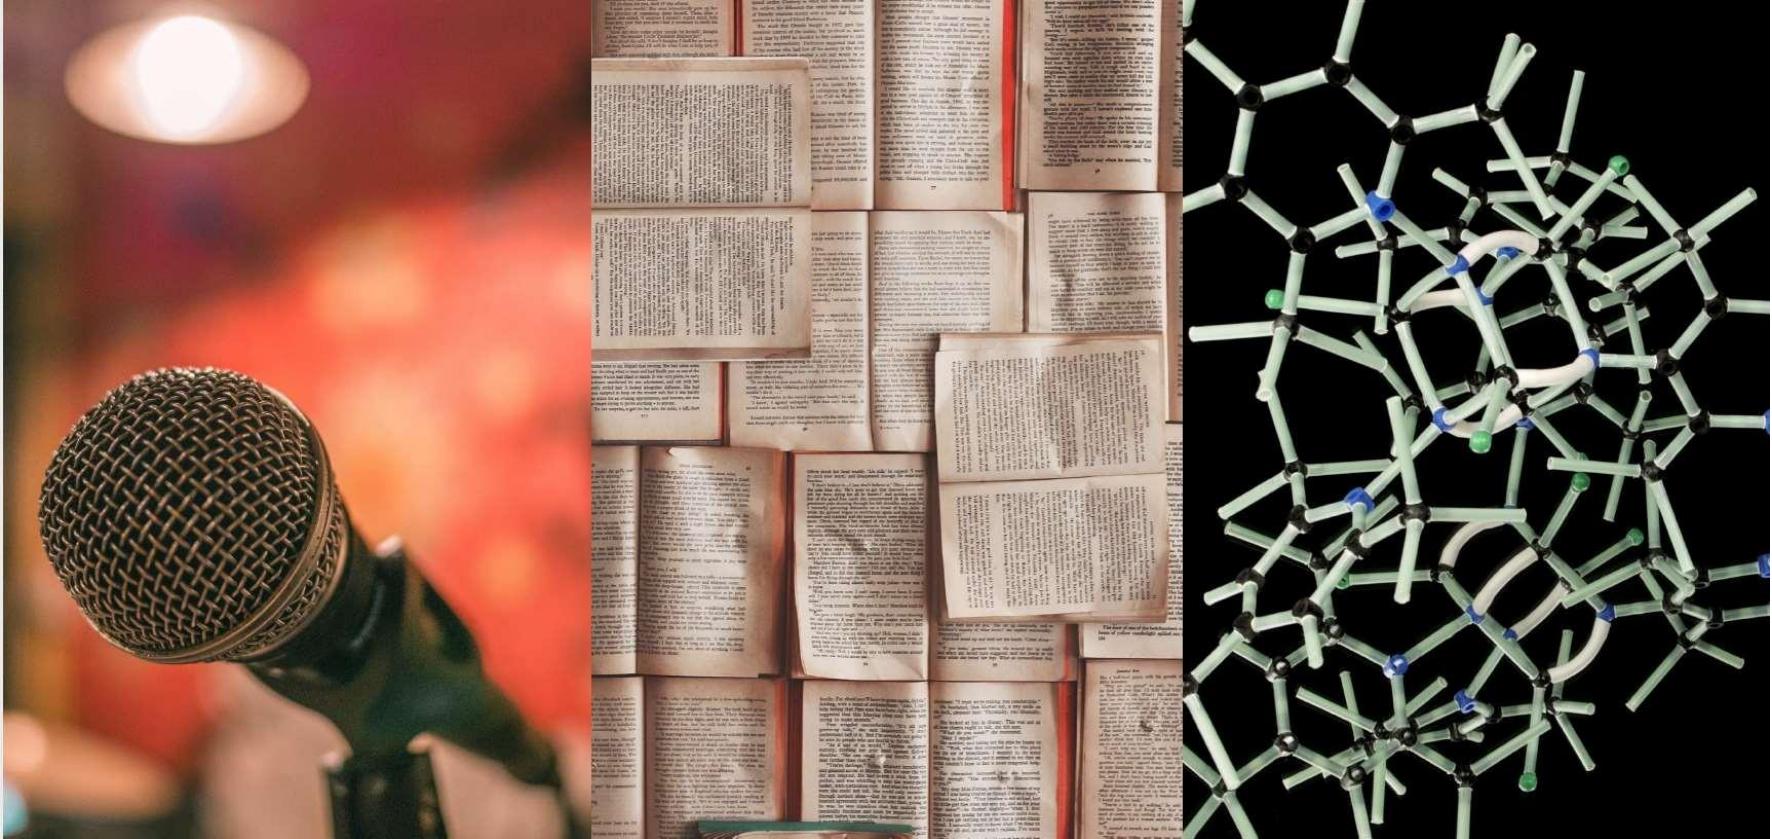
BACKGROUND

Locality and traslation invariance



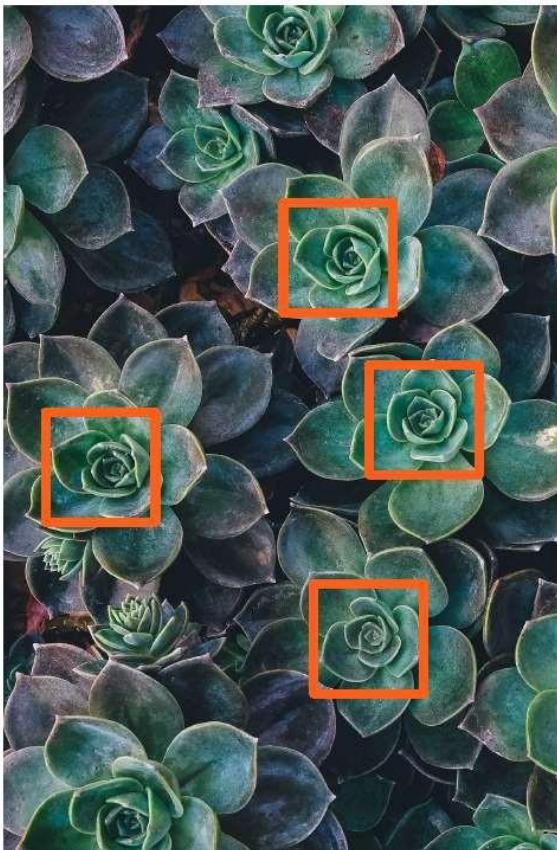
BACKGROUND

Locality and traslation invariance



BACKGROUND

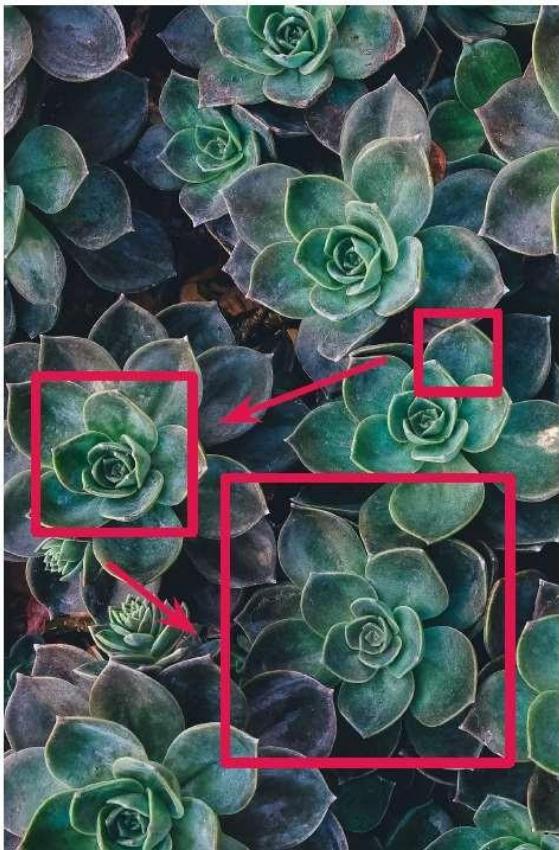
Taking Advantage of topological structure



Weight sharing: use the same network parameters to detect local patterns at many locations in the image

BACKGROUND

Taking Advantage of topological structure



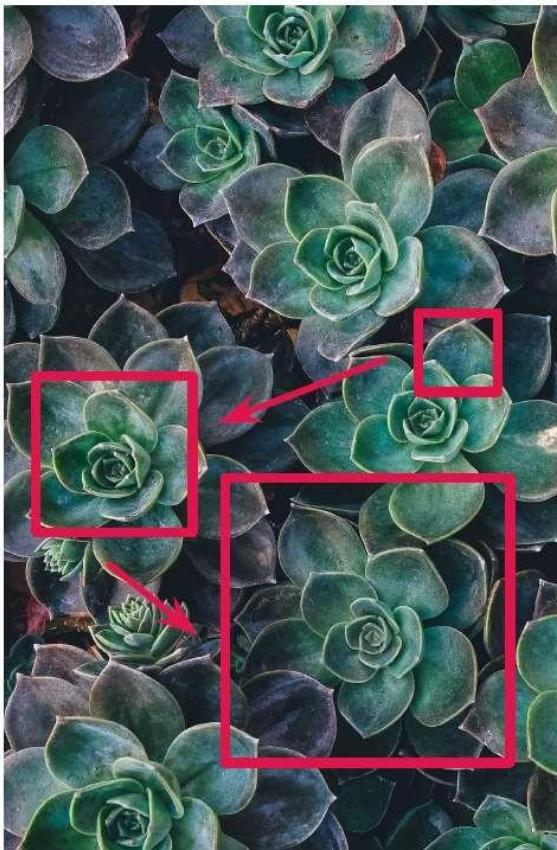
Weight sharing: use the same network parameters to detect local patterns at many locations in the image

Hierarchy: local low-level features are composed into larger, more abstract features



BACKGROUND

Taking Advantage of topological structure



Weight sharing: use the same network parameters to detect local patterns at many locations in the image

Hierarchy: local low-level features are composed into larger, more abstract features



edges and textures

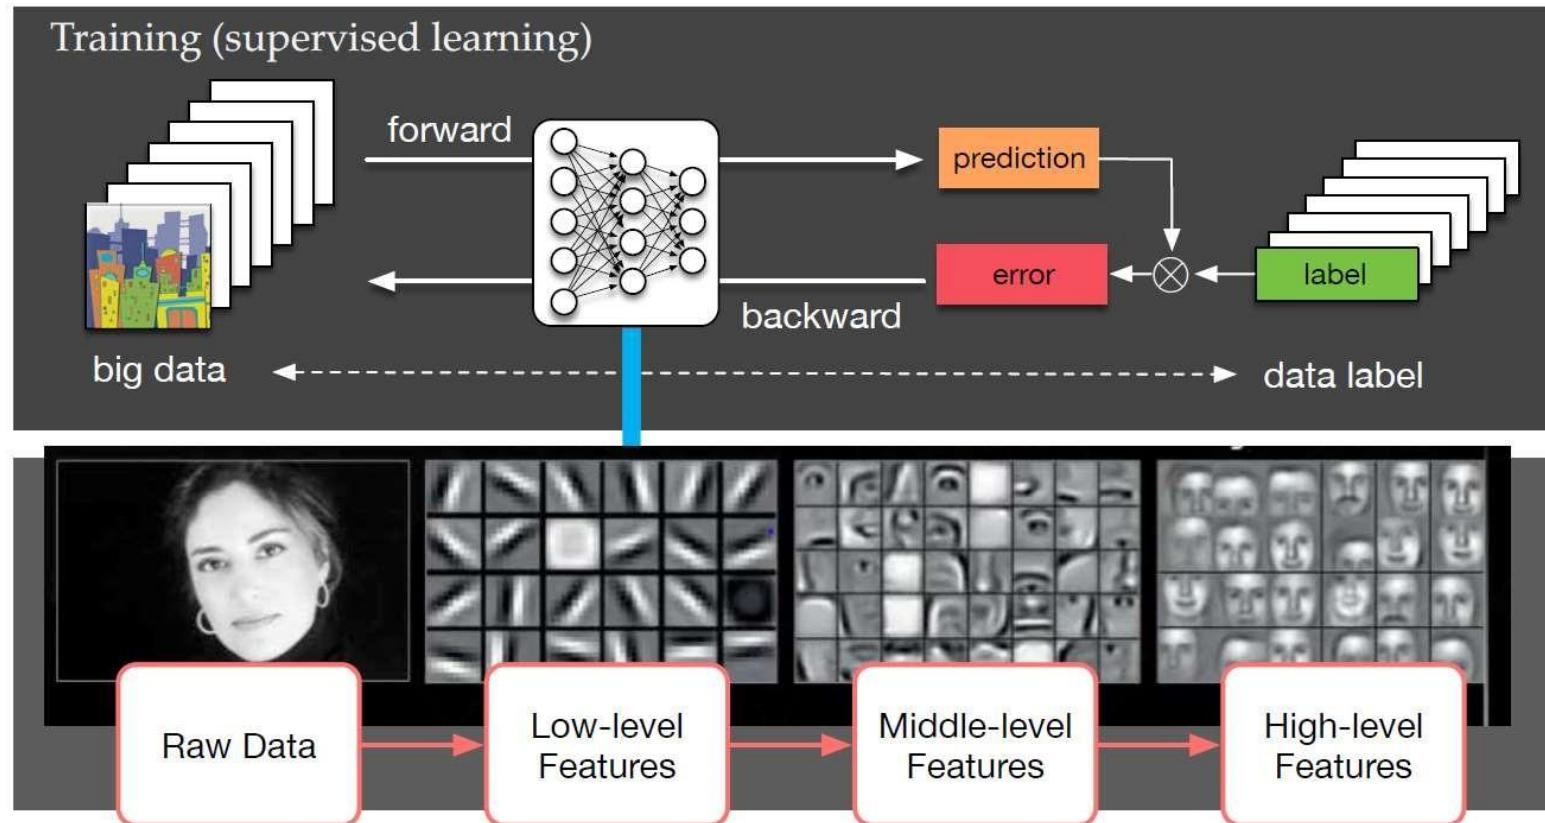


object parts

objects

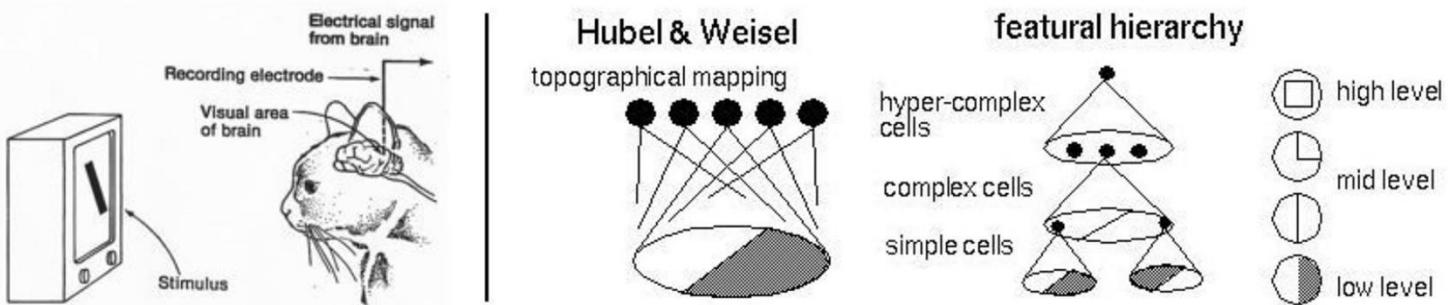
BACKGROUND

Taking Advantage of topological structure



BACKGROUND

experiments of Hubel and Wiesel



D.H. Hubel and T.N. Wiesel (1962) Receptive fields, binocular interaction and functional architecture in the cat's visual cortex *J Physiol.* 160(1): 106154.2.

BACKGROUND

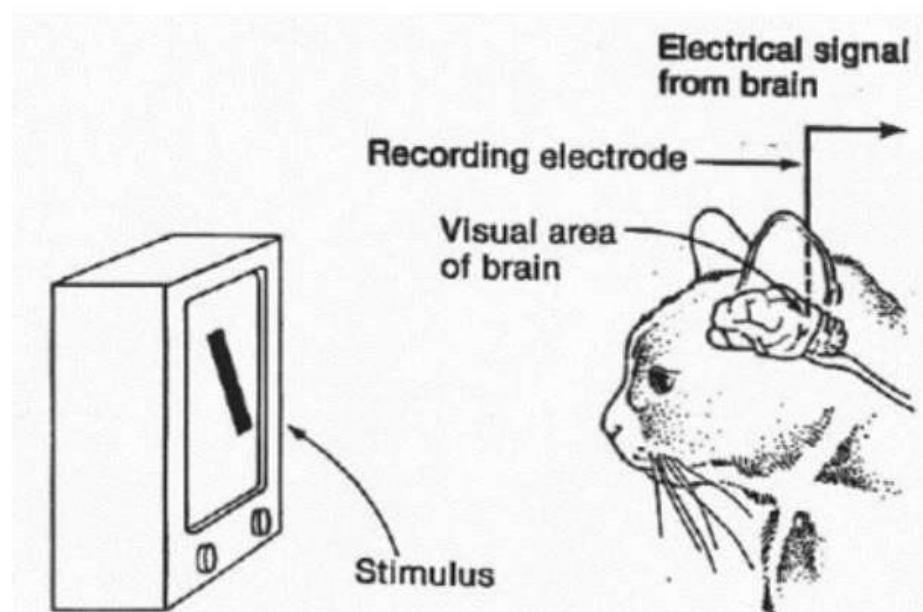
experiments of Hubel and Wiesel



BACKGROUND

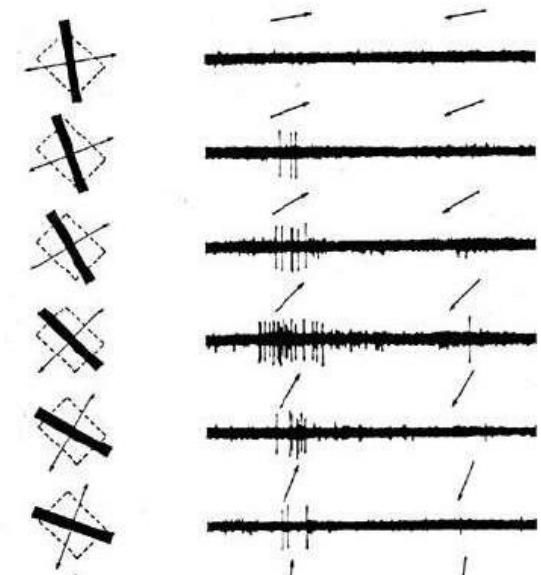
experiments of Hubel and Wiesel

Experiment Set-up:



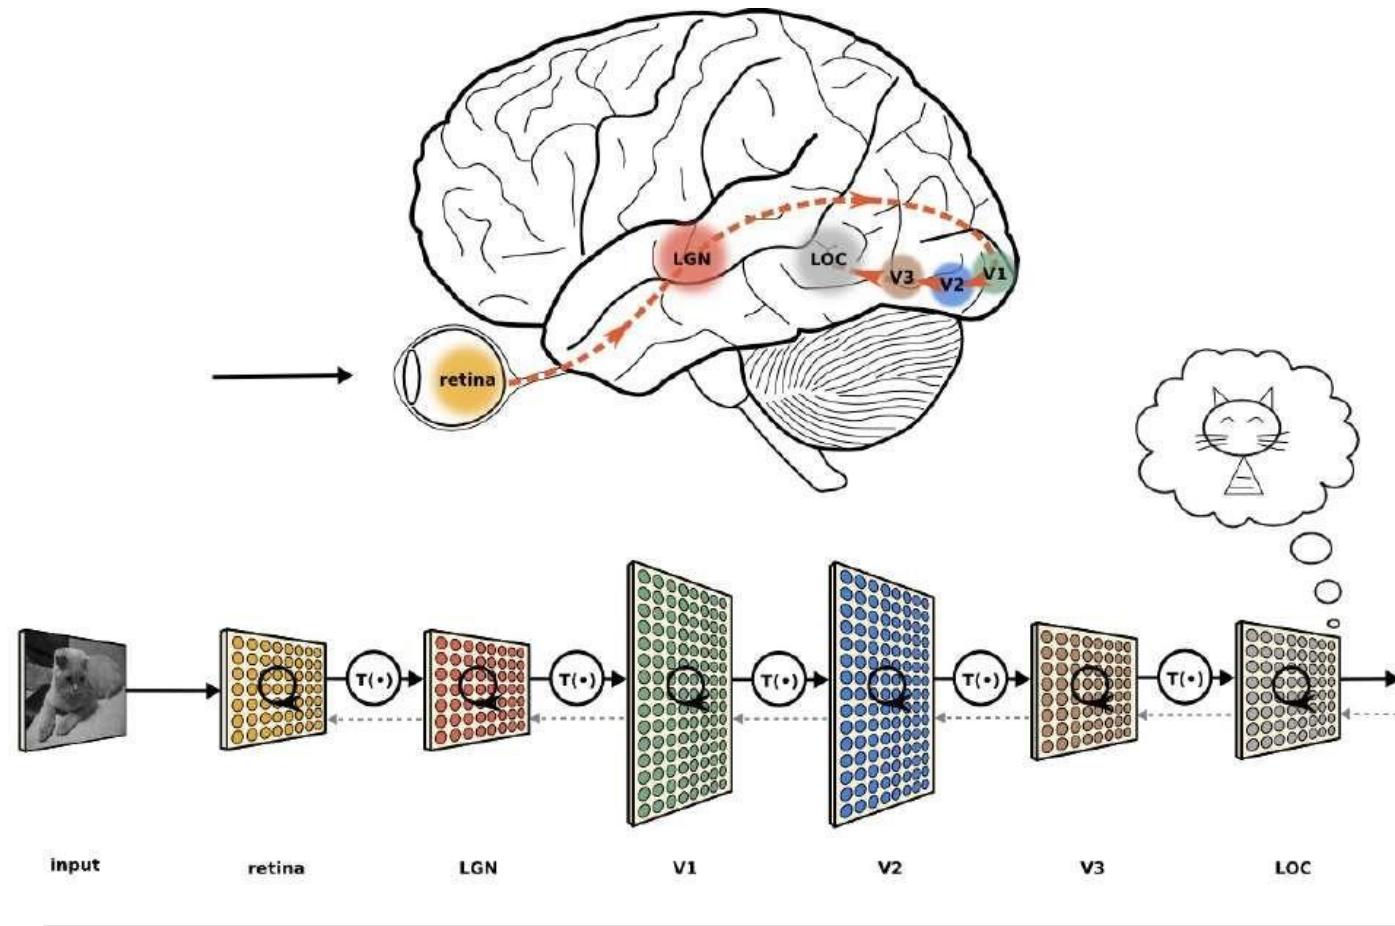
Key Finding: initial neurons responded strongly only when light was shown in certain orientations

V1 physiology:
direction
selectivity



BACKGROUND

experiments of Hubel and Wiesel

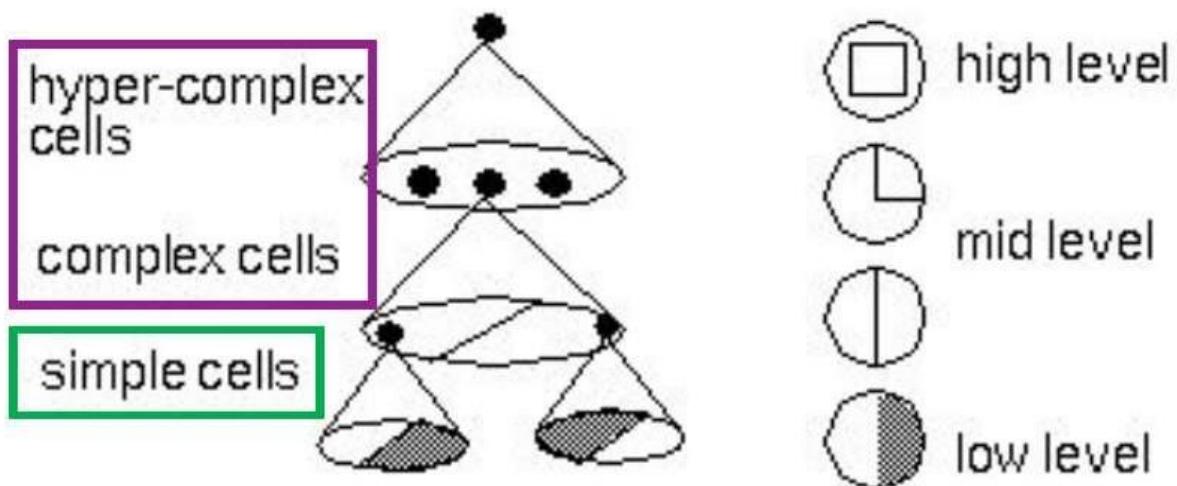


BACKGROUND

experiments of Hubel and Wiesel

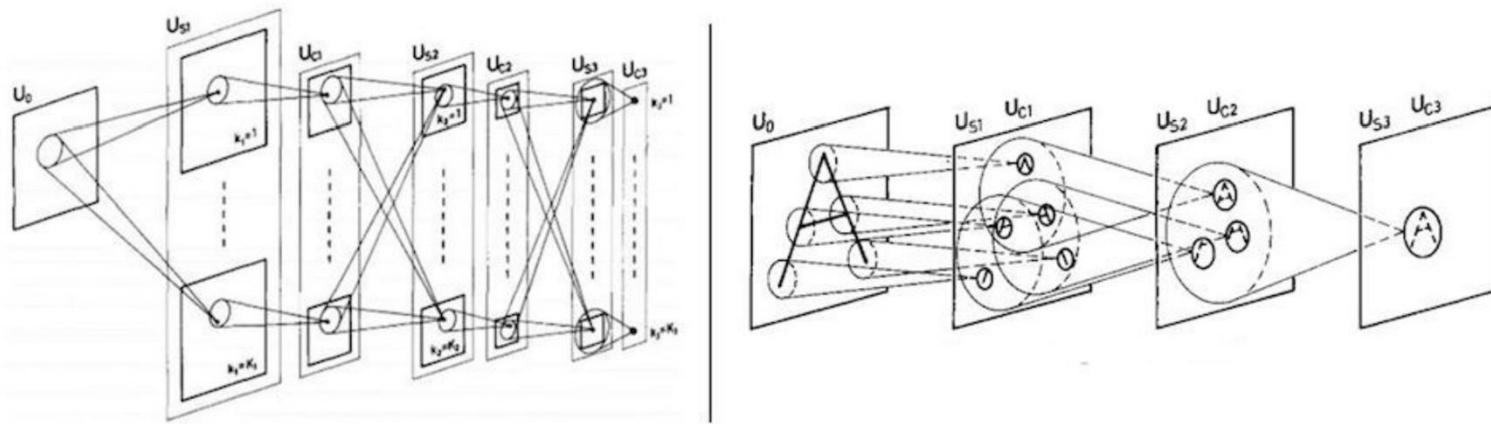
Key Idea: cells are organized as a hierarchy of feature detectors, with **higher level features** responding to patterns of activation in **lower level cells**

featural hierarchy



BACKGROUND

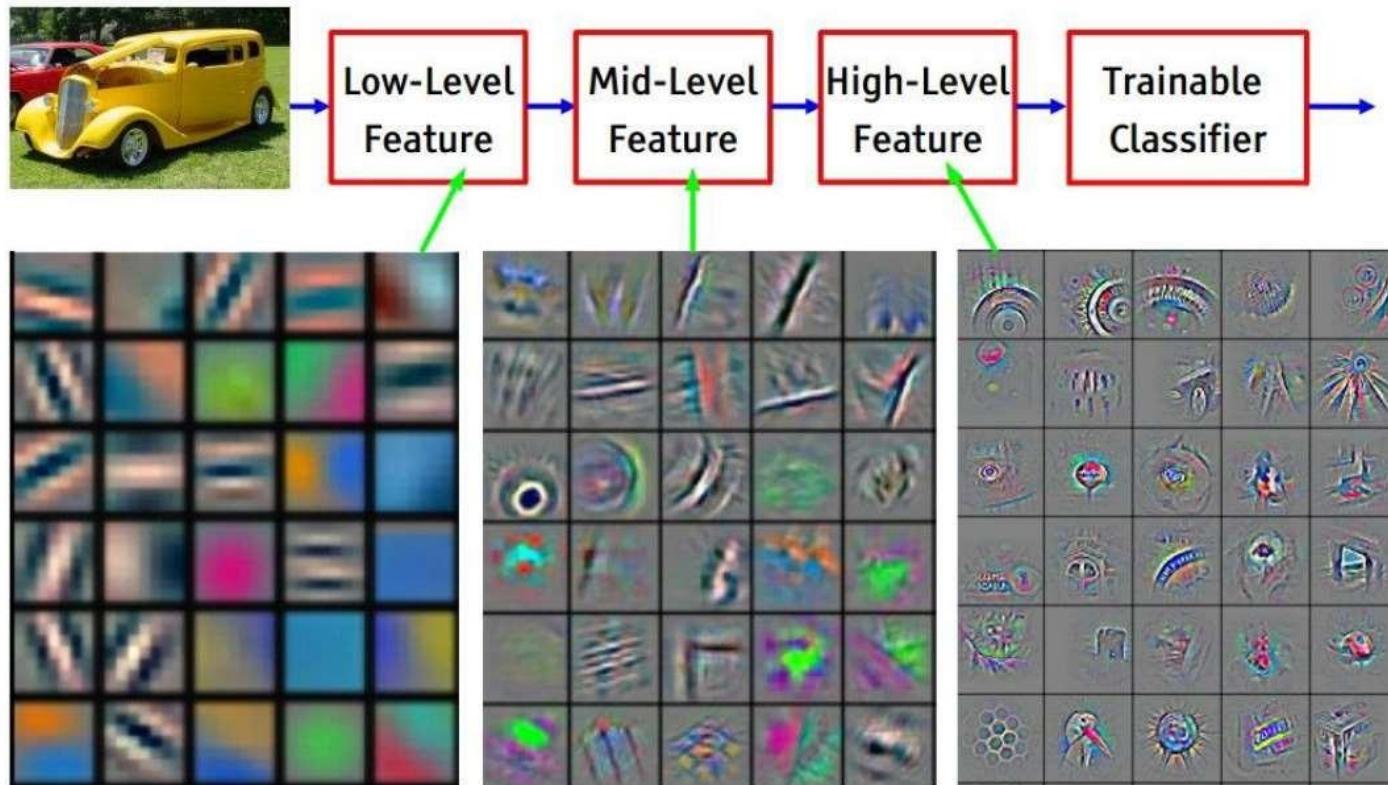
neocognitron



K. Fukushima (1980) Neocognitron: A Self-organizing Neural Network Model for a Mechanism of Pattern Recognition Unaffected by Shift in Position. *Biol. Cybernetics* 36, 193–202

BACKGROUND

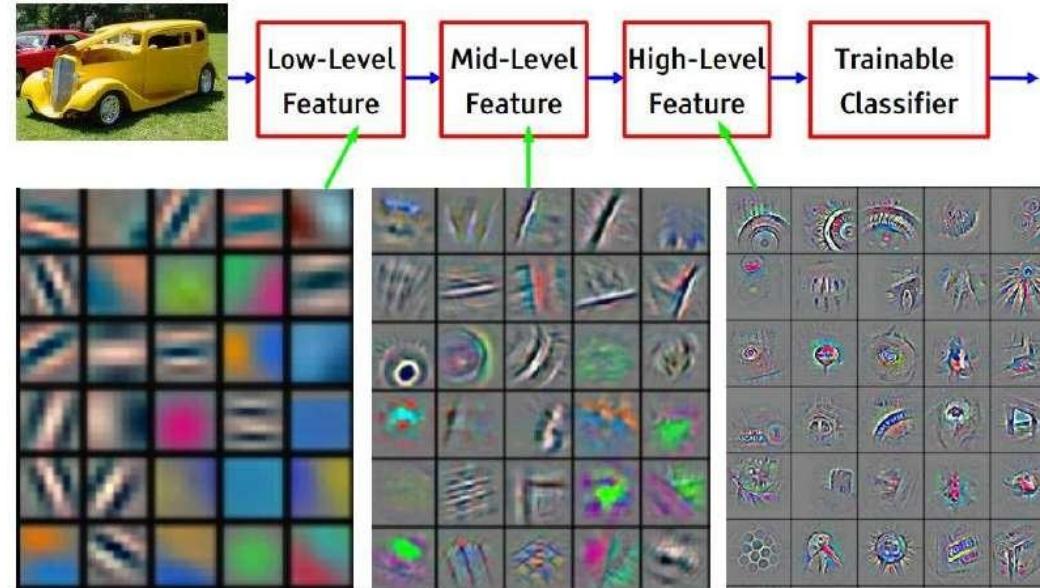
LeNet



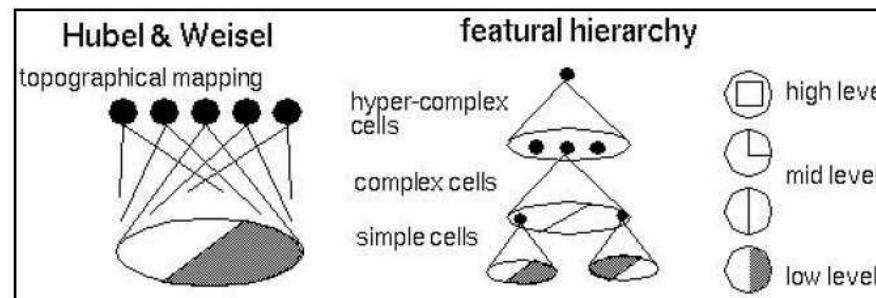
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

BACKGROUND

LeNet

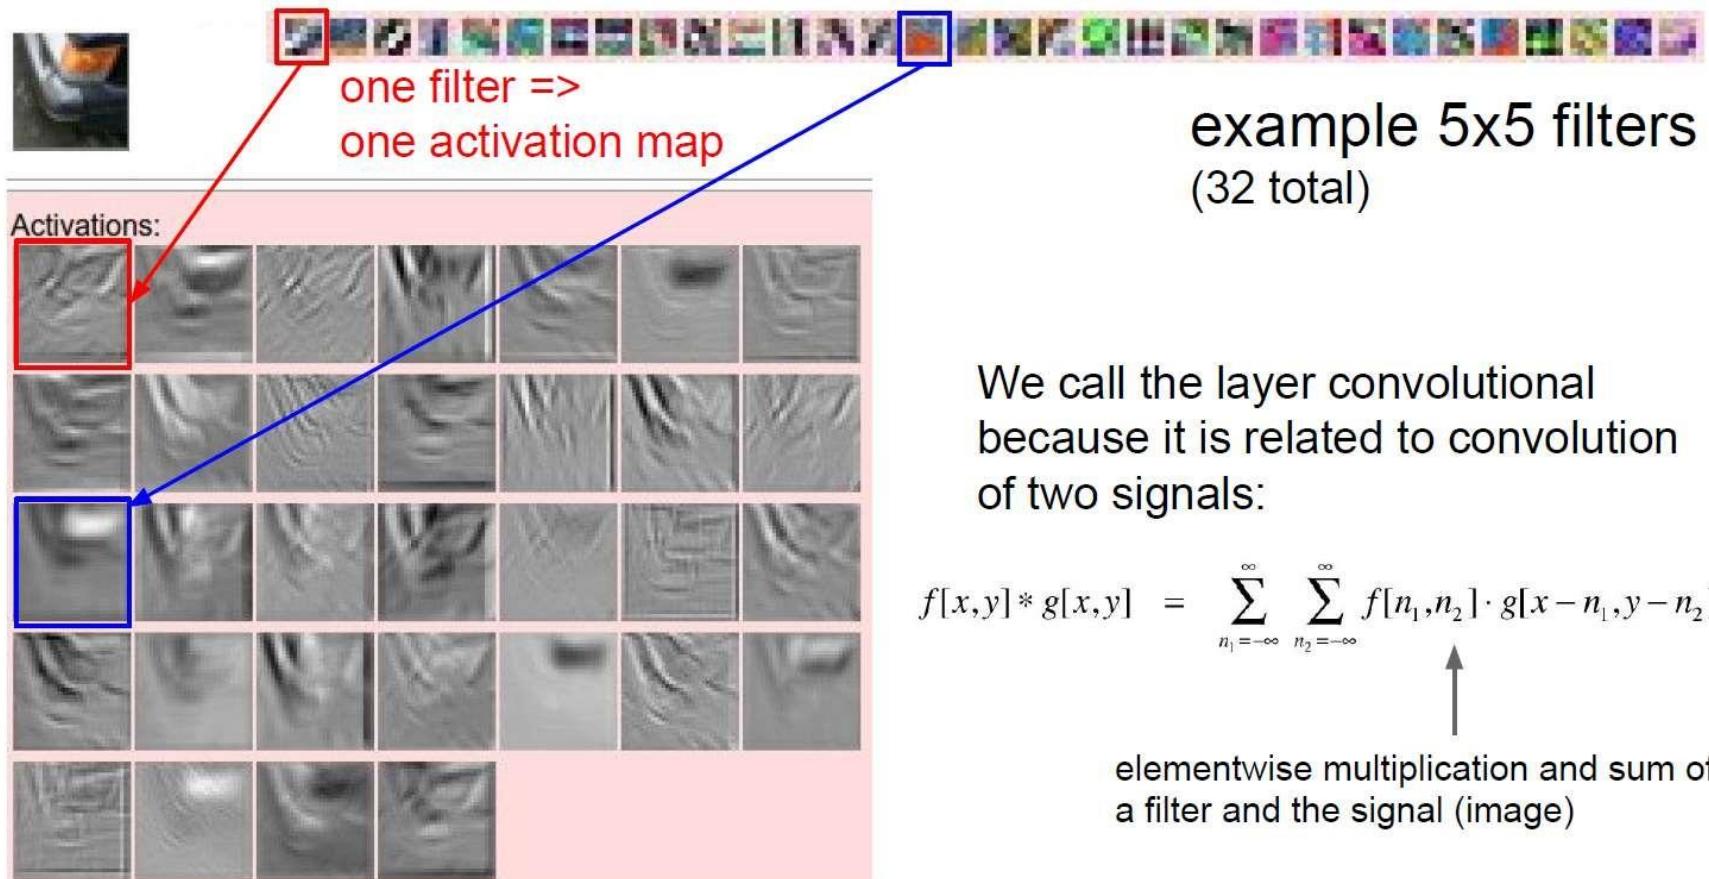


Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]



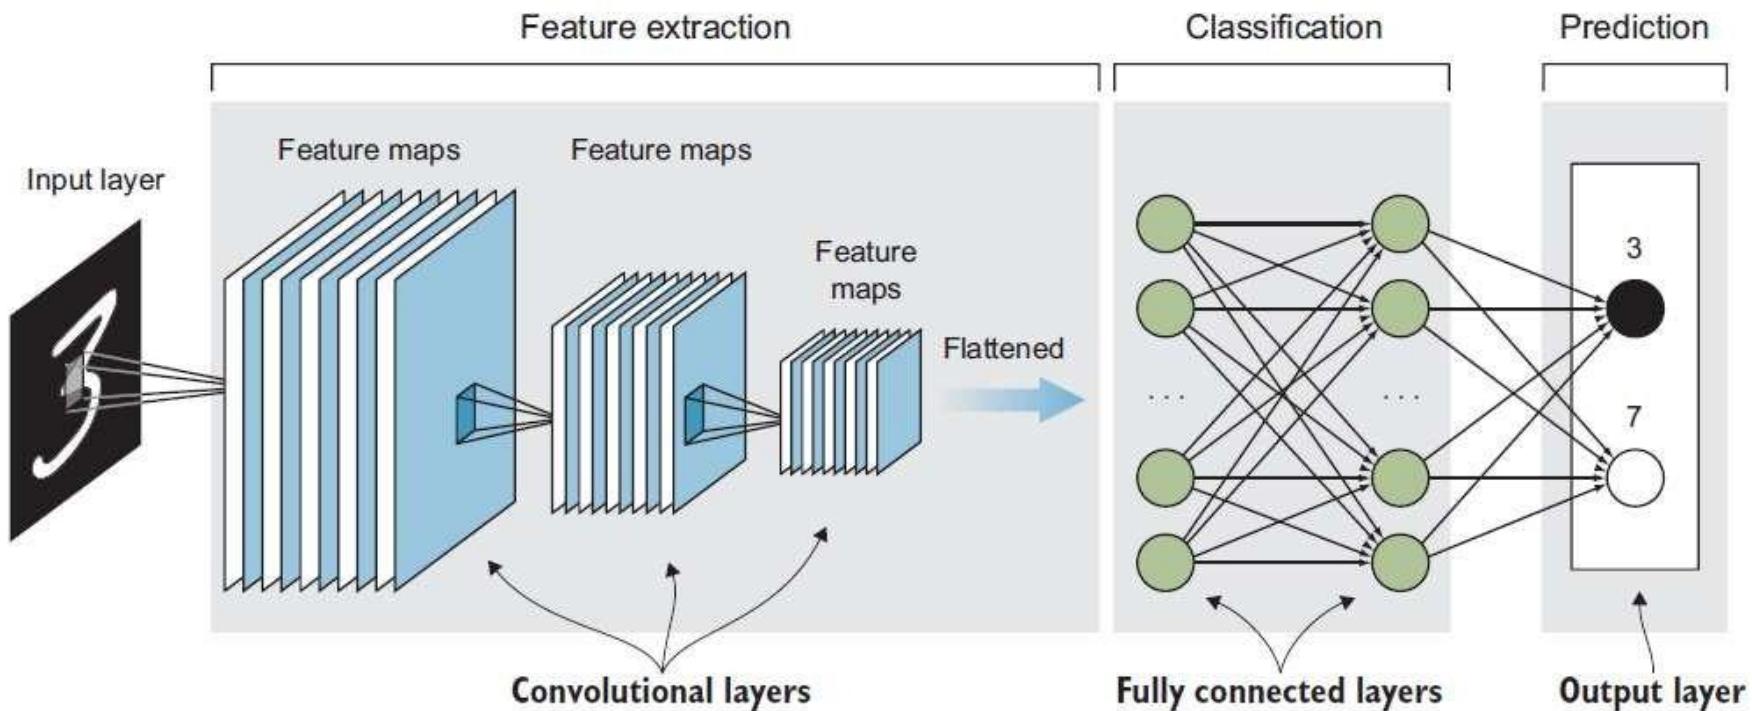
BACKGROUND

LeNet



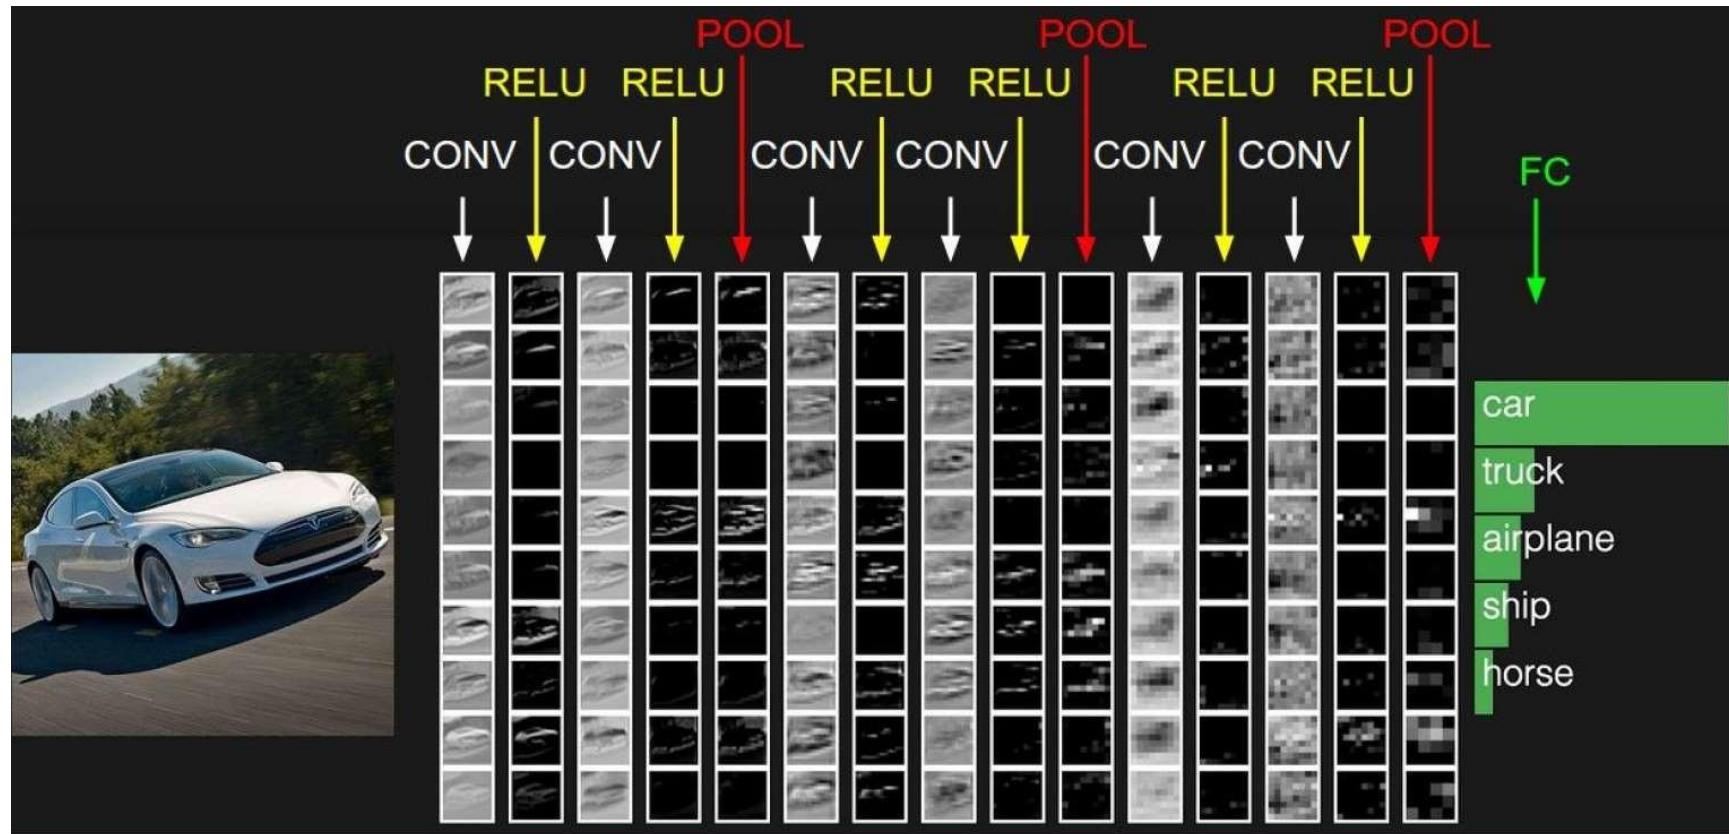
BACKGROUND

LeNet



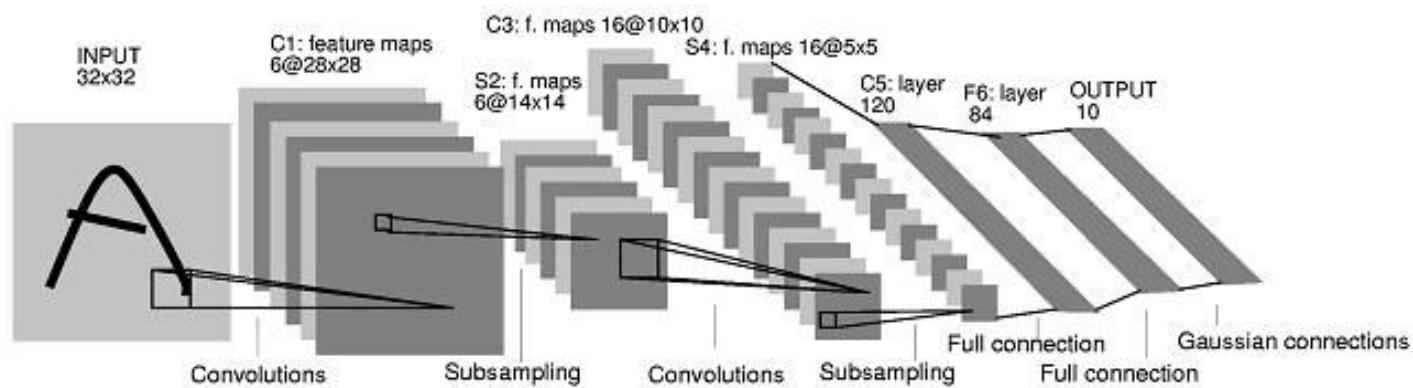
BACKGROUND

LeNet



BACKGROUND

LeNet



Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner (1998) Gradient-based learning applied to document recognition, *Proc. IEEE* 86(11): 2278-2324

BACKGROUND

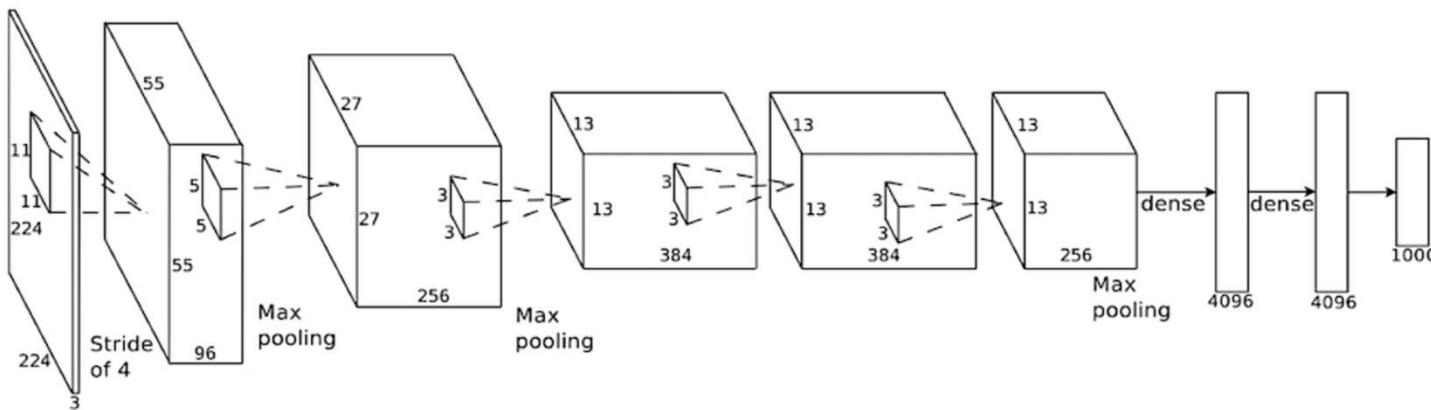
LeNet

LeNet-5 model summary

| Layer (type) | Output Shape | Param # |
|-------------------------------|--------------------|---------|
| <hr/> | | |
| conv2d_1 (Conv2D) | (None, 28, 28, 6) | 156 |
| average_pooling2d_1 (Average) | (None, 14, 14, 6) | 0 |
| <hr/> | | |
| conv2d_2 (Conv2D) | (None, 10, 10, 16) | 2416 |
| average_pooling2d_2 (Average) | (None, 5, 5, 16) | 0 |
| <hr/> | | |
| conv2d_3 (Conv2D) | (None, 1, 1, 120) | 48120 |
| flatten_1 (Flatten) | (None, 120) | 0 |
| <hr/> | | |
| dense_1 (Dense) | (None, 84) | 10164 |
| dense_2 (Dense) | (None, 10) | 850 |
| <hr/> | | |
| Total params: | 61,706 | |
| Trainable params: | 61,706 | |
| Non-trainable params: | 0 | |

BACKGROUND

AlexNet



A. Krizhevsky, I. Sutskever, G. E. Hinton (2012) ImageNet Classification with Deep Convolutional Neural Networks. Advances in Neural Information Processing Systems 25

BACKGROUND

Data Drives Research

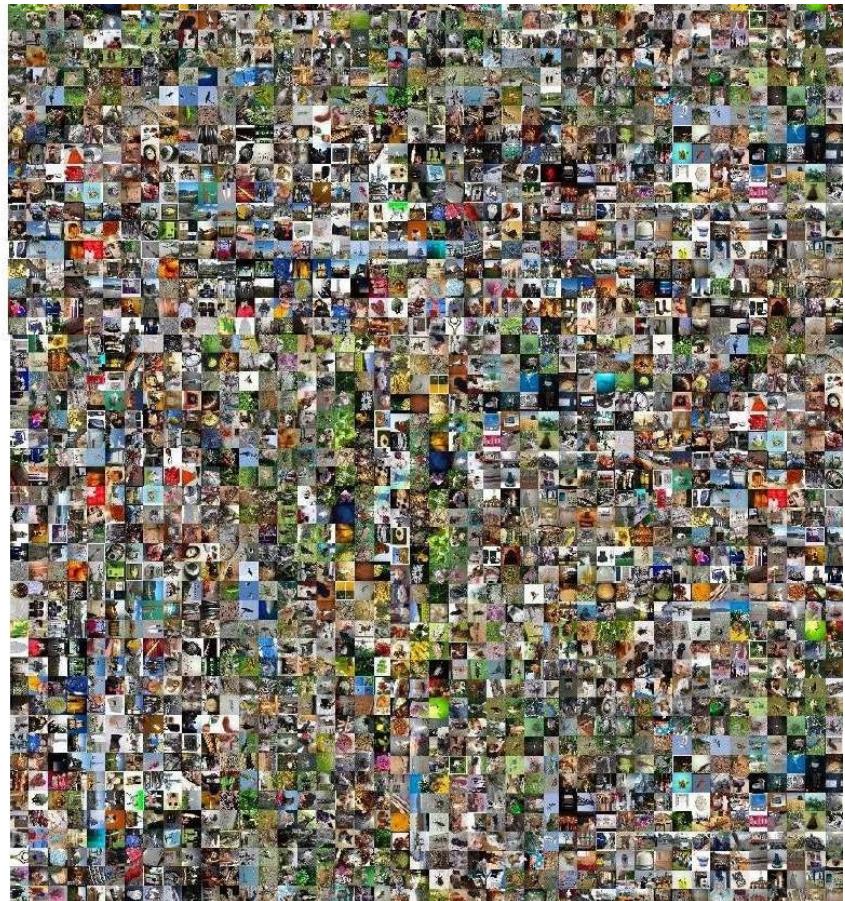
The ImageNet challenge

- Major computer vision benchmark
- Ran from 2010 to 2017
- 1.4M images, 1000 classes
- Image classification

Want to learn more?

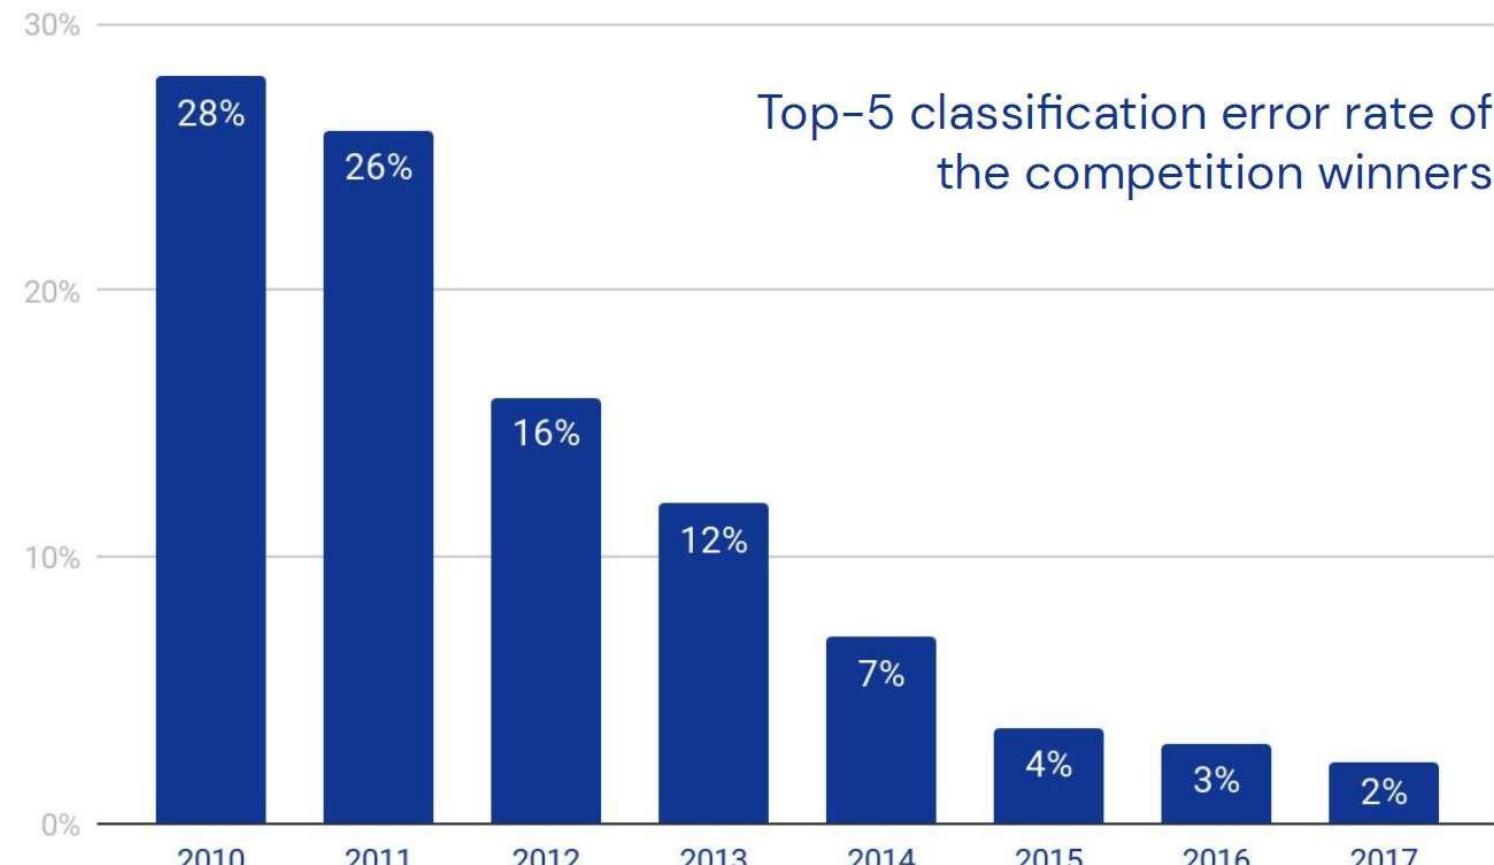


Russakovsky, Olga et al. *ImageNet Large Scale Visual Recognition Challenge* International Journal of Computer Vision 115.3 (2015)



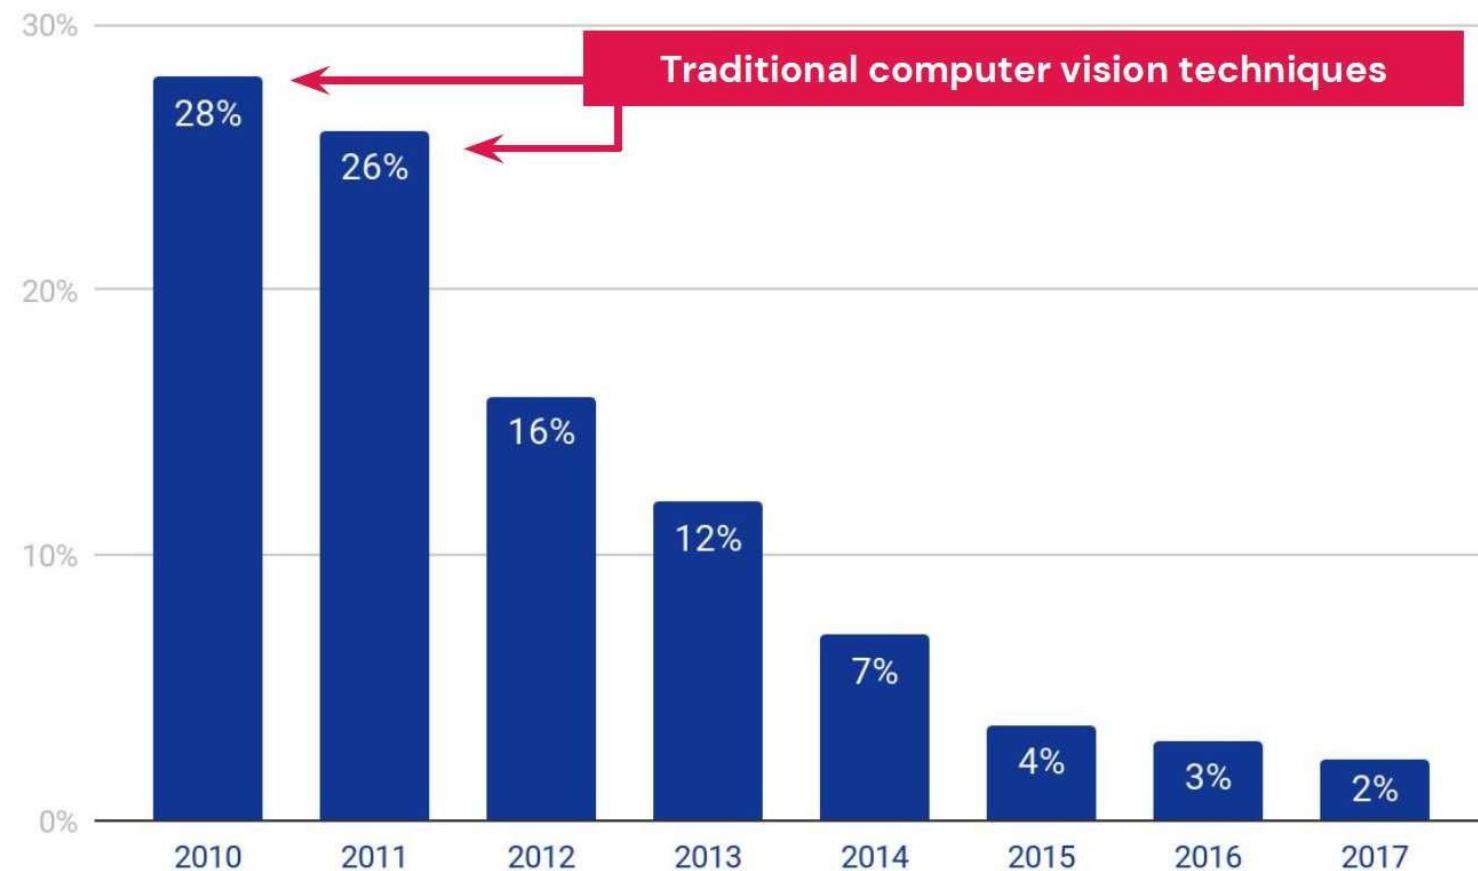
BACKGROUND

Data Drives Research



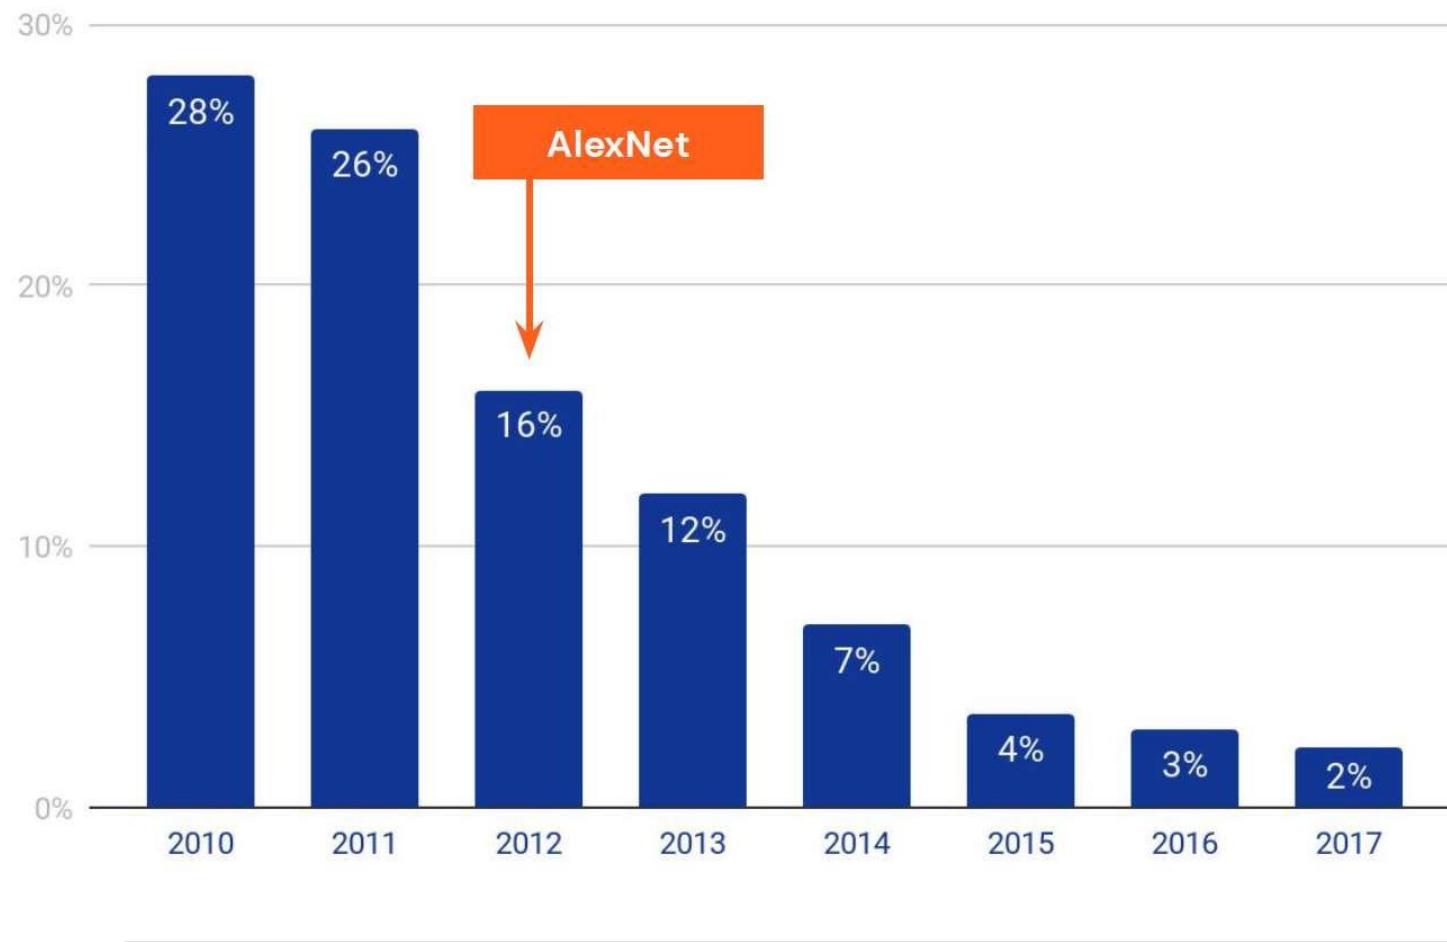
BACKGROUND

Data Drives Research



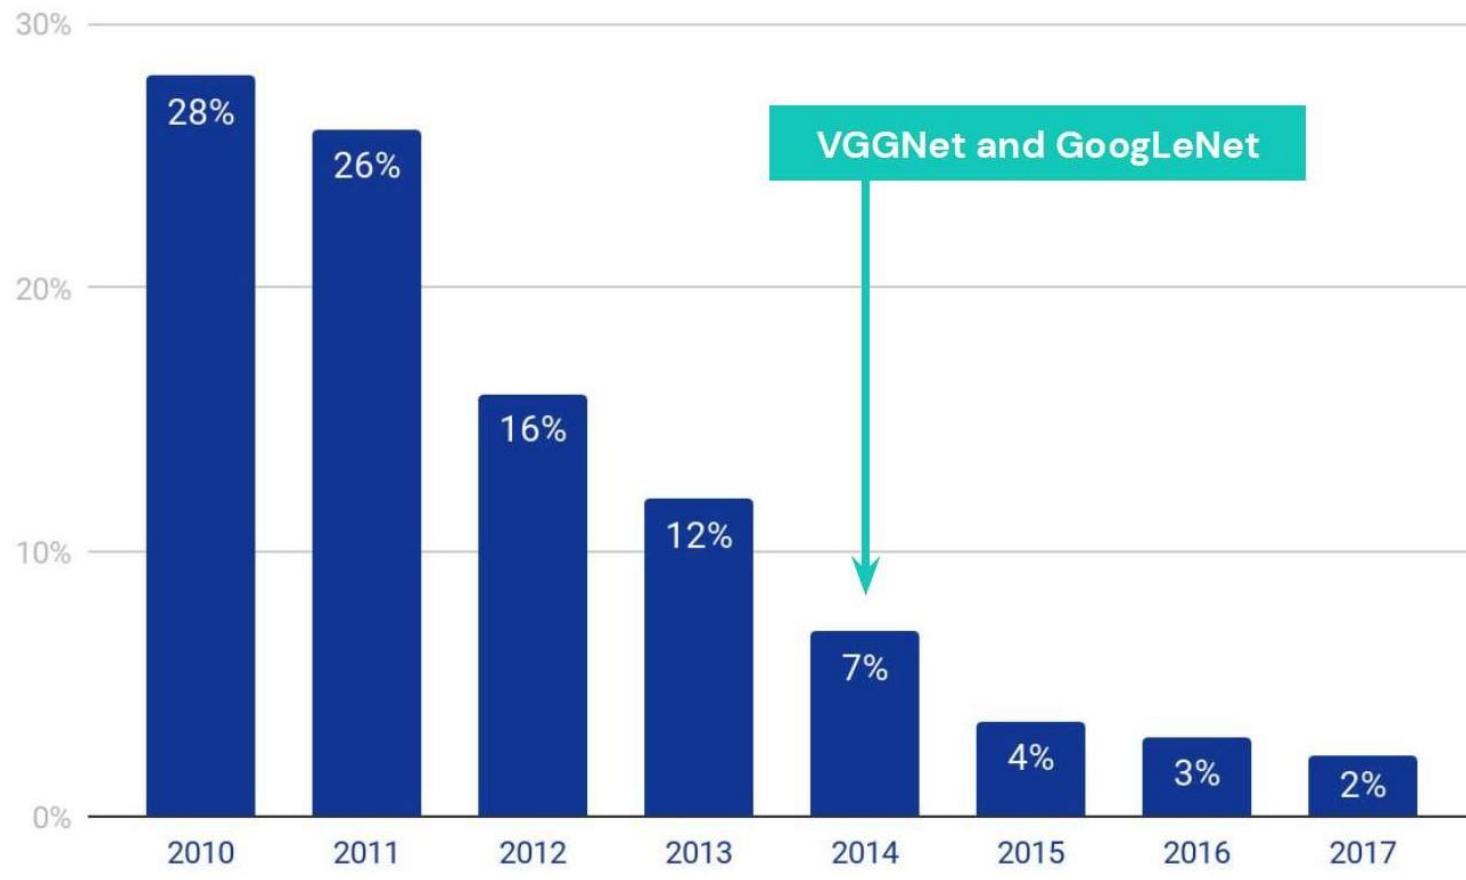
BACKGROUND

Data Drives Research



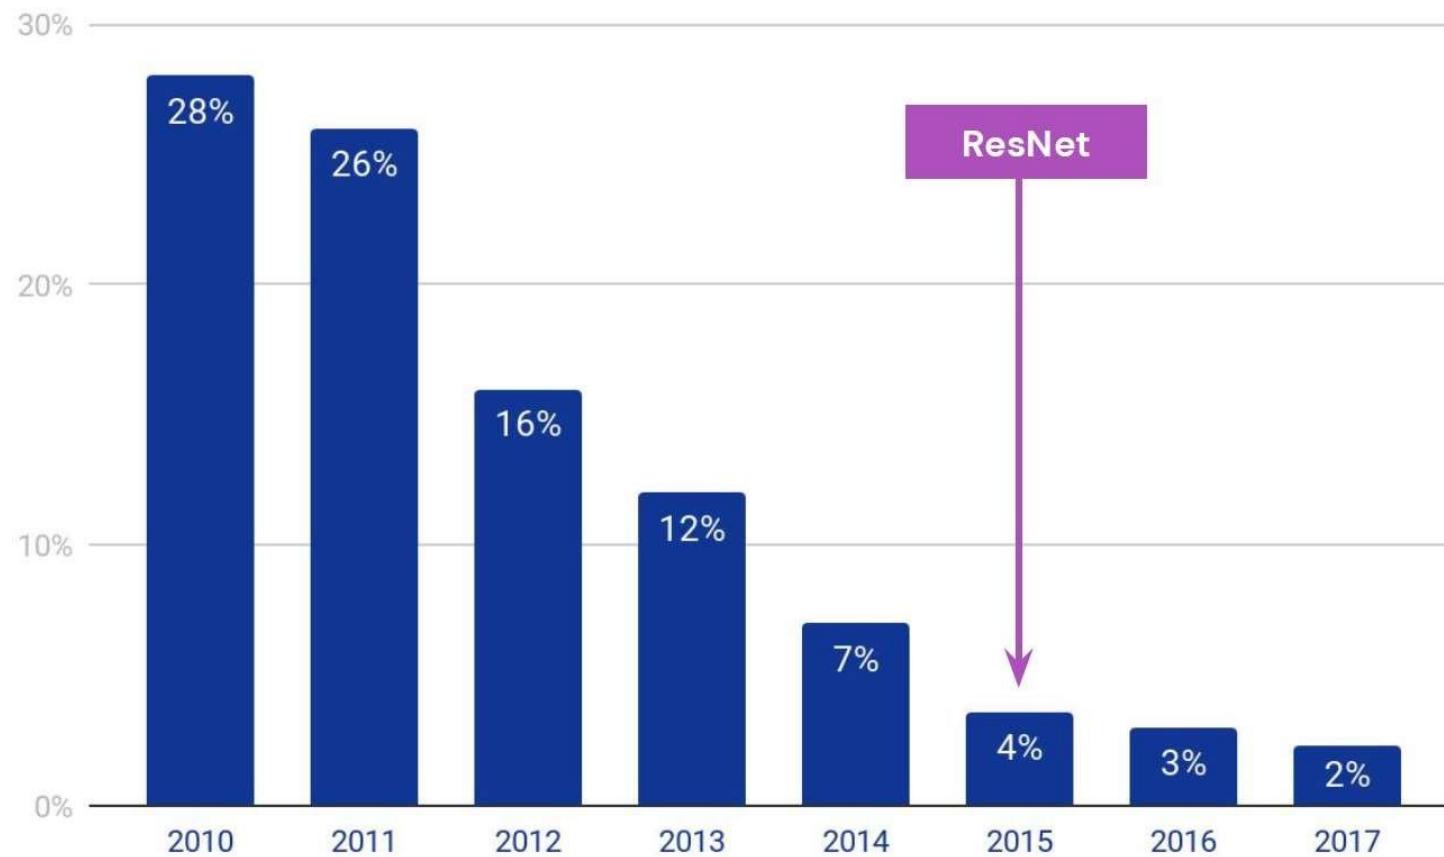
BACKGROUND

Data Drives Research



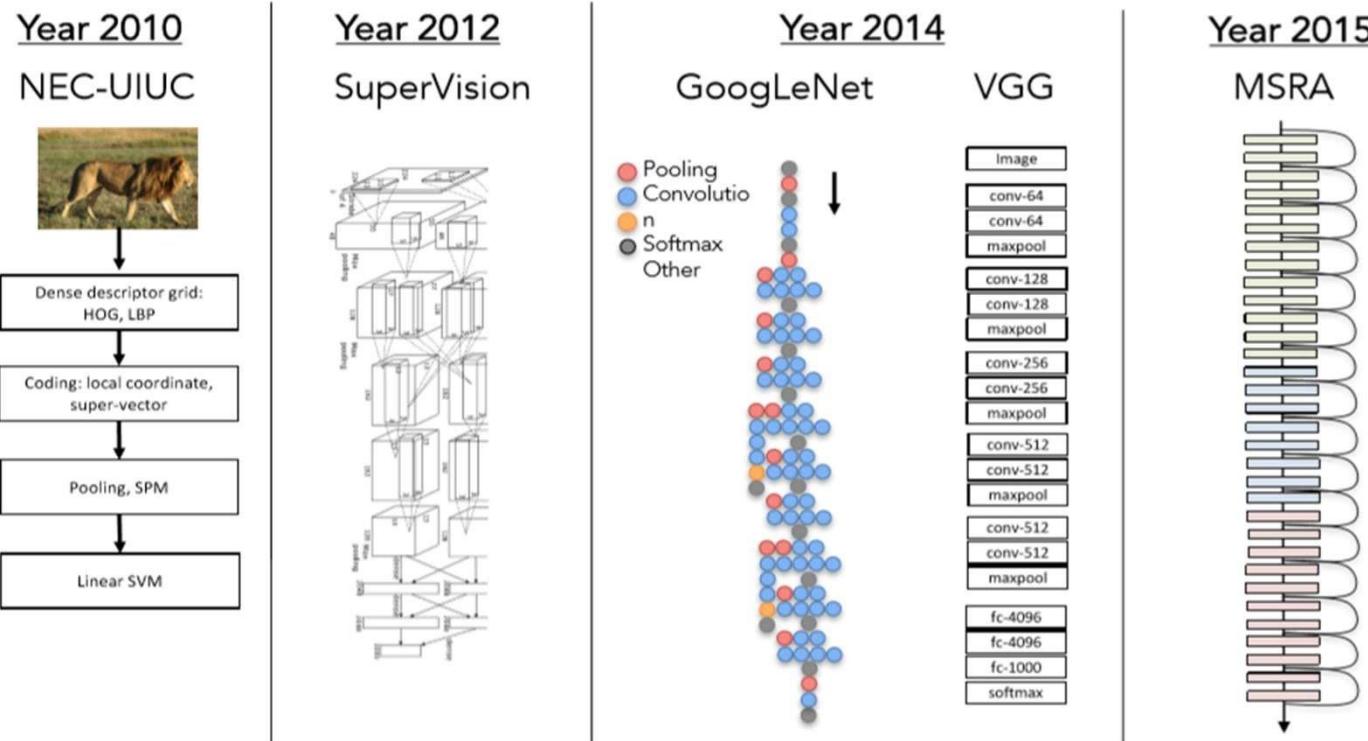
BACKGROUND

Data Drives Research



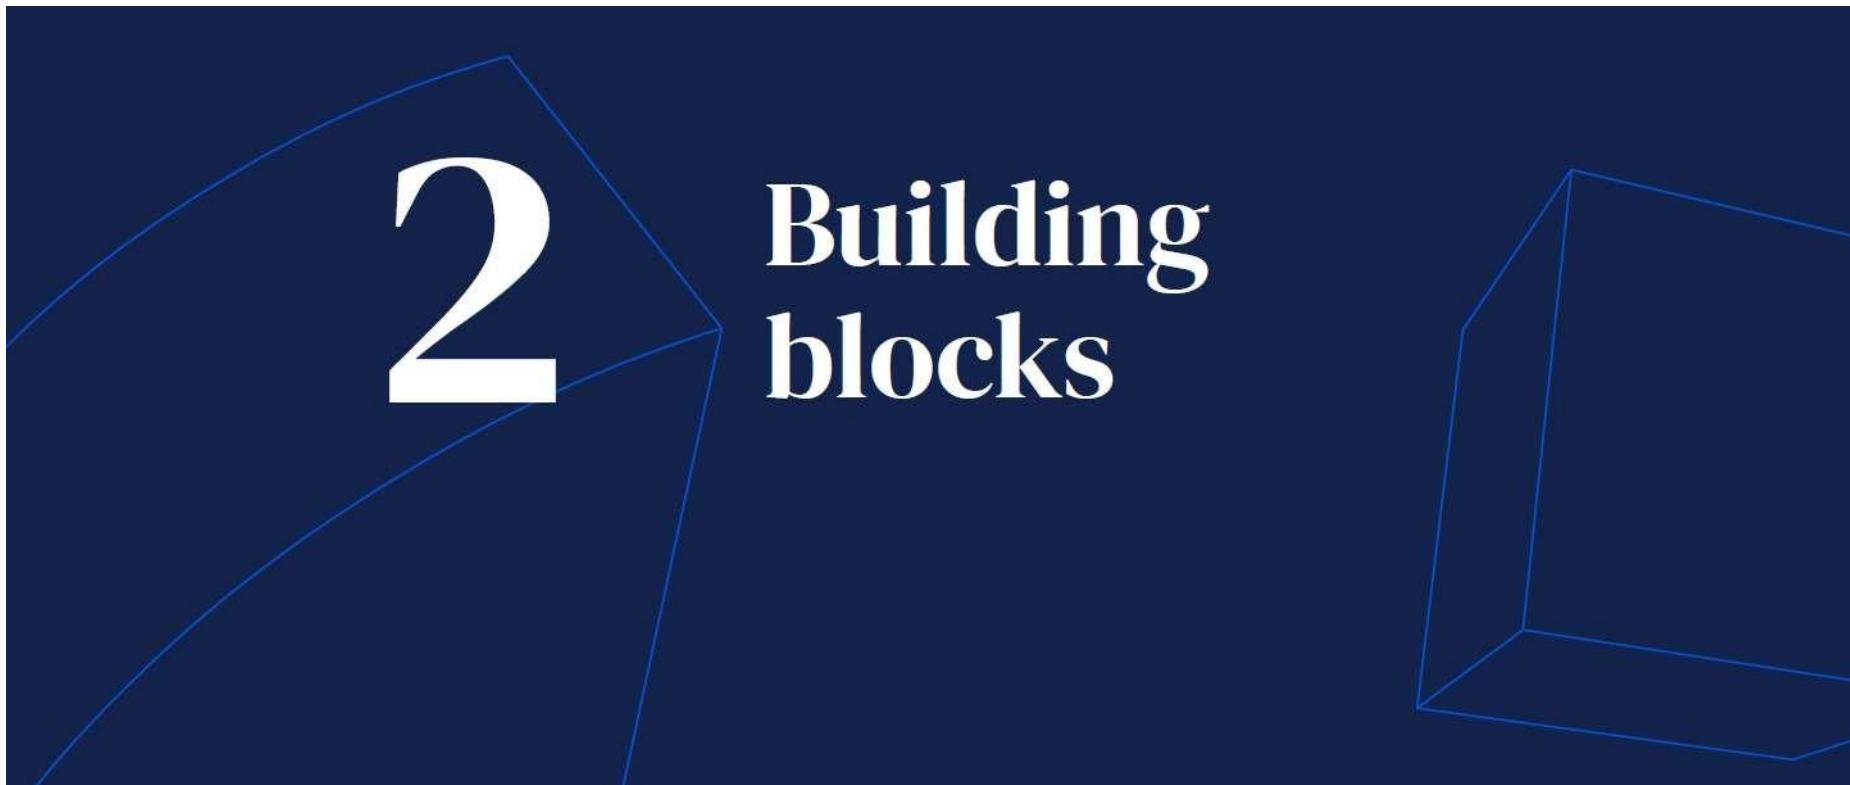
BACKGROUND

Data Drives Research



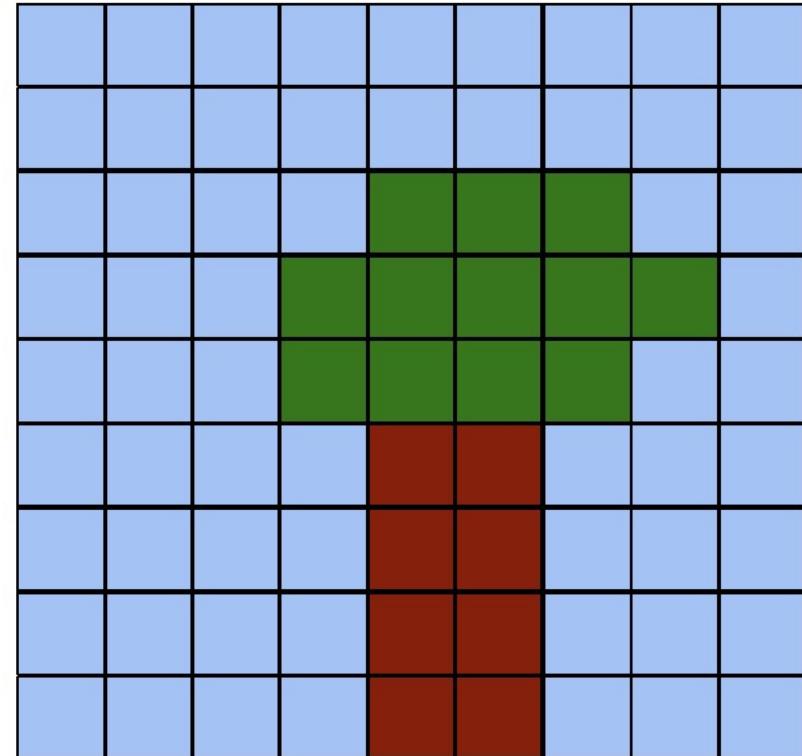
Convolutional Neural Networks

Building Blocks



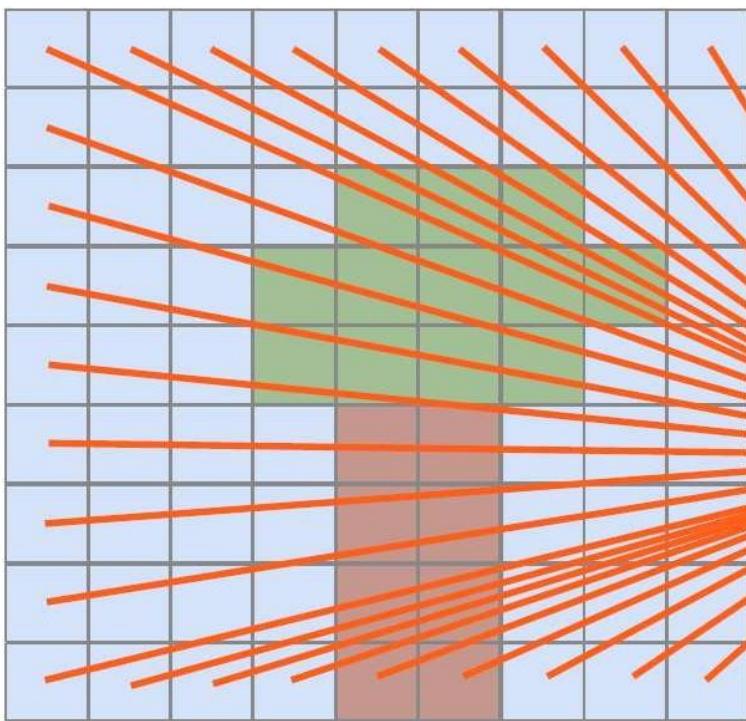
Convolutional Neural Networks

Building Blocks



Convolutional Neural Networks

Building Blocks



From fully connected to locally connected



fully-connected unit

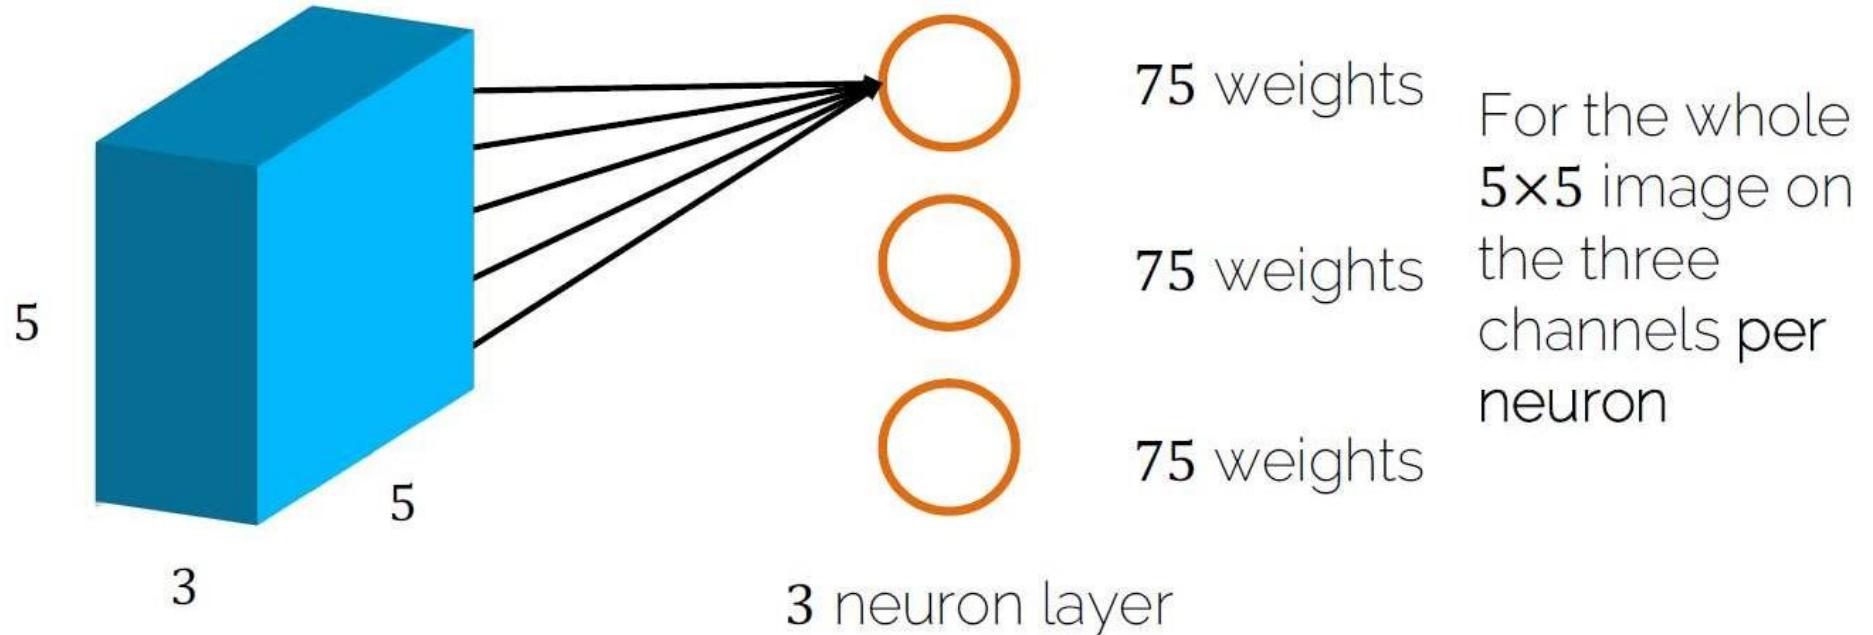
$$y = \sum_{i \in \text{image}} \mathbf{w}_i \mathbf{x}_i + b$$

Convolutional Neural Networks

Building Blocks

Problems using FC Layers on Images

- How to process a tiny image with FC layers

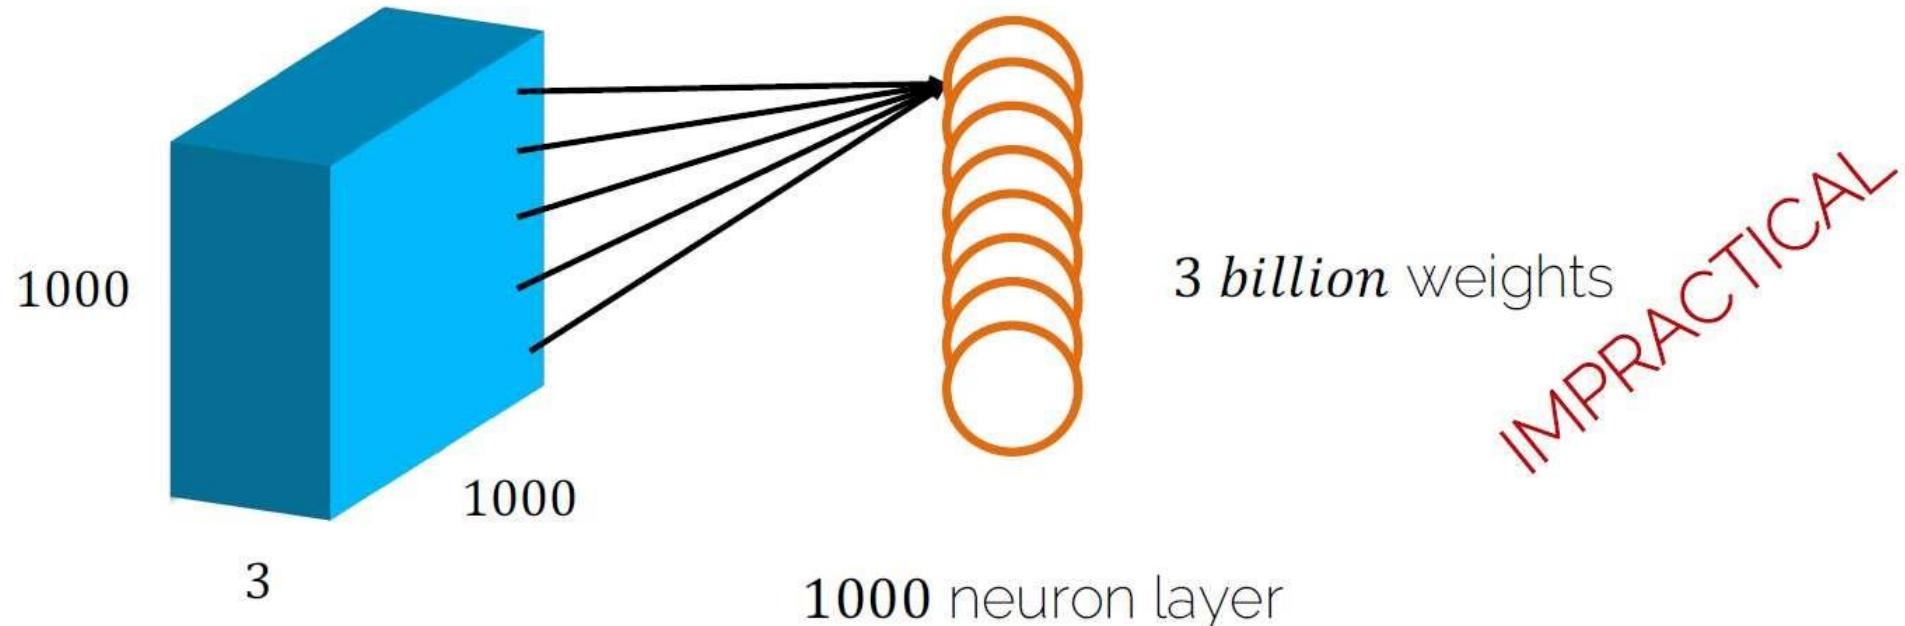


Convolutional Neural Networks

Building Blocks

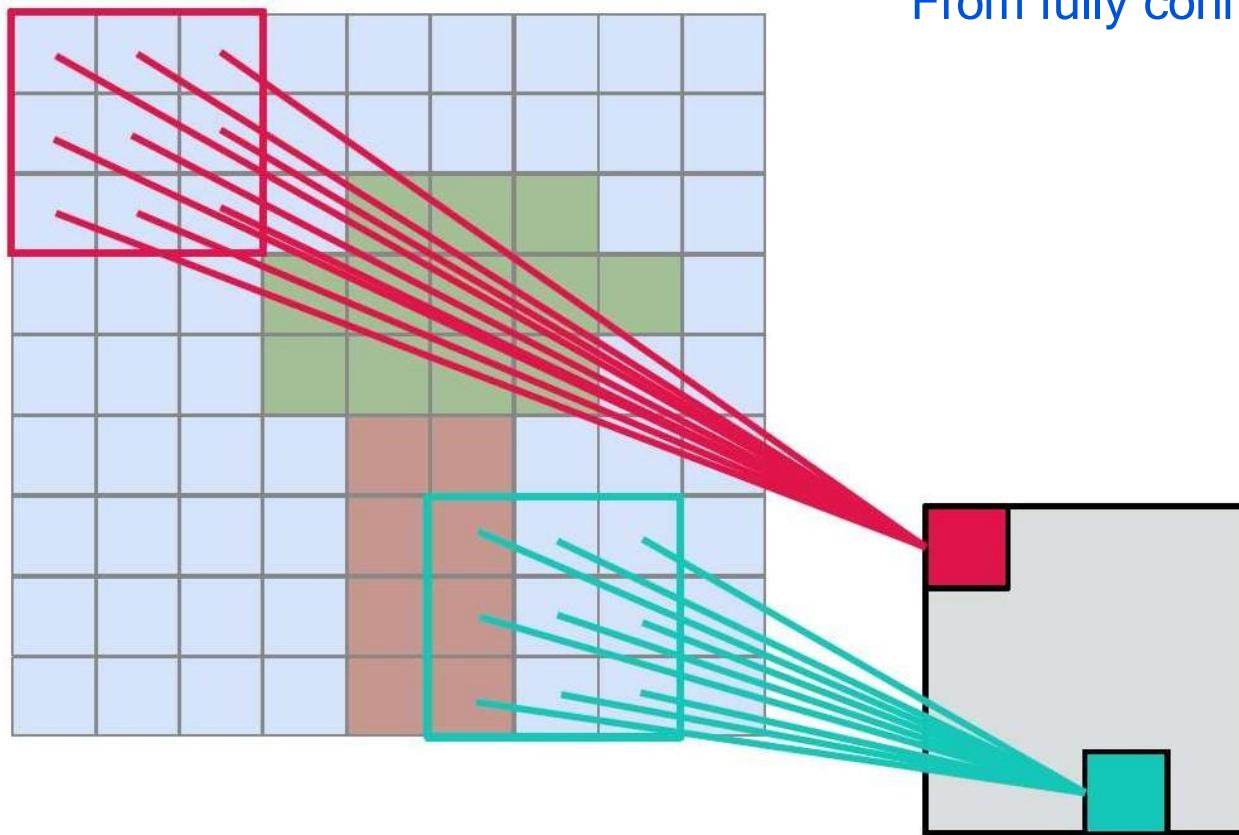
Problems using FC Layers on Images

- How to process a tiny image with FC layers



Convolutional Neural Networks

Building Blocks



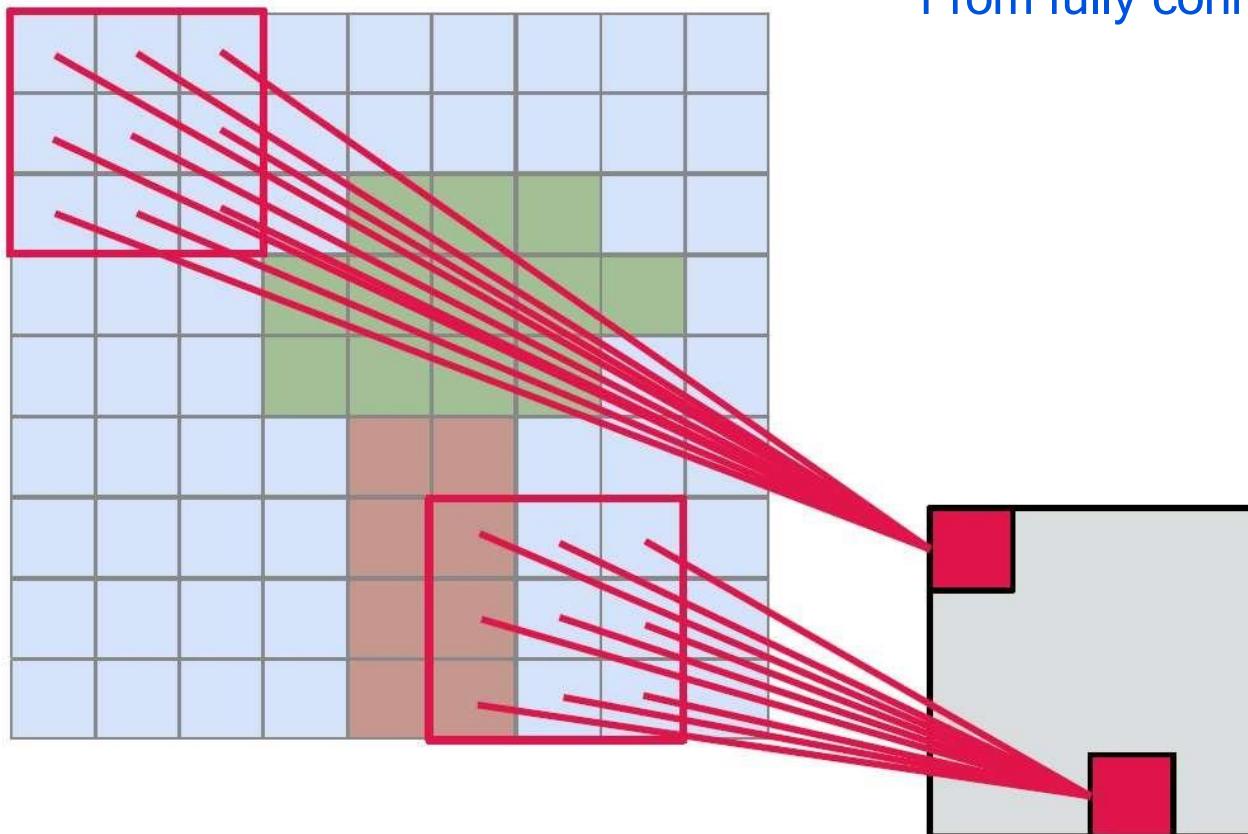
From fully connected to locally connected

$$y = \sum_{i \in 3 \times 3} \mathbf{w}_i \mathbf{x}_i + b$$

locally-connected units
3×3 receptive field

Convolutional Neural Networks

Building Blocks



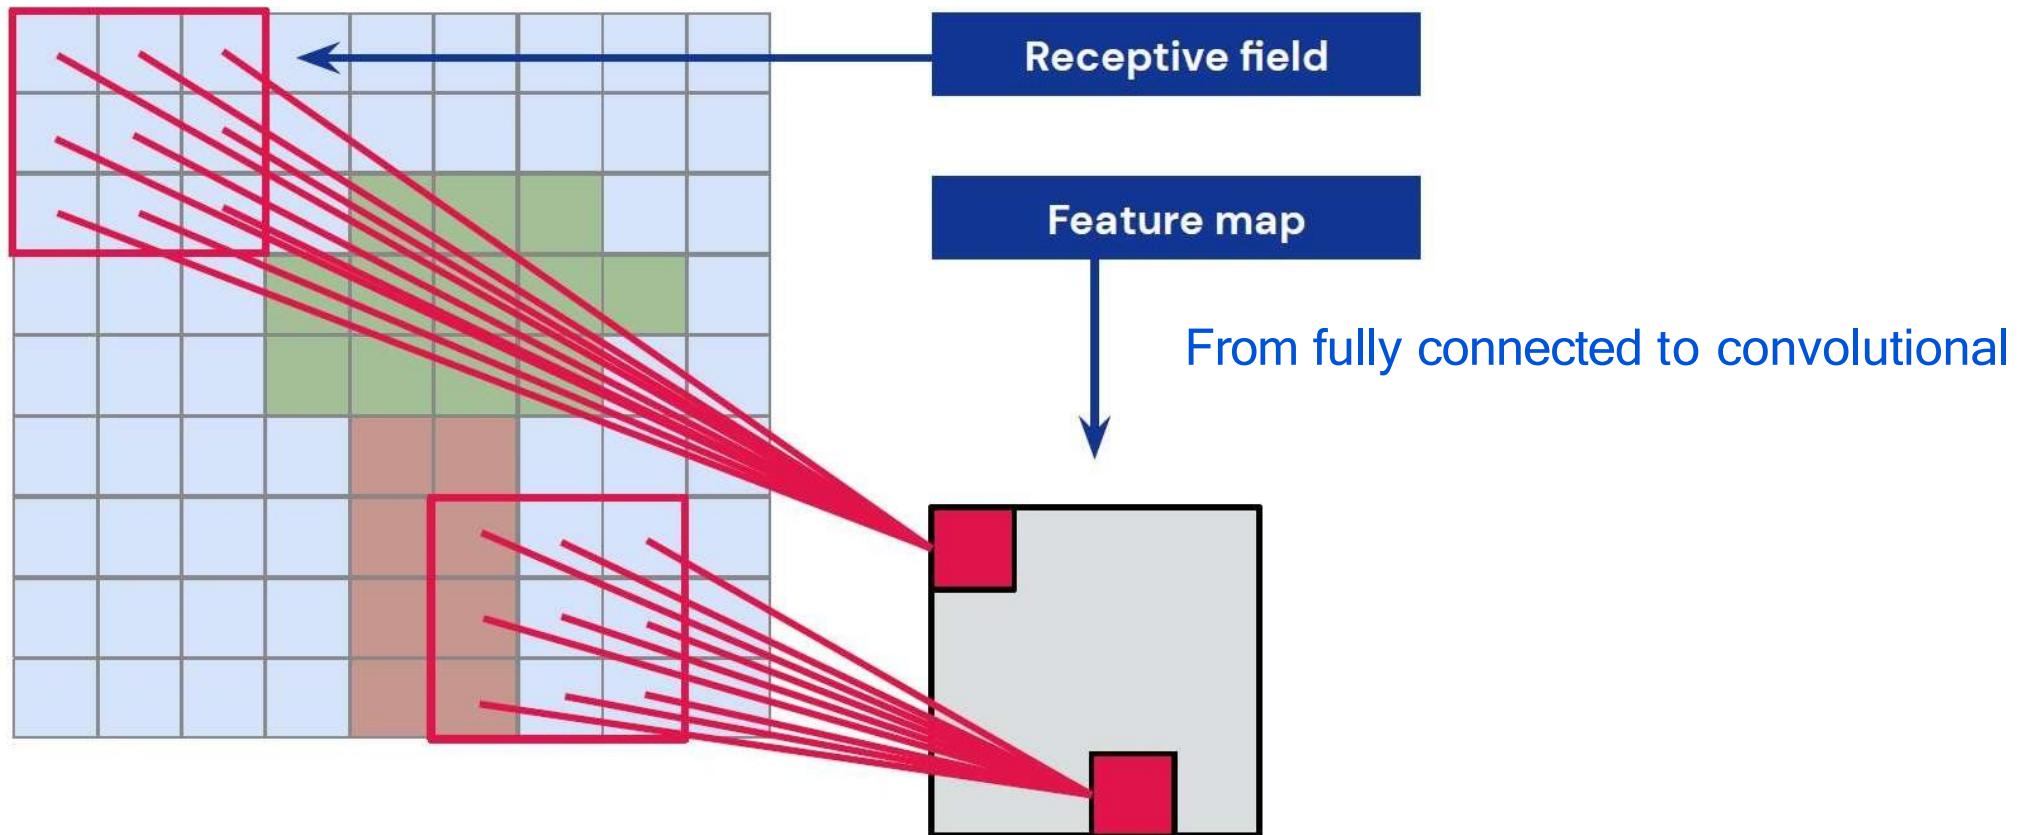
From fully connected to convolutional

$$y = \mathbf{w} * \mathbf{x} + b$$

convolutional units
3×3 receptive field

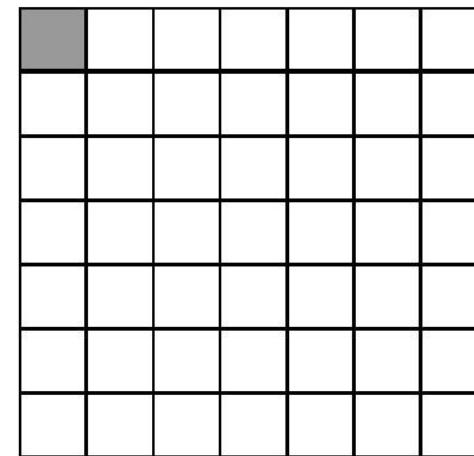
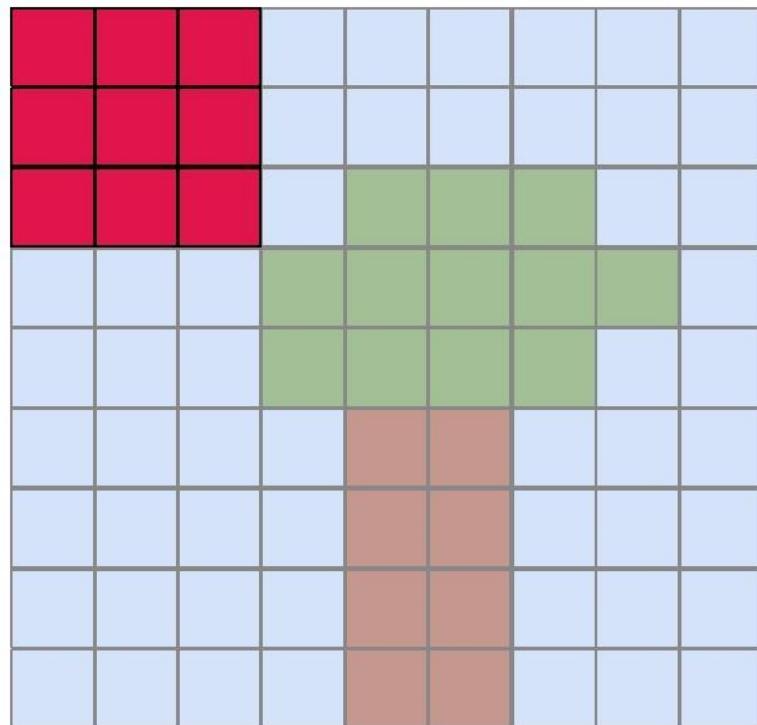
Convolutional Neural Networks

Building Blocks



Convolutional Neural Networks

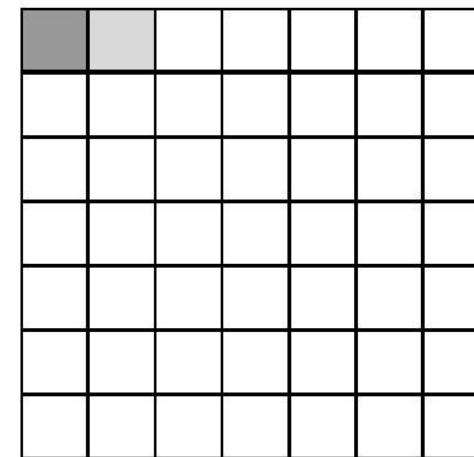
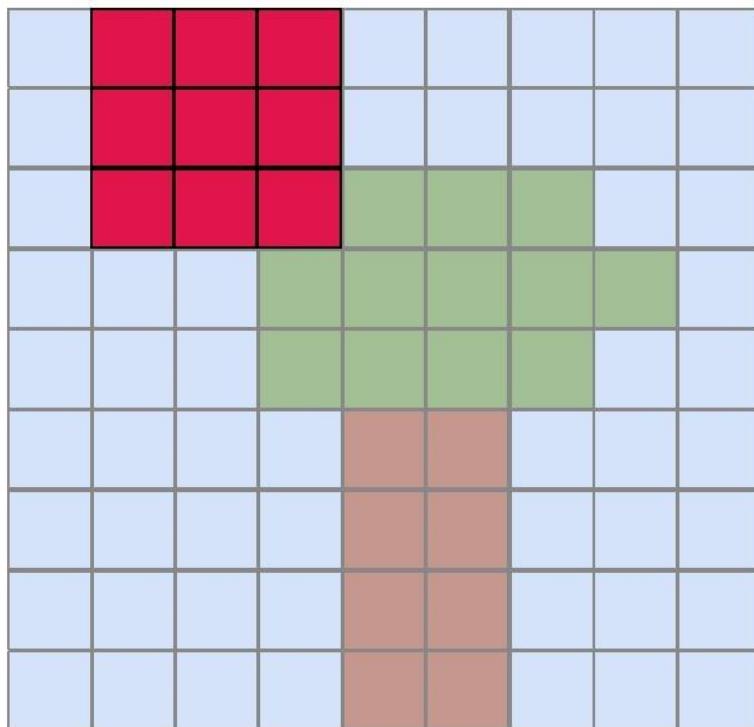
Convolution Operation



The **kernel** slides across the image and produces an output value at each position

Convolutional Neural Networks

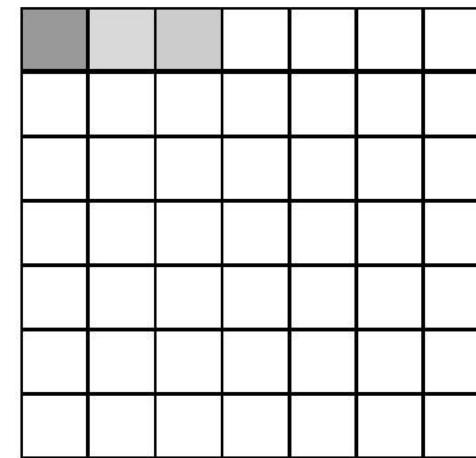
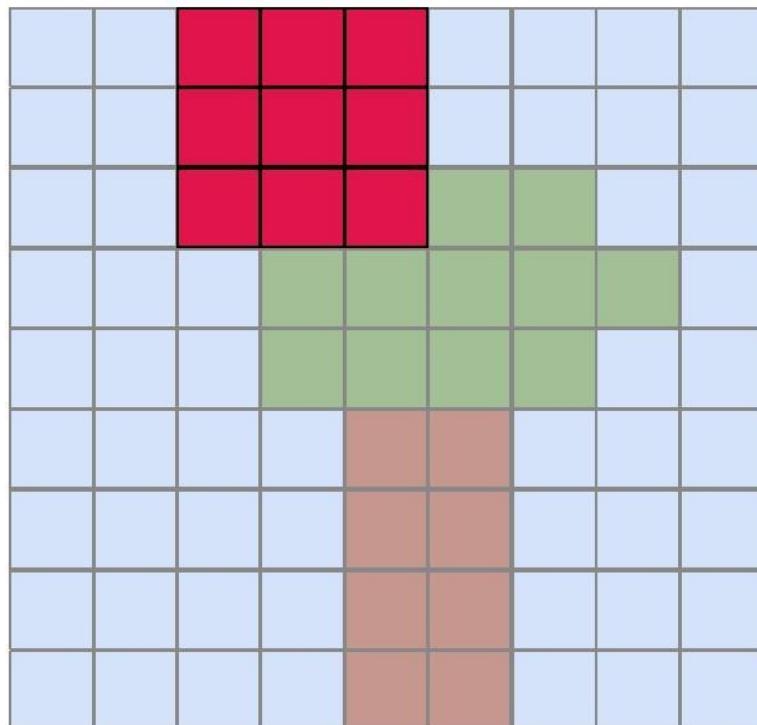
Convolution Operation



The **kernel** slides across the image and produces an output value at each position

Convolutional Neural Networks

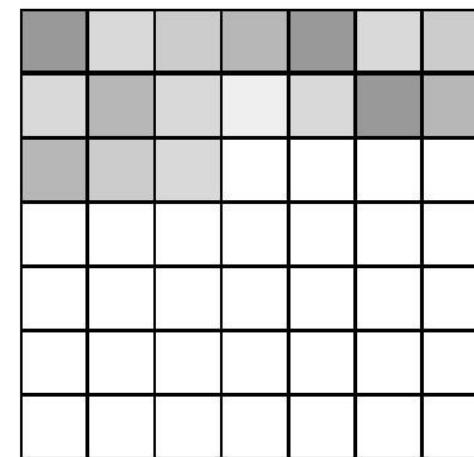
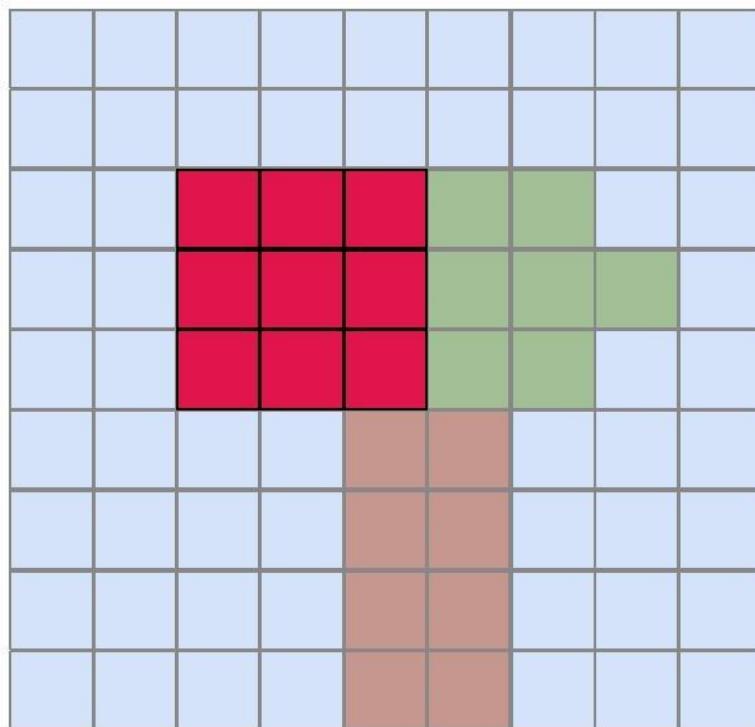
Convolution Operation



The **kernel** slides across the image and produces an output value at each position

Convolutional Neural Networks

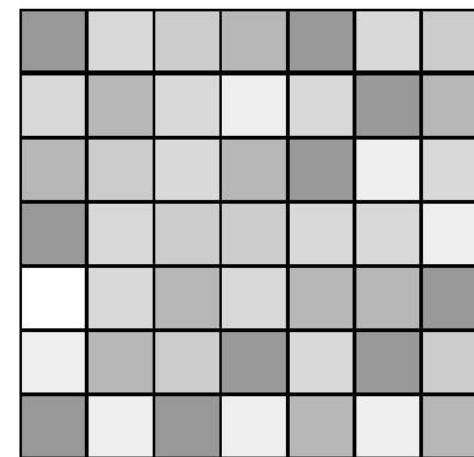
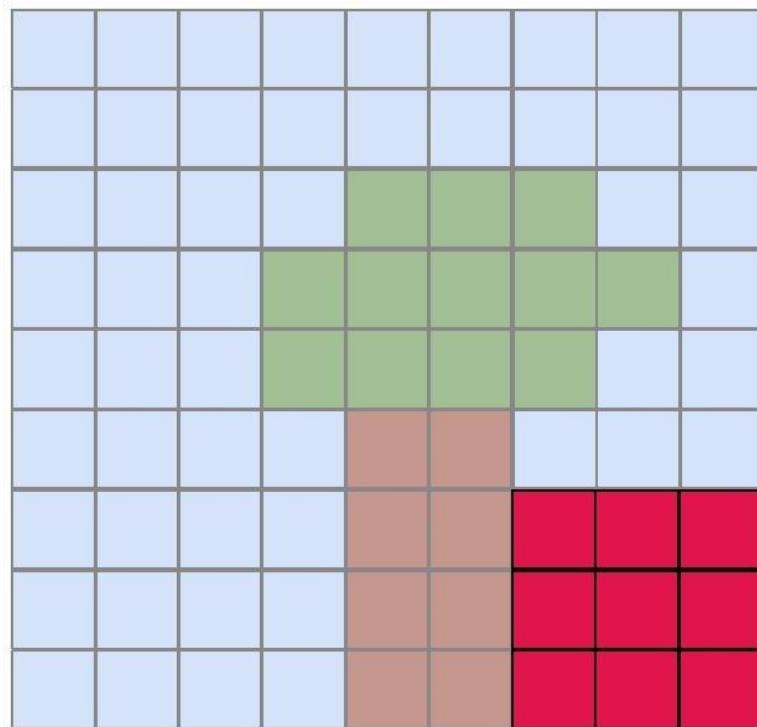
Convolution Operation



The **kernel** slides across the image and produces an output value at each position

Convolutional Neural Networks

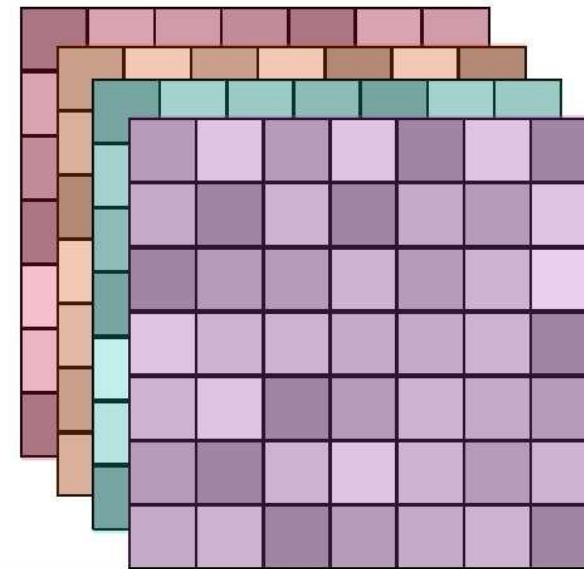
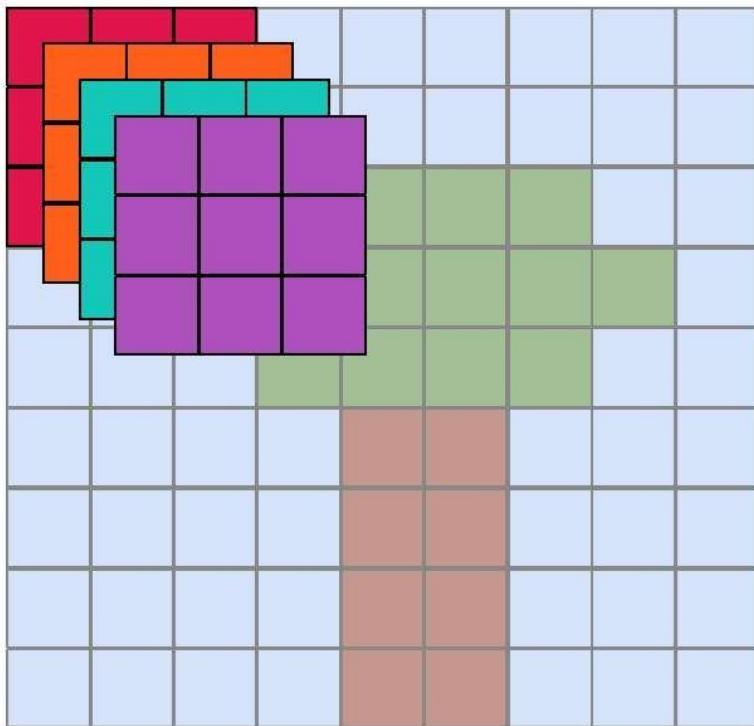
Convolution Operation



The **kernel** slides across the image and produces an output value at each position

Convolutional Neural Networks

Convolution Operation

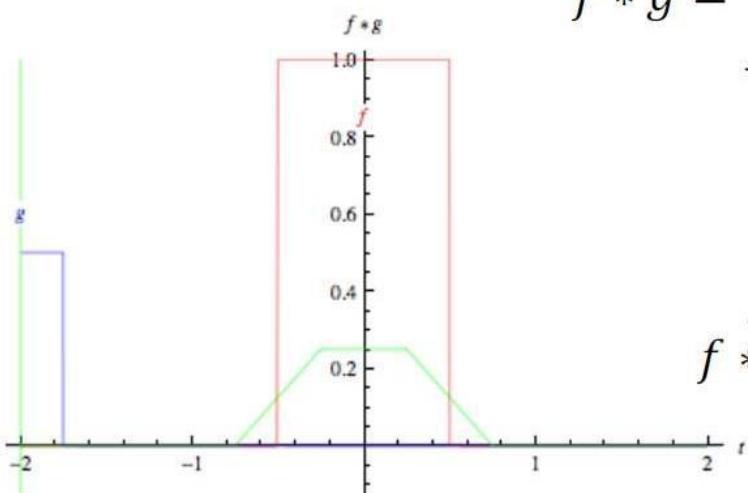


We convolve multiple kernels and obtain multiple feature maps or **channels**

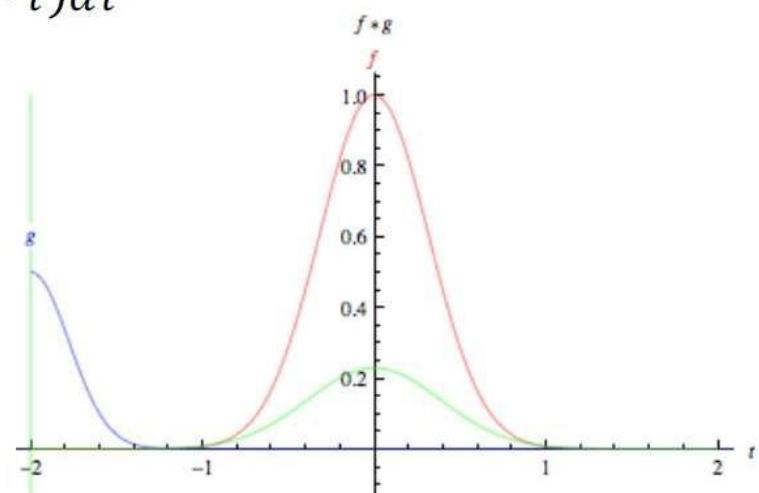
Convolutional Neural Networks

Convolution Operation

$$f * g = \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau$$



Convolution of two box functions



Convolution of two Gaussians

Application of a filter to a function

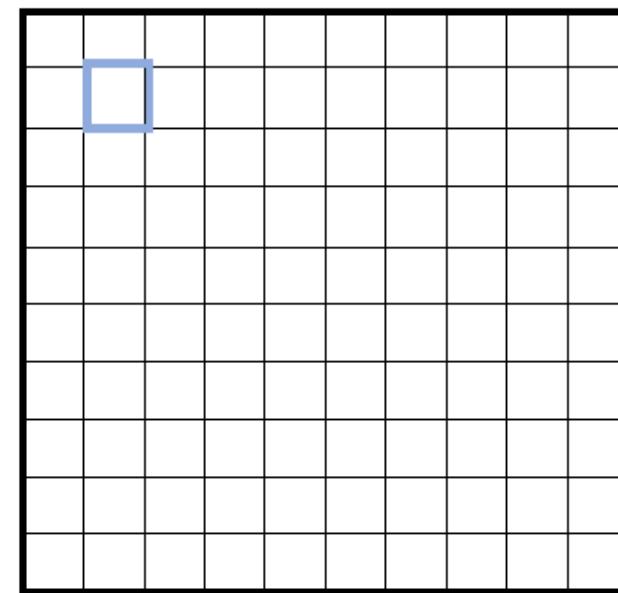
- The ‘smaller’ one is typically called the filter kernel

Convolutional Neural Networks

Convolution Operation

$$h[m, n] = \sum_{k,l} f[k, l] I[m+k, n+l]$$

| | | | | | | | | | | | |
|---|---|----|----|----|----|----|----|----|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |



Convolutional Neural Networks

Convolution Operation

$$h[m, n] = \sum_{k,l} f[k, l] I[m+k, n+l]$$

| | | | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | | | | | |
|---|----|--|--|--|--|--|--|--|--|--|--|
| 0 | 10 | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |

Convolutional Neural Networks

Convolution Operation

$$h[m, n] = \sum_{k,l} f[k, l] I[m+k, n+l]$$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|----|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 0 | 90 | 90 | 90 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 90 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | | | |
|--|--|--|---|----|----|--|--|--|--|
| | | | | | | | | | |
| | | | | | | | | | |
| | | | 0 | 10 | 20 | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Convolutional Neural Networks

Convolution Operation

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | | | |
|--|--|--|---|----|----|----|--|--|--|
| | | | | | | | | | |
| | | | | | | | | | |
| | | | 0 | 10 | 20 | 30 | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Convolutional Neural Networks

Convolution Operation

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Convolutional Neural Networks

Convolution Operation

$$h[m, n] = \sum_{k,l} f[k, l] I[m+k, n+l]$$

| | | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | | | | |
|--|--|--|----|----|----|----|----|----|----|----|
| | | | | | | | | | | |
| | | | 0 | 10 | 20 | 30 | 30 | 30 | 20 | 10 |
| | | | 0 | 20 | 40 | 60 | 60 | 60 | 40 | 20 |
| | | | 0 | 30 | 60 | 90 | 90 | 90 | 60 | 30 |
| | | | 0 | 30 | 50 | 80 | 80 | 90 | 60 | 30 |
| | | | 0 | 30 | 50 | 80 | 80 | 90 | 60 | 30 |
| | | | 0 | 20 | 30 | 50 | 50 | 60 | 40 | 20 |
| | | | 10 | 20 | 30 | 30 | 30 | 30 | 20 | 10 |
| | | | 10 | 10 | 10 | 0 | 0 | 0 | 0 | 0 |
| | | | | | | | | | | |

Convolutional Neural Networks

Convolution Operation

Image filtering

The “kernels” compute a function of the surrounding pixels

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Very important in image-processing applications

Improve or enhance properties of the image

Noise, sizes changes, contrast, etc

Information extraction

Texture, lines, keypoints, etc. Pattern detection and matching:

Find correspondences between templates

“Gradient Filter”

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -1 | 0 | 1 |
| -1 | 0 | 1 |

“Sobel Filter”

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

Convolutional Neural Networks

Convolution Operation

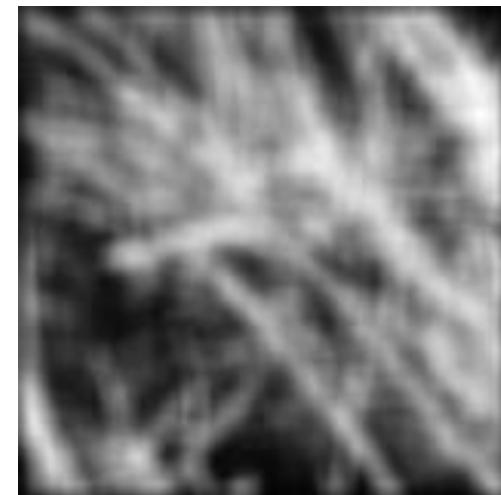
Image filtering : Smoothing

- This “kernel” replaces each pixel with the means of its neighbors
- It yields a smoothing effect (low-pass filter)
- A type of image processing highly used in computer vision



$$f[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

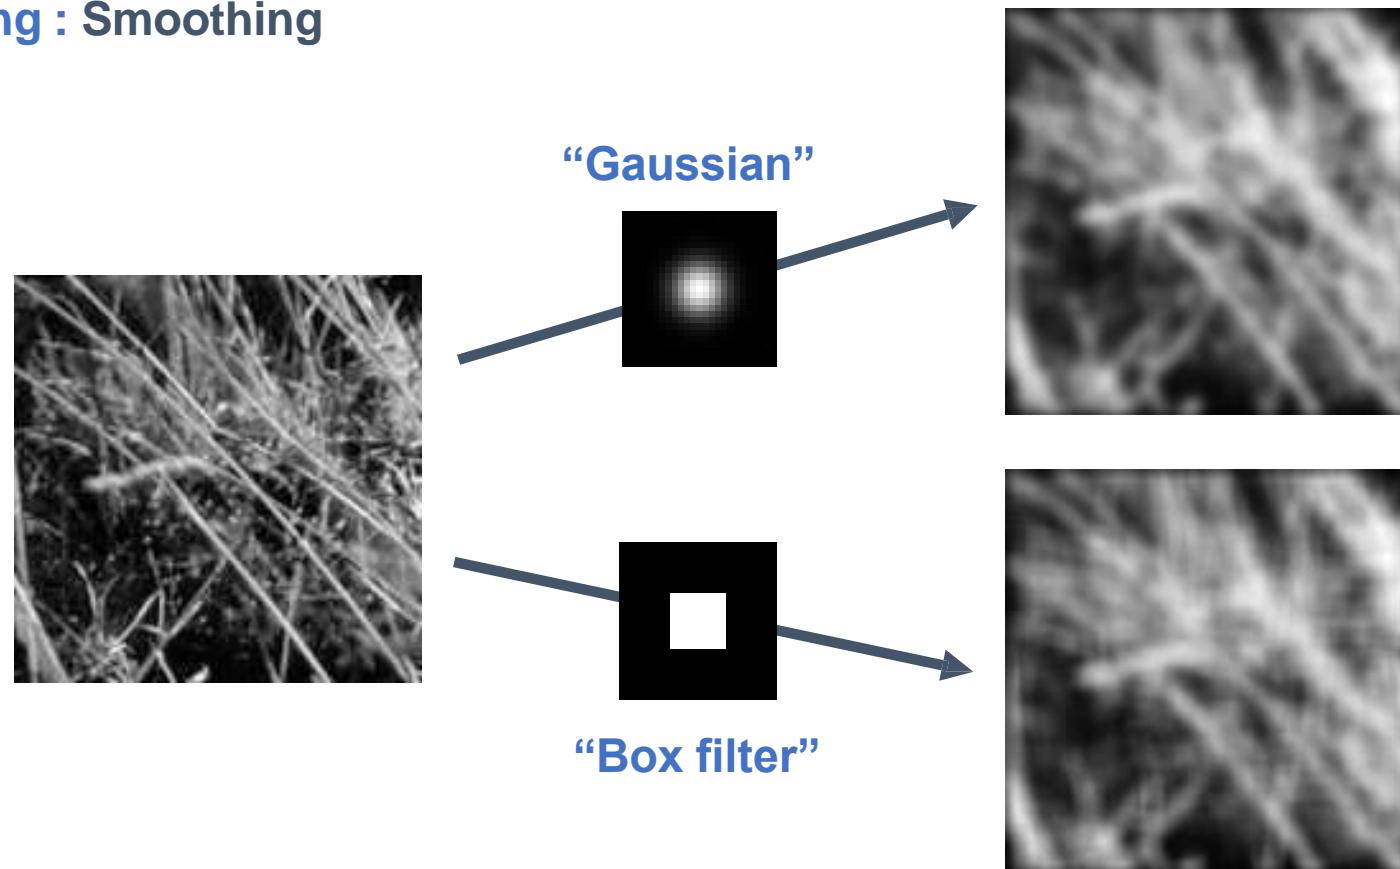
“Box filter”



Convolutional Neural Networks

Convolution Operation

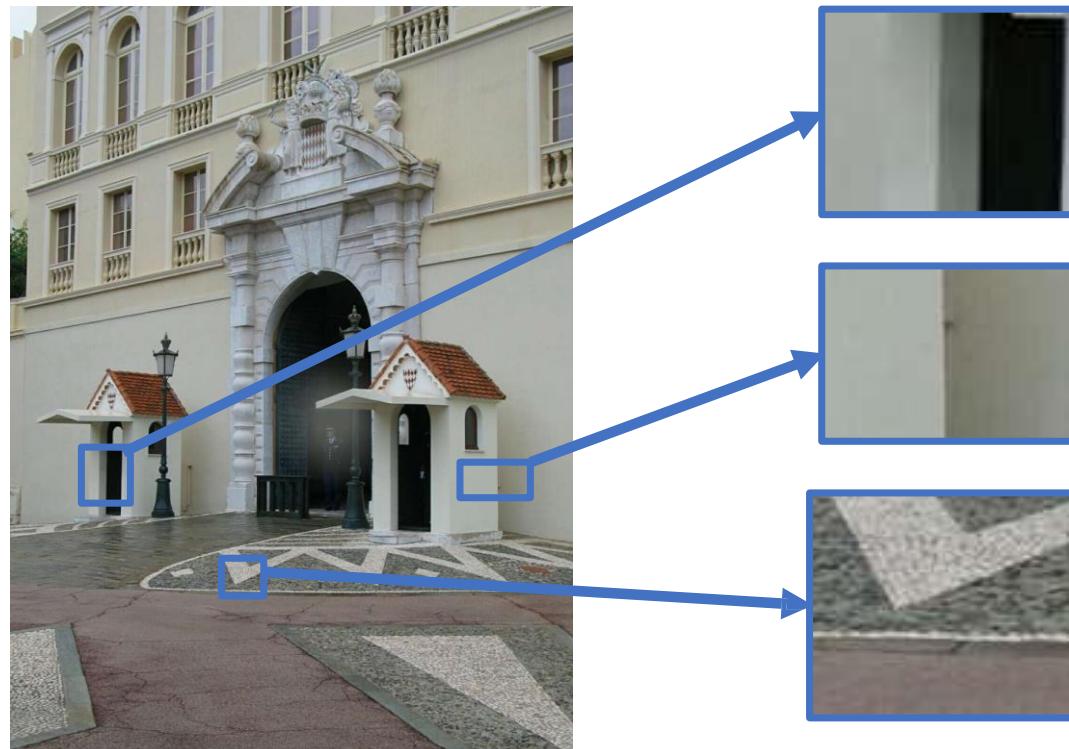
Image filtering : Smoothing



Convolutional Neural Networks

Convolution Operation

Image filtering : Edge Detection



Discontinuity in
normal of the
the surface

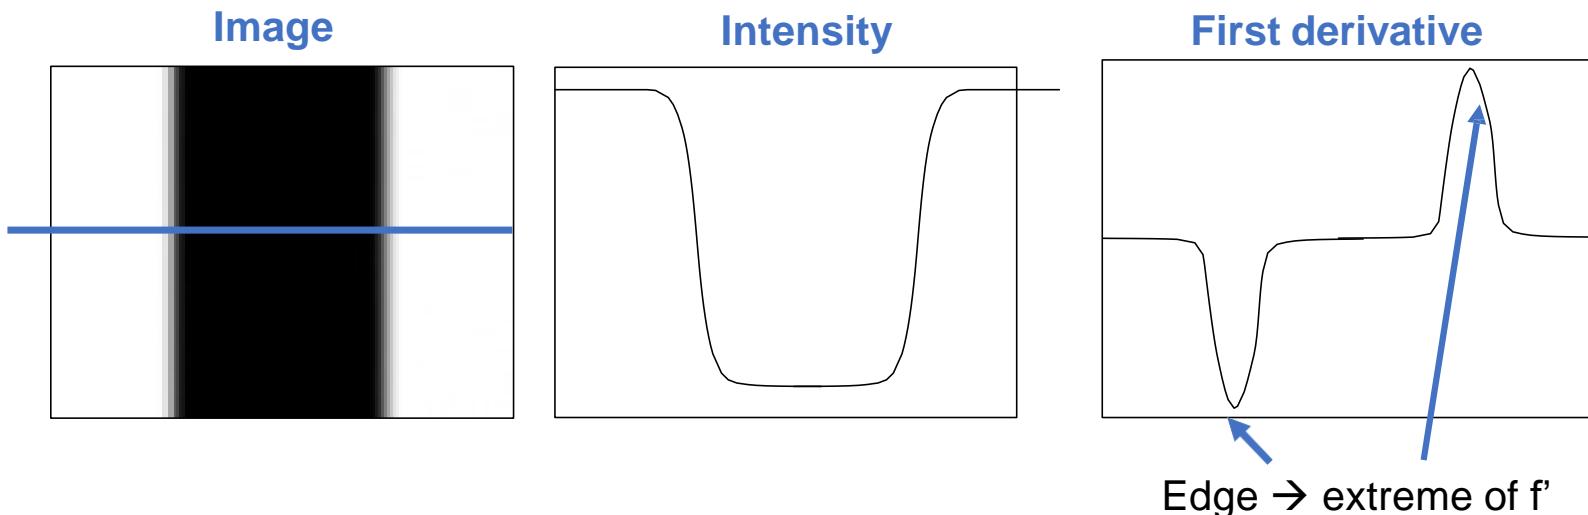
Depth
Discontinuity

Color or texture
discontinuity

Convolutional Neural Networks

Convolution Operation

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x) = f_x$$



Convolutional Neural Networks

Convolution Operation inside

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x)$$

$$\frac{df}{dx} = \frac{f(x) - f(x - 1)}{1} = f'(x)$$

$$\frac{df}{dx} = f(x) - f(x - 1) = f'(x)$$

Discrete derivative in 1D

Convolutional Neural Networks

Convolution Operation

Backward

$$\frac{df}{dx} = f(x) - f(x-1) = f'(x)$$

Forward

$$\frac{df}{dx} = f(x) - f(x+1) = f'(x)$$

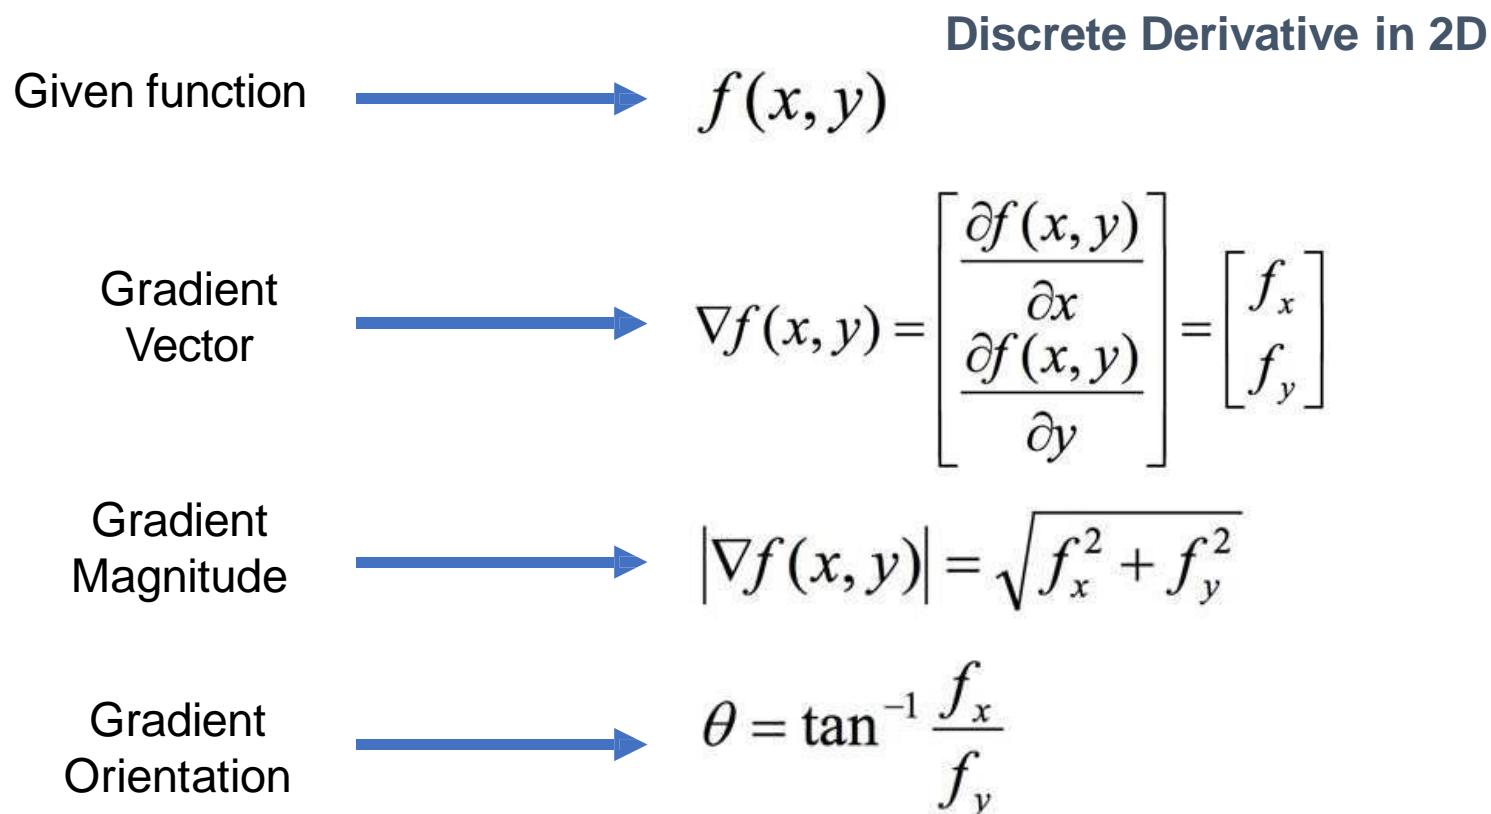
Central

$$\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$$

Discrete derivative in 1D

Convolutional Neural Networks

Convolution Operation



Convolutional Neural Networks

Convolution Operation

Backward filter: $f(x) - f(x-1) = f'(x)$

$$[\begin{matrix} 0 & 1 & -1 \end{matrix}]$$

$f(x) - f(x+1) = f'(x)$

Forward:

$$[\begin{matrix} -1 & 1 & 0 \\ 1 & \end{matrix}]$$

Central: $f(x+1) - f(x-1) = f'(x)$

$$[\begin{matrix} 0 & -1 \\ 1 & 1 \end{matrix}]$$

Types of discrete derivatives in 1D

Convolutional Neural Networks

Convolution Operation

| | | |
|----|---|--|
| -1 | 0 | |
| 0 | 1 | |
| | | |

| | | |
|----|---|----|
| +1 | 0 | -1 |
| +1 | 0 | -1 |
| +1 | 0 | -1 |

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

| | | |
|---|----|--|
| 0 | -1 | |
| 1 | 0 | |
| | | |

| | | |
|----|----|----|
| +1 | +1 | +1 |
| 0 | 0 | 0 |
| -1 | -1 | -1 |

| | | |
|----|---|---|
| - | - | - |
| -1 | 2 | 1 |
| 0 | 0 | 0 |
| 1 | 2 | 1 |

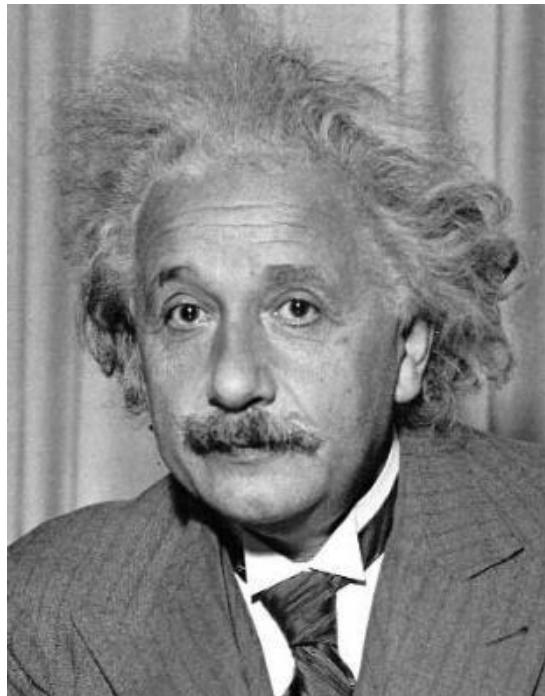
Robert's cross-gradient operators

Prewitt Gradient Operators

Sobel Gradient Operators

Convolutional Neural Networks

Convolution Operation



Sobel

A horizontal blue arrow pointing from left to right, positioned above the kernel matrix.

| | | |
|---|---|---|
| 1 | 2 | 1 |
| 0 | 0 | 0 |
| - | - | - |
| 1 | 2 | 1 |



Convolutional Neural Networks

Convolution Operation

Detección de líneas

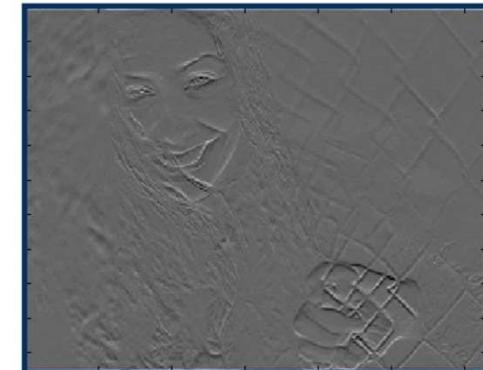
Discrete
Derivative
in 2D



$$\frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

Horizontal Lines

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



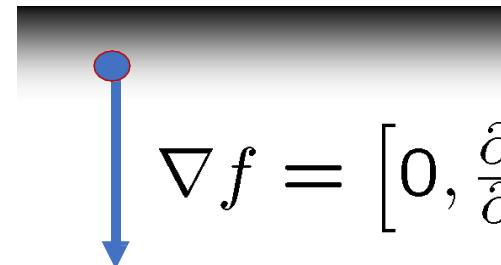
Convolutional Neural Networks

Convolution Operation

The **gradient** of an image $\rightarrow \nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$



$$\rightarrow \nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

The **vector of the gradient** points in the direction of the major increment of the intensity of the input function

The “**strength of the line**” is given by the **magnitude of the gradient**

Convolutional Neural Networks

Convolution Operation

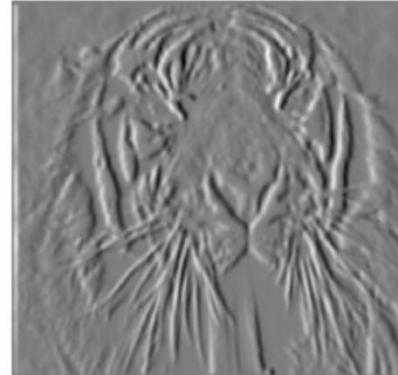
Original
Image



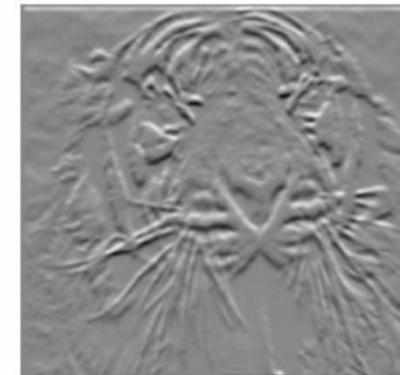
Gradient
Magnitude



Gradient
Direction x

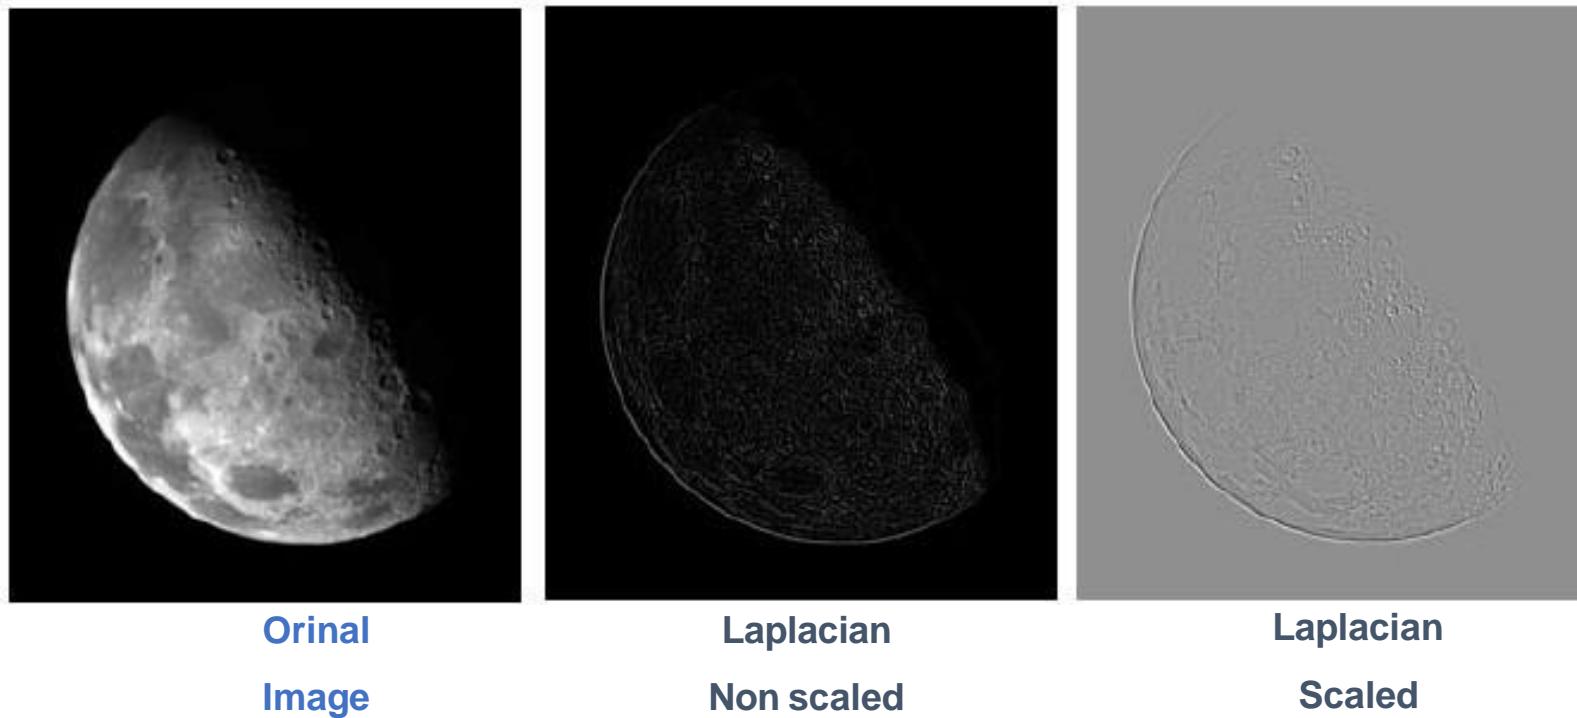


Gradient
Direction y



Convolutional Neural Networks

Convolution Operation



Convolutional Neural Networks

Convolution Operation



Convolutional Neural Networks

Convolution Operation

The demo interface shows two versions of a bell image side-by-side. Below the images are controls for a convolution operation:

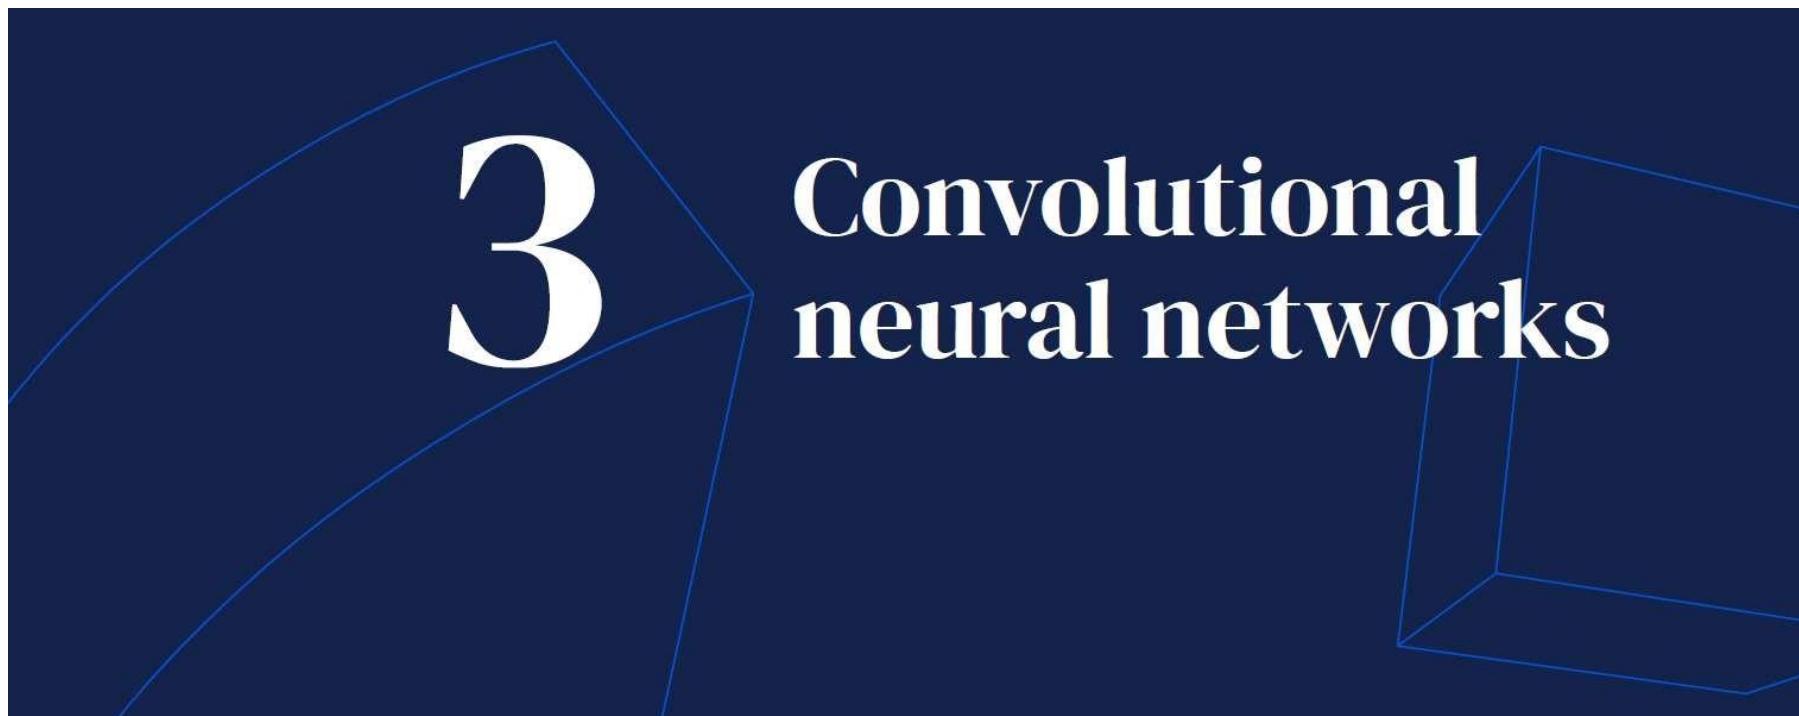
- Filter:** Sharpen
- Image:** Bell
- The Matrix:** A 3x3 grid of numbers:

| | | |
|----|----|----|
| 0 | -3 | 0 |
| -3 | 21 | -3 |
| 0 | -3 | 0 |
- Divisor:** 9

Demo: <http://beej.us/blog/data/convolution-image-processing/>

Convolutional Neural Networks

Convolution Operation



Convolutional Neural Networks

Convolution Operation

Stacking the building blocks

CNNs or ConvNets

Up to 100s of layers

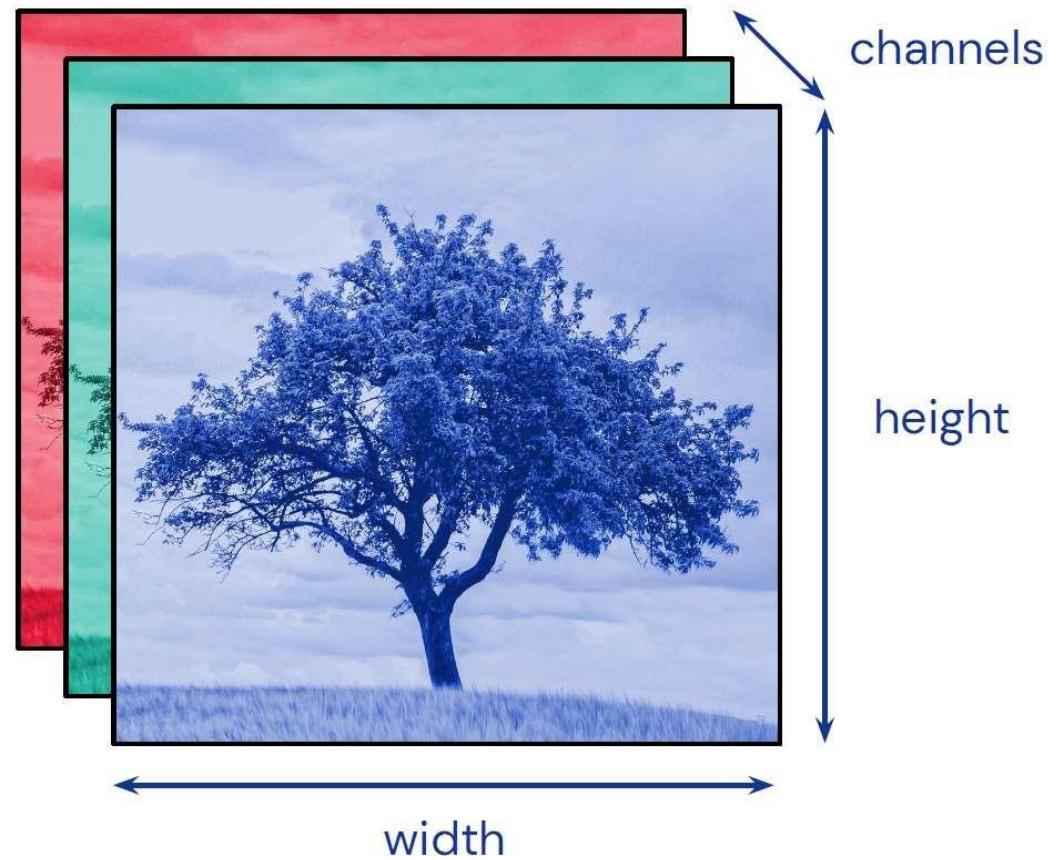
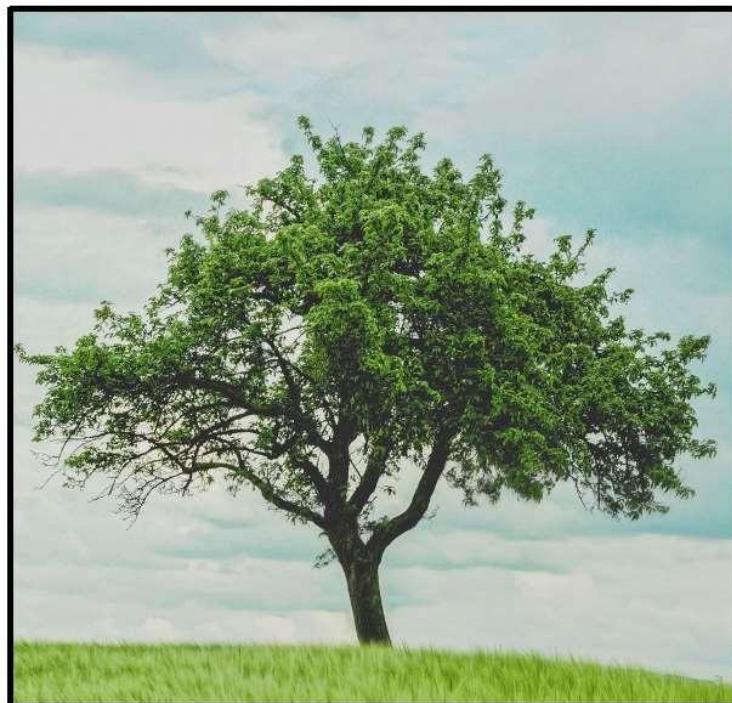
Alternate convolutions and pooling
to create a hierarchy



Convolutional Neural Networks

Convolution Operation

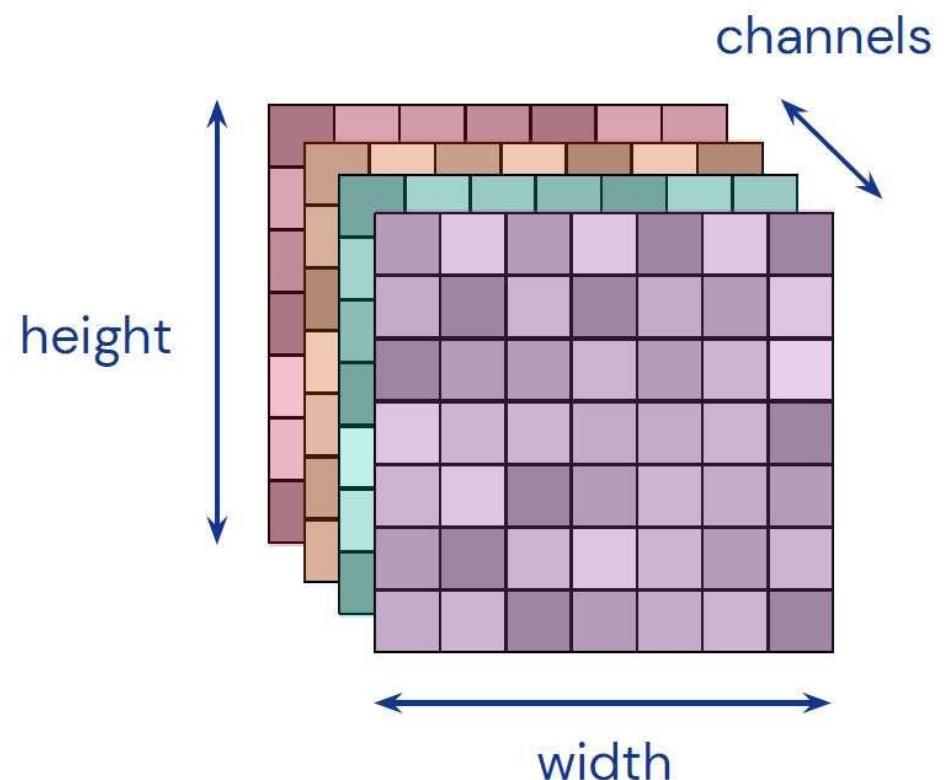
Inputs and outputs are tensors



Convolutional Neural Networks

Convolution Operation

Inputs and outputs are tensors



Convolutional Neural Networks

Convolution Operation



Phil Noble / AP



Convolutional Neural Networks

Convolution Operation

An image contains a discrete number of pixels

Example:

Value of the pixels

“grayscale”

(or “intensity”):

[0,255]

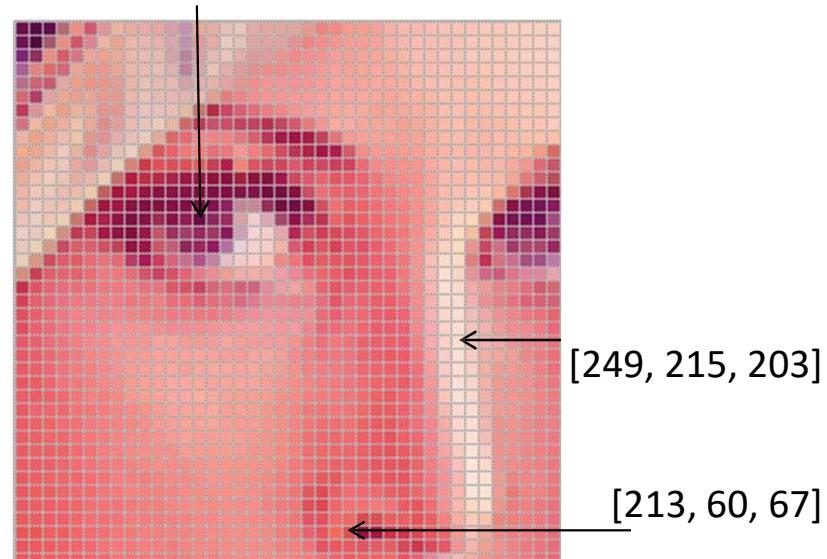
“color”

RGB: [R, G, B]

Lab: [L, a, b]

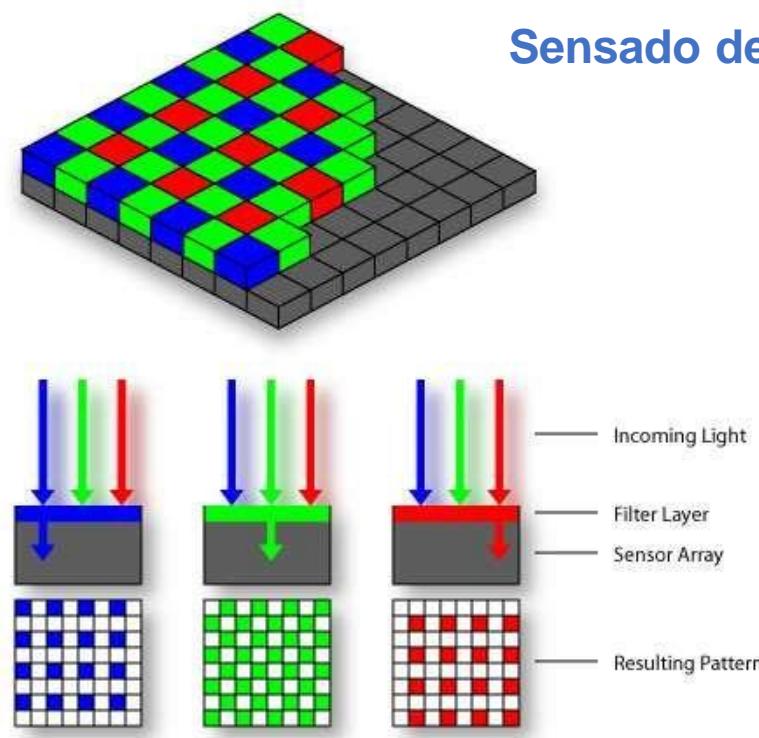
HSV: [H, S, V]

Images are
Sampled and Quantized
[90, 0, 53]



Convolutional Neural Networks

Convolution Operation



Sensado de colores practico: arreglo de Bayer

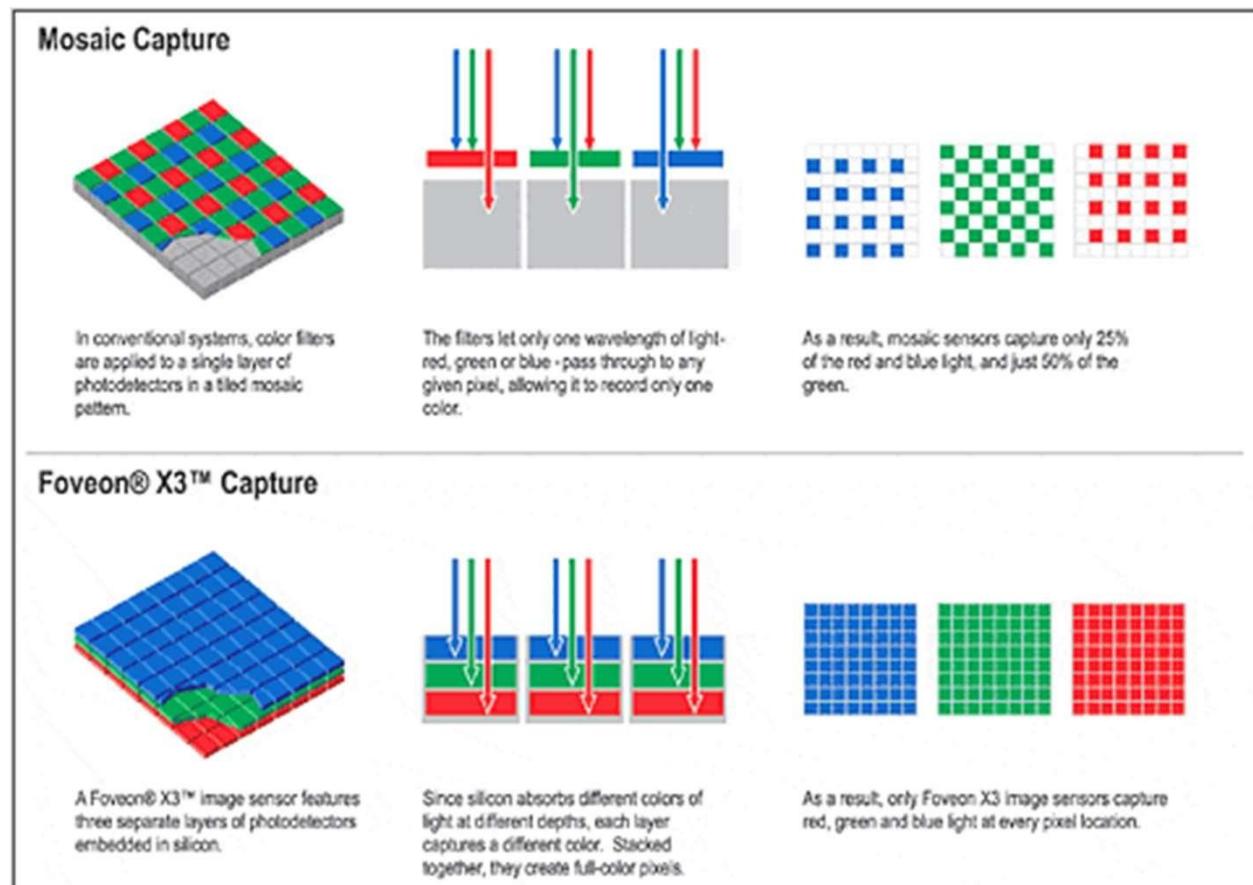
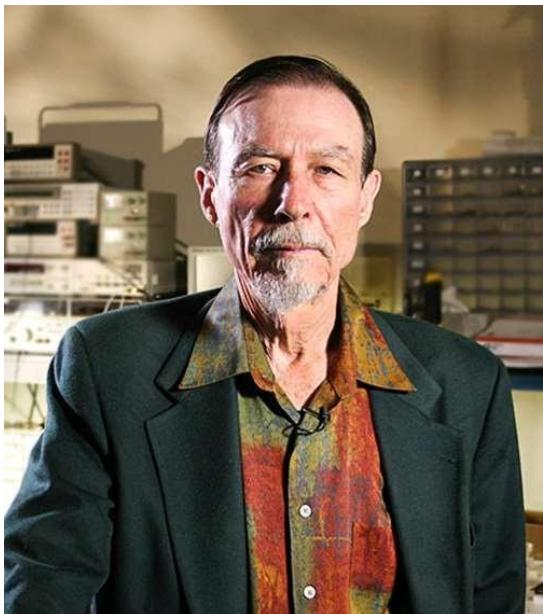


Estimar valores R, G, B
en celda G a partir de
pixeles aledaños

Convolutional Neural Networks

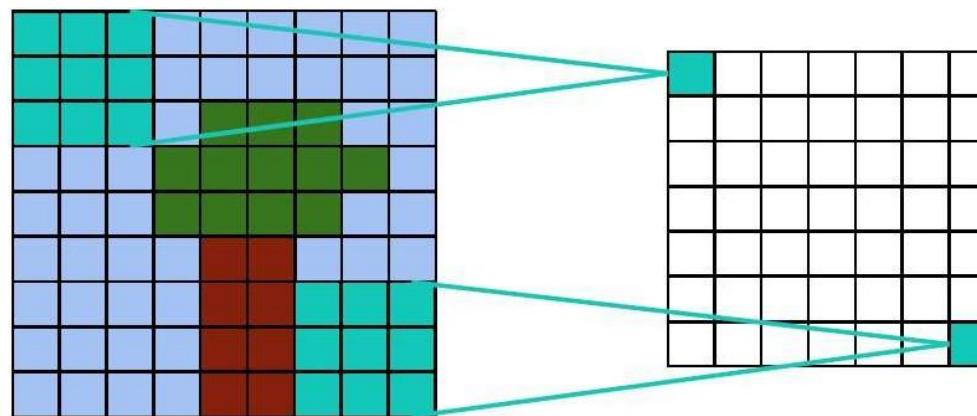
Convolution Operation

Carver Mead



Convolutional Neural Networks

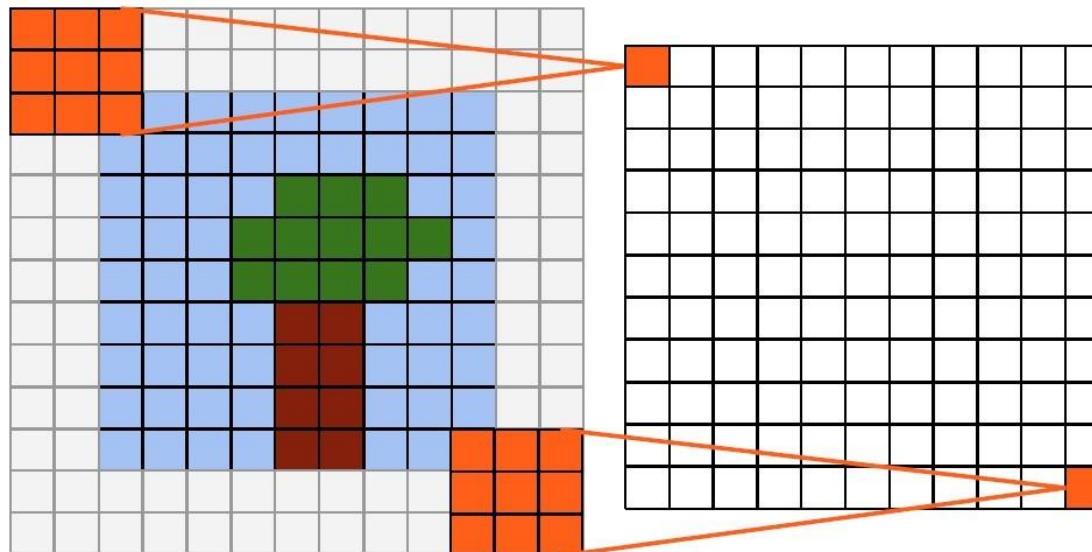
Convolution Operation: Variants



Valid convolution: $\text{output size} = \text{input size} - \text{kernel size} + 1$

Convolutional Neural Networks

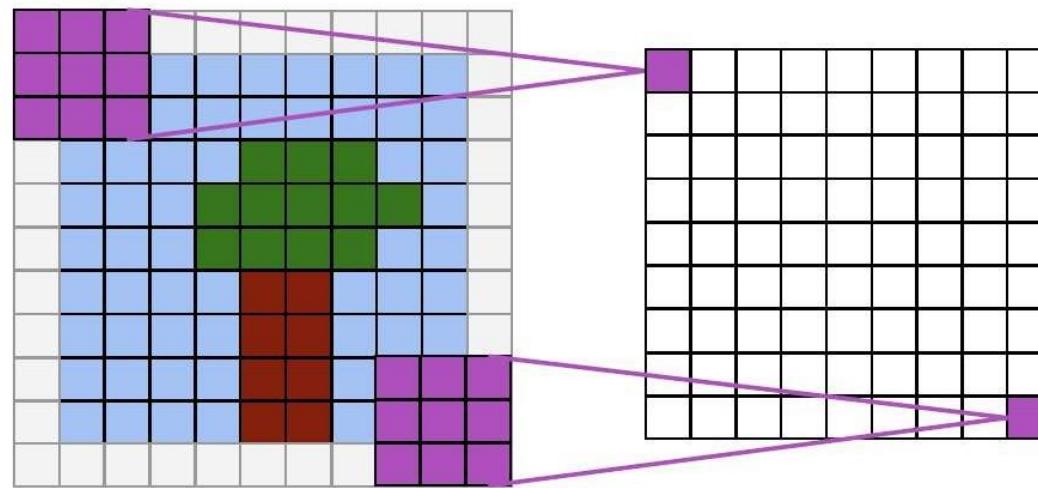
Convolution Operation: Variants



Full convolution: output size = input size + kernel size - 1

Convolutional Neural Networks

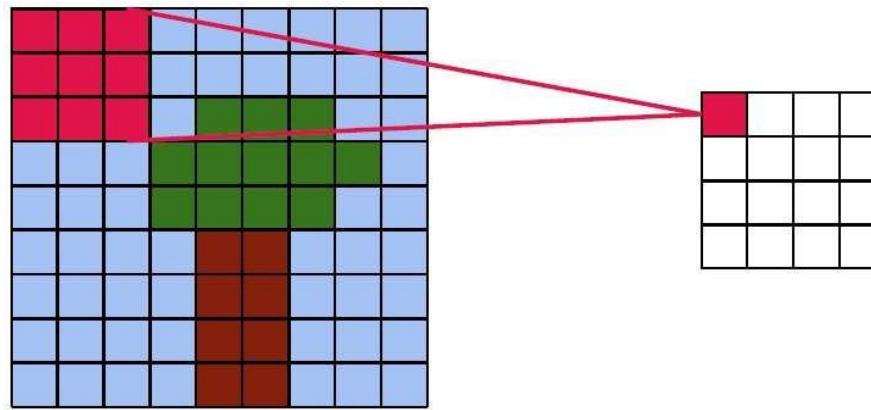
Convolution Operation: Variants



Same convolution: output size = input size

Convolutional Neural Networks

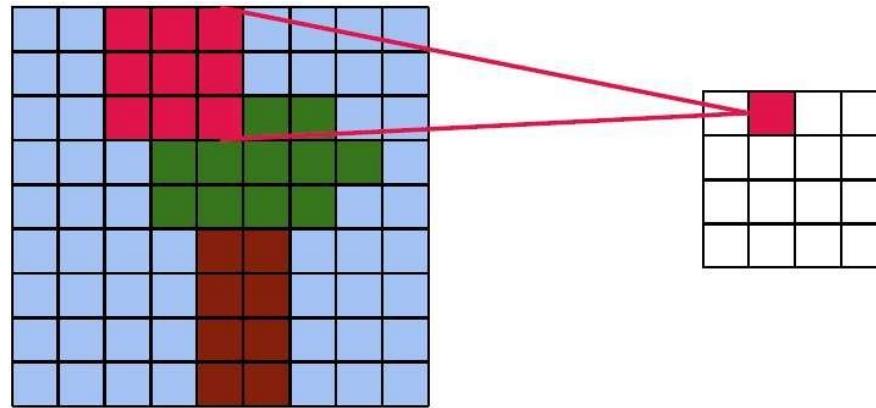
Convolution Operation: Variants



Strided convolution: kernel slides along the image with a step > 1

Convolutional Neural Networks

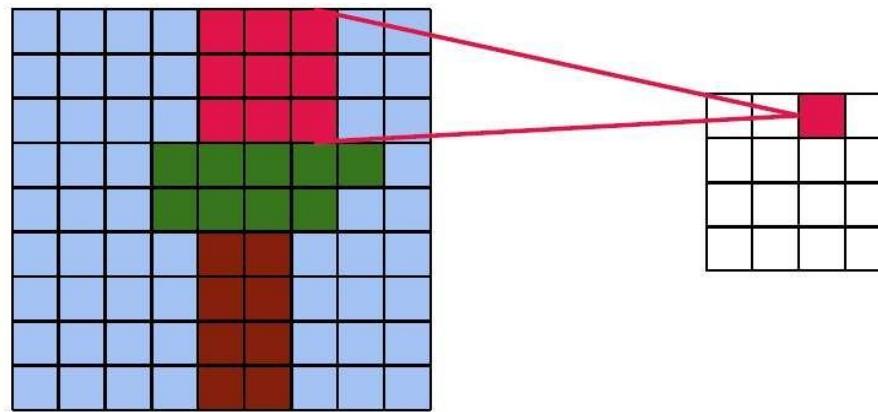
Convolution Operation: Variants



Strided convolution: kernel slides along the image with a step > 1

Convolutional Neural Networks

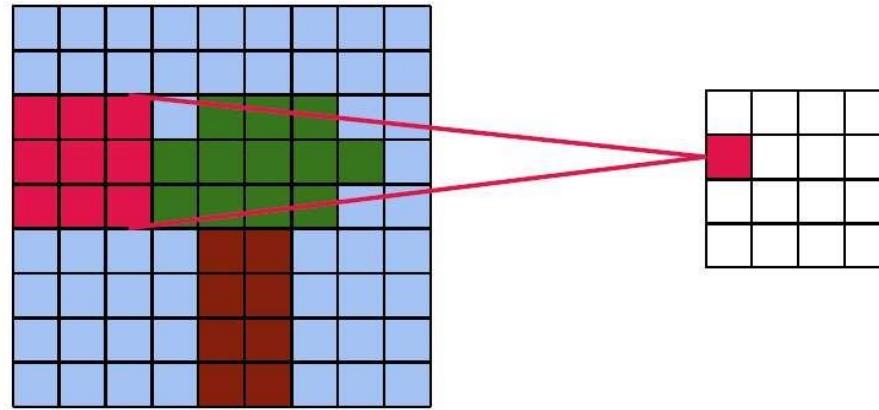
Convolution Operation: Variants



Strided convolution: kernel slides along the image with a step > 1

Convolutional Neural Networks

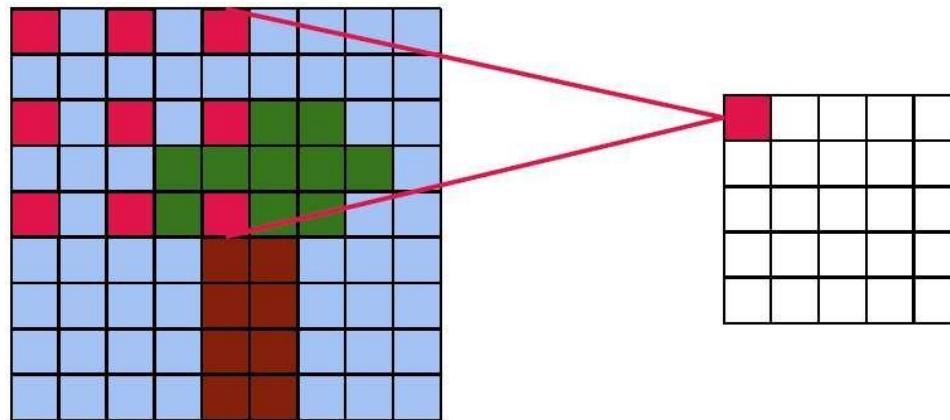
Convolution Operation: Variants



Strided convolution: kernel slides along the image with a step > 1

Convolutional Neural Networks

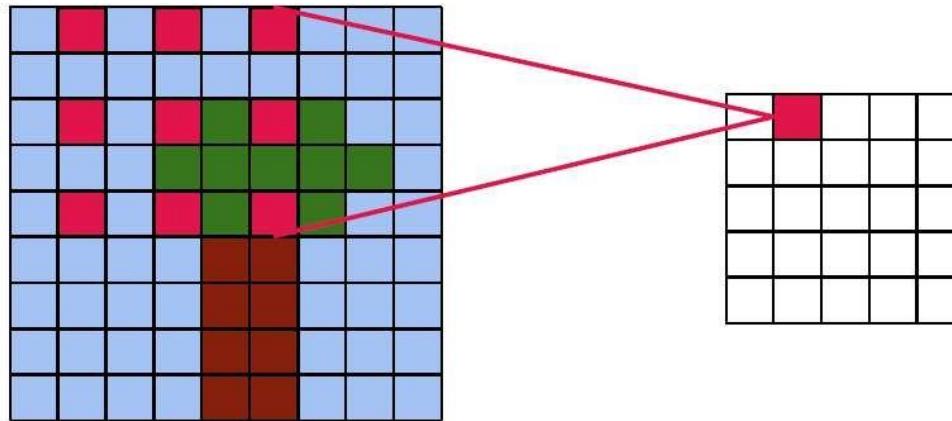
Convolution Operation: Variants



Dilated convolution: kernel is spread out, step > 1 between kernel elements

Convolutional Neural Networks

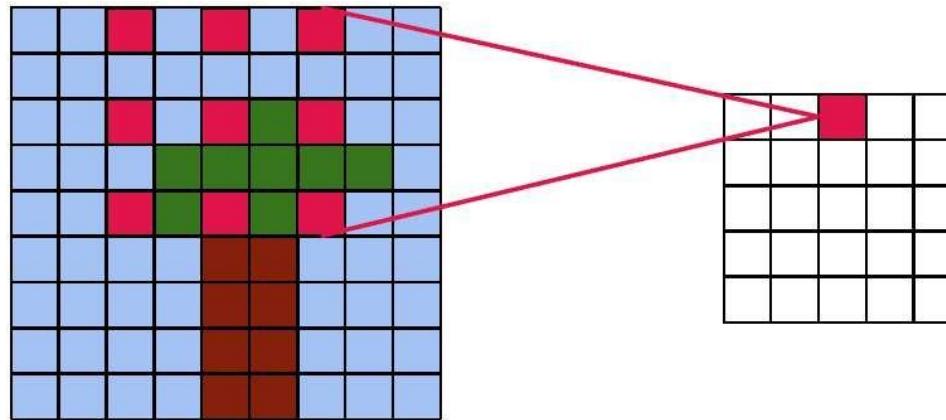
Convolution Operation: Variants



Dilated convolution: kernel is spread out, step > 1 between kernel elements

Convolutional Neural Networks

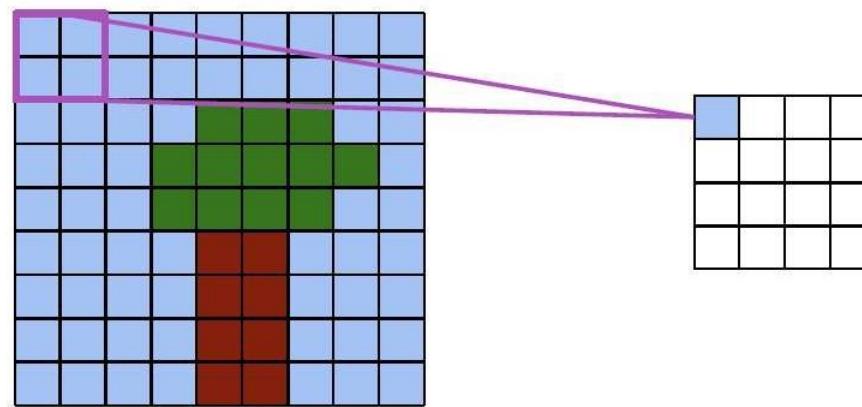
Convolution Operation: Variants



Dilated convolution: kernel is spread out, step > 1 between kernel elements

Convolutional Neural Networks

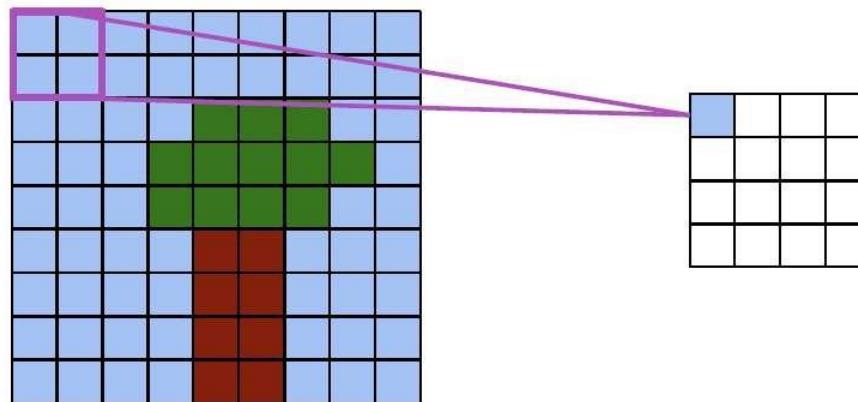
Convolution Operation: Variants



Pooling: compute mean or max over small windows to reduce resolution

Convolutional Neural Networks

Convolution Operation: Variants

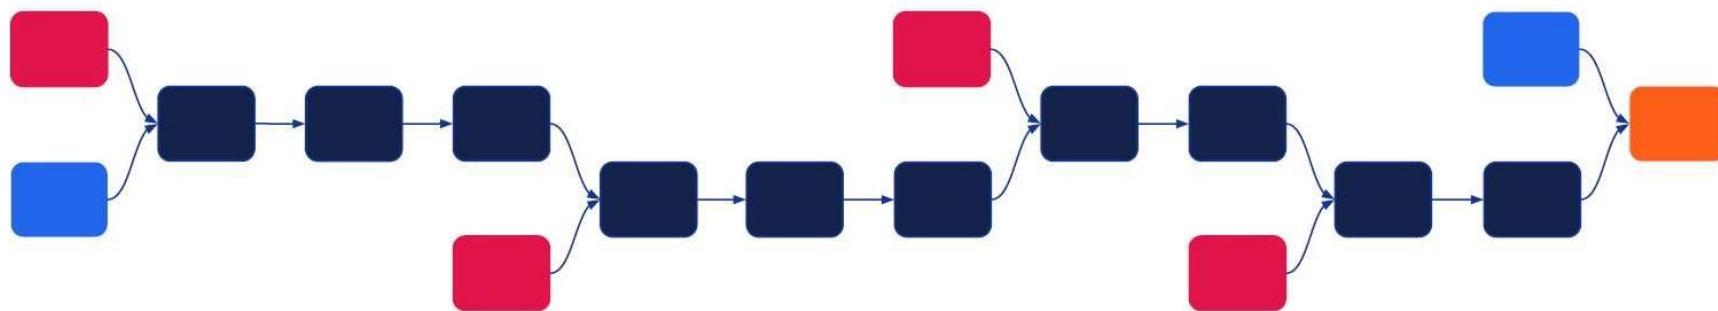


Pooling: compute mean or max over small windows to reduce resolution

Convolutional Neural Networks

Stacking the building blocks

Neural networks as computational graphs



input

loss

computation

parameters

Convolutional Neural Networks

Stacking the building blocks



Convolutional Neural Networks

Some Case Studies:Le-Net5 (1998)



Convolutional Neural Networks

Some Case Studies:Le-Net5 (1998)



Architecture of **LeNet-5**, a convnet
for handwritten digit recognition

Want to learn more?

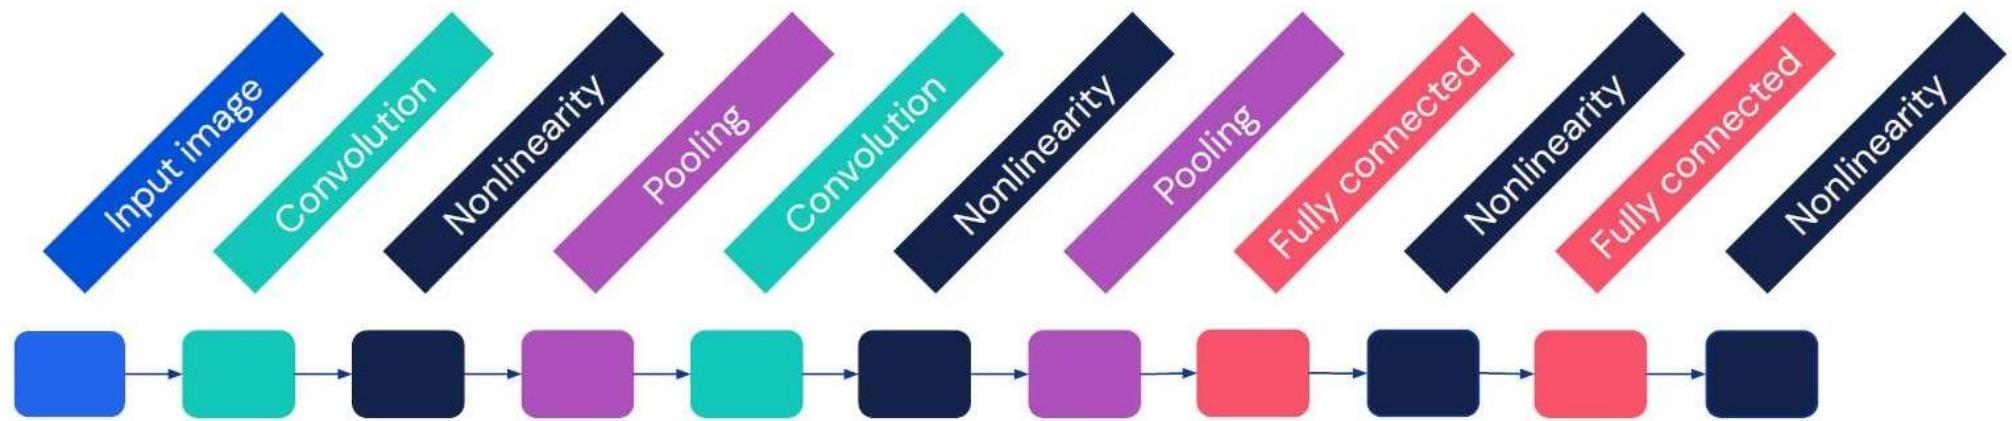


Lecun, Y.; Bottou, L.; Bengio, Y.; Haffner, P.
Gradient-based learning applied to document recognition
Proceedings of the IEEE 86(11) (1998)



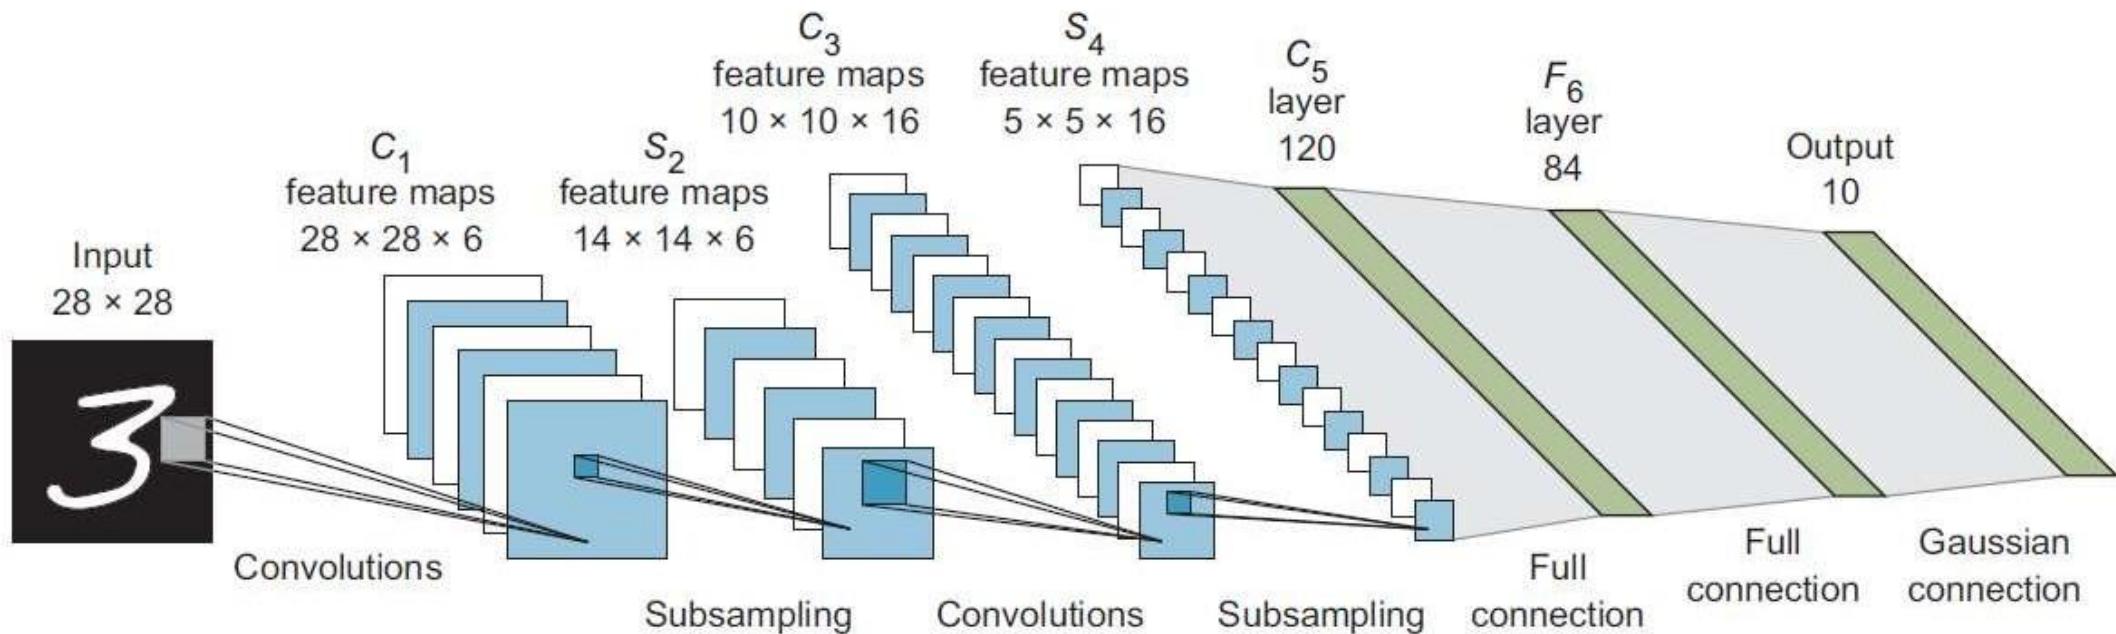
Convolutional Neural Networks

Some Case Studies:Le-Net5 (1998)



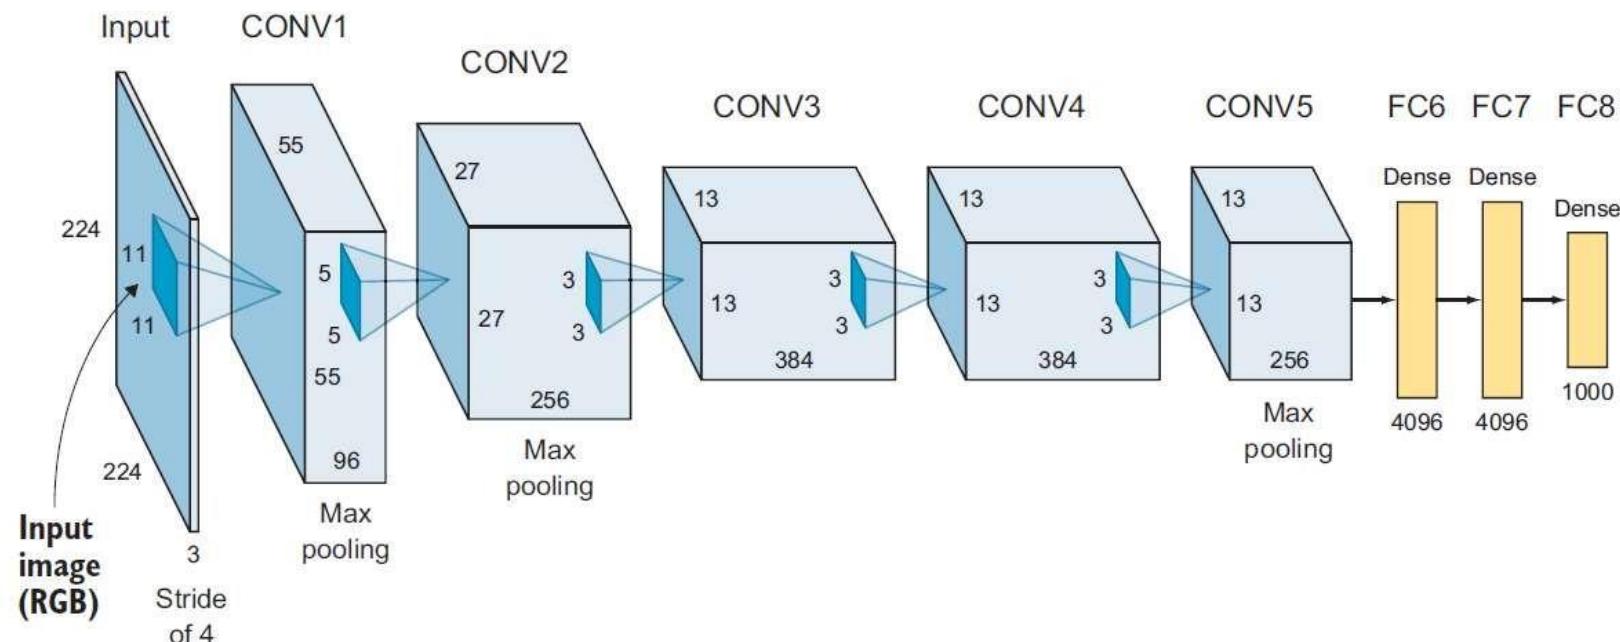
Convolutional Neural Networks

Some Case Studies:Le-Net5 (1998)



Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



Want to learn more?



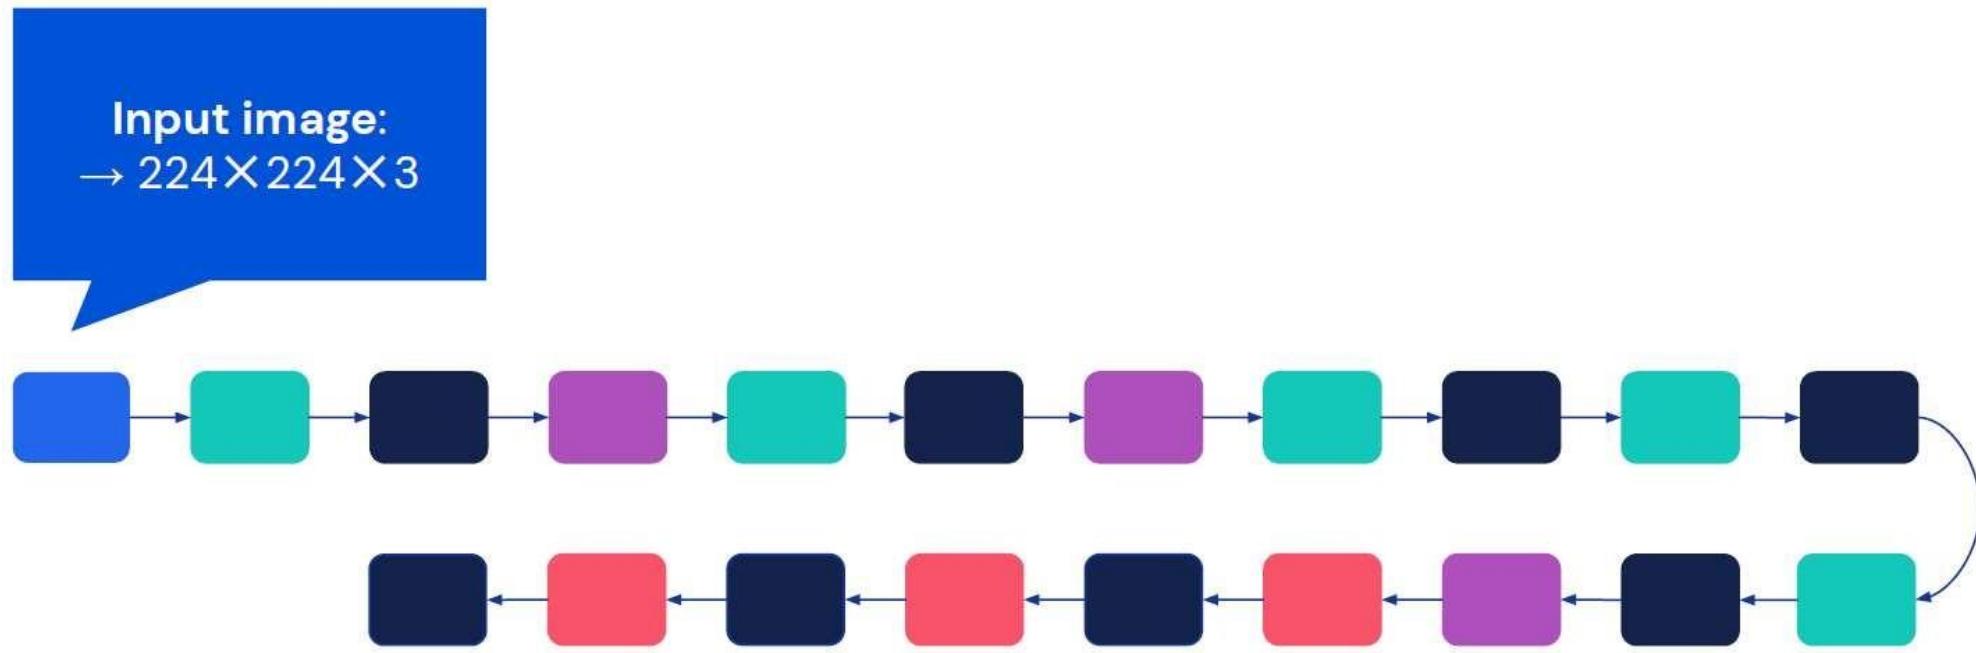
Krizhevsky, A.; Sutskever, I.; Hinton, G.E.
ImageNet classification with deep convolutional neural networks
Neural Information Processing Systems (2012)

Architecture: 8 layers, ReLU, dropout, weight decay

Infrastructure: large dataset, trained 6 days on 2 GPUs

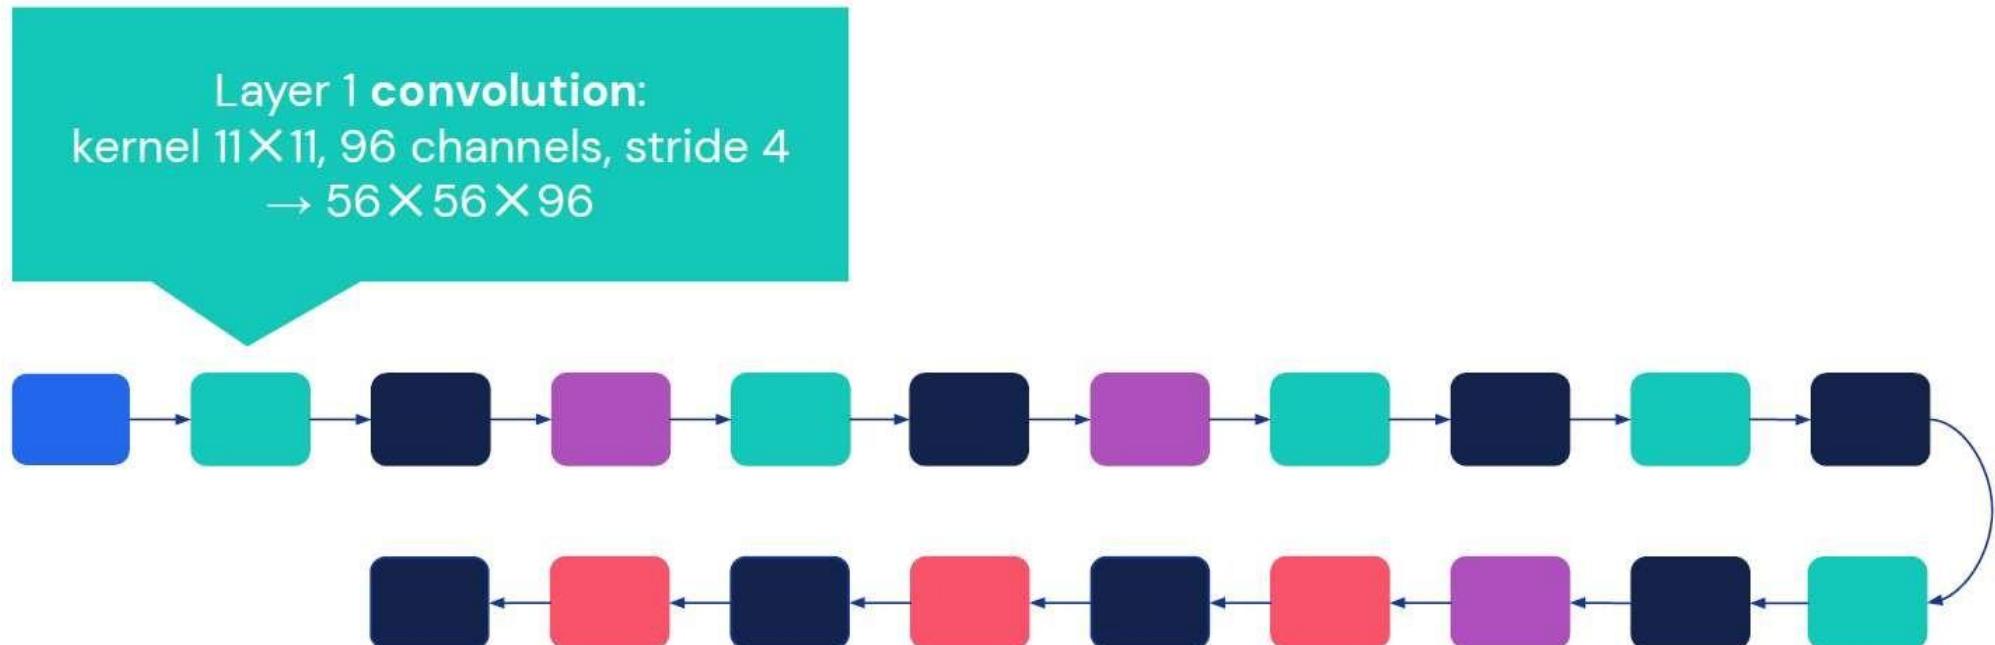
Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



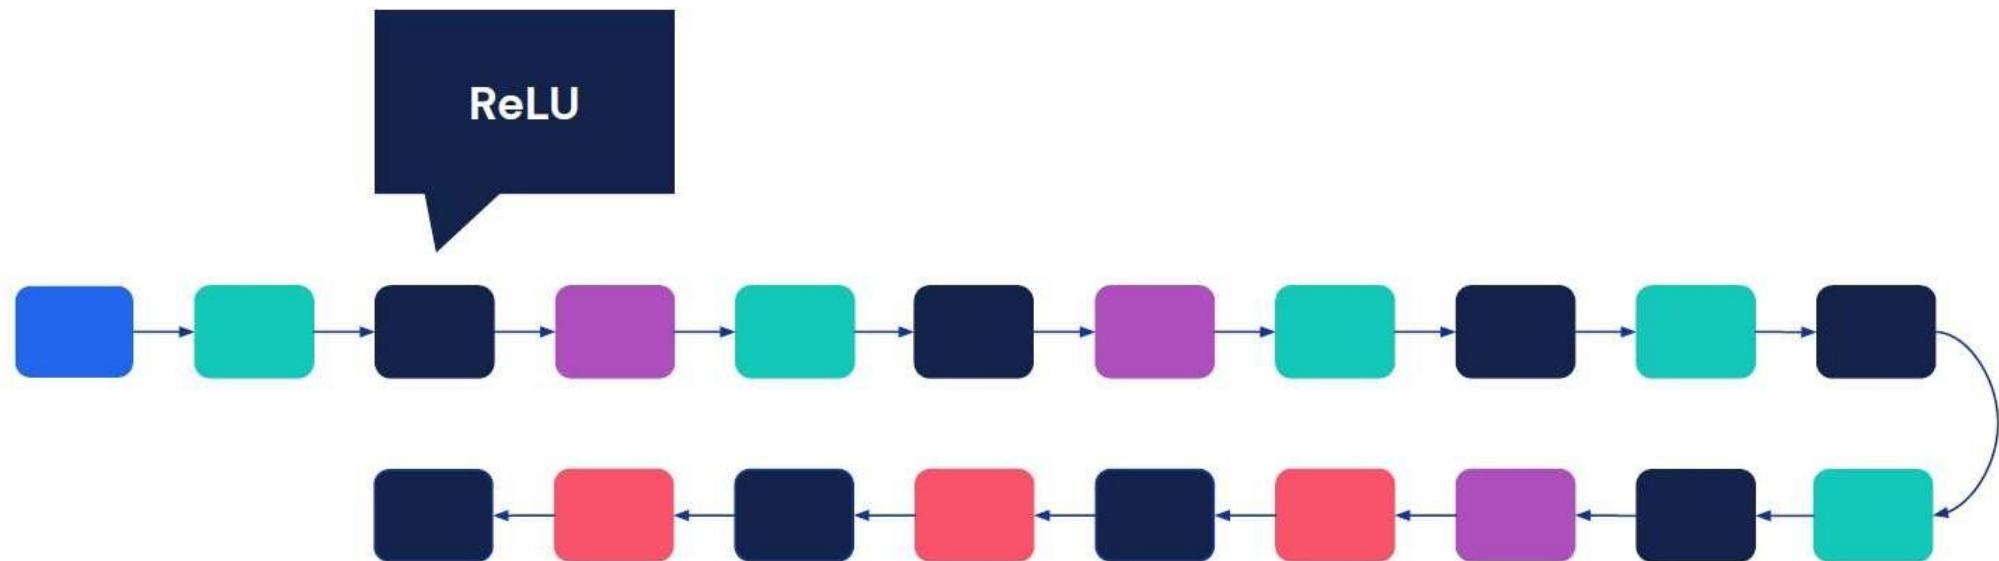
Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



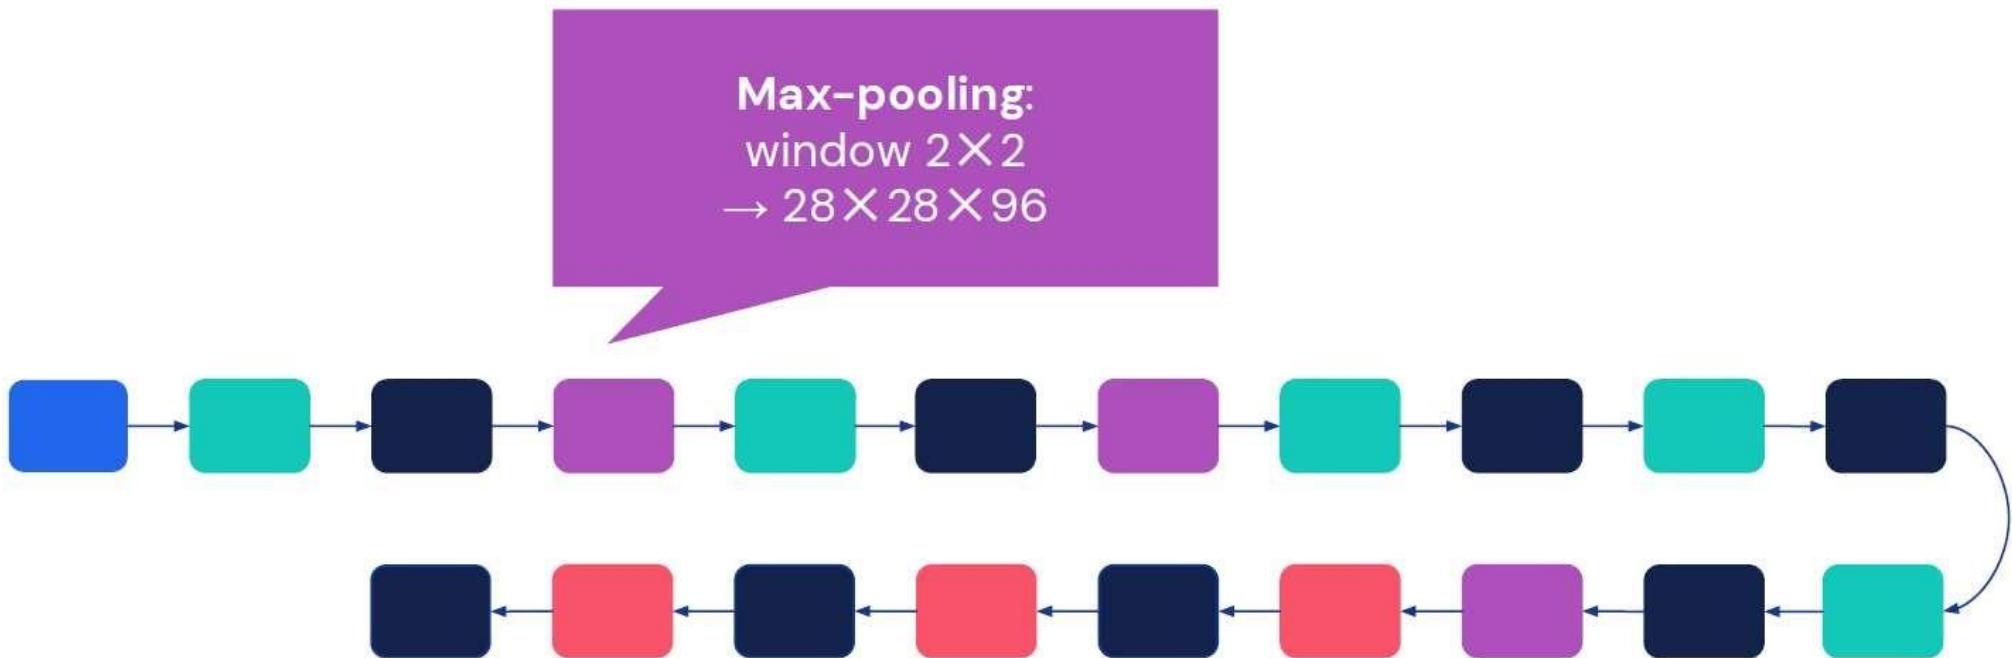
Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



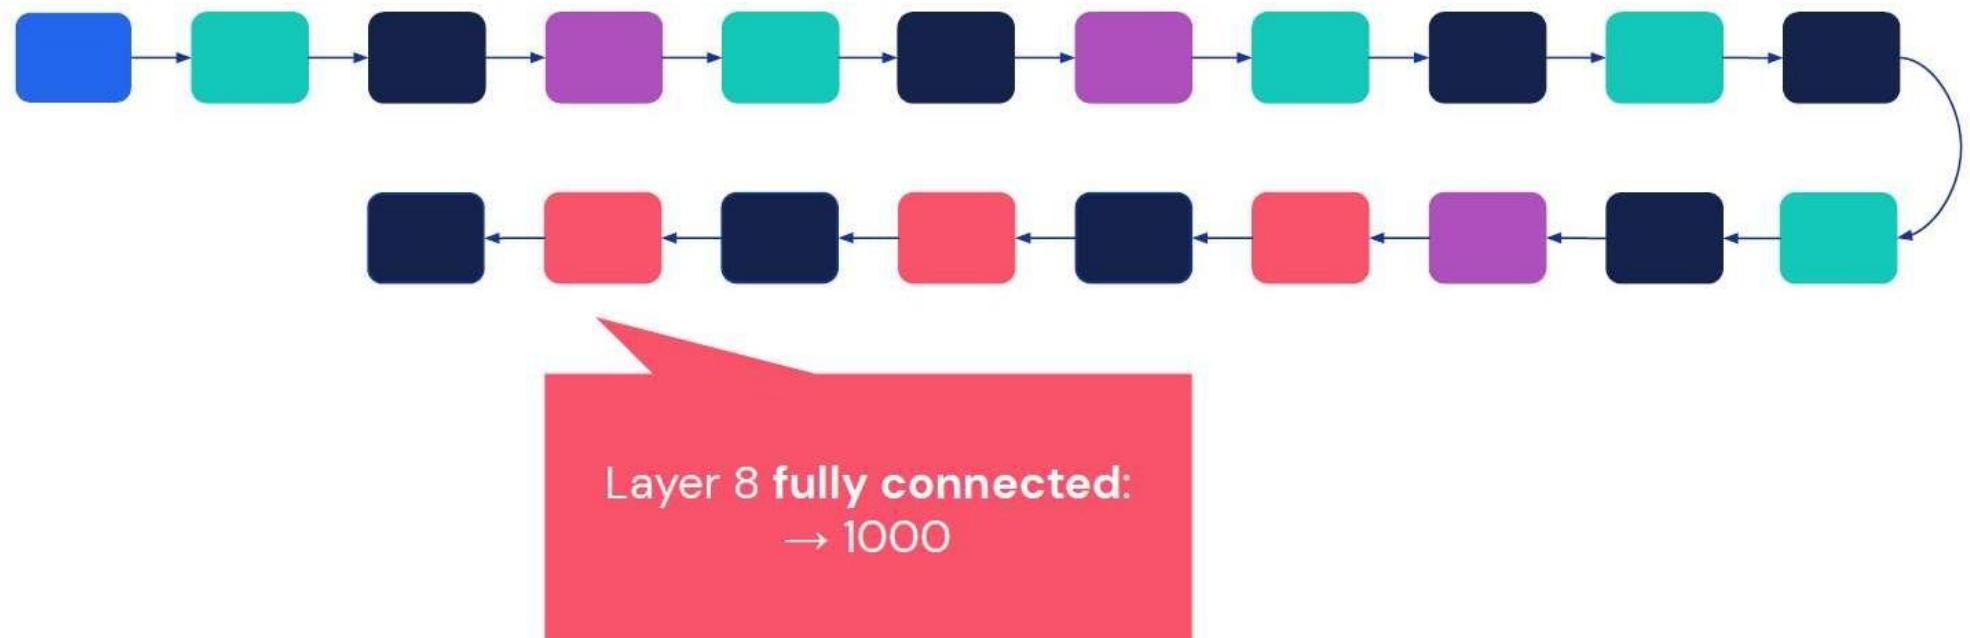
Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



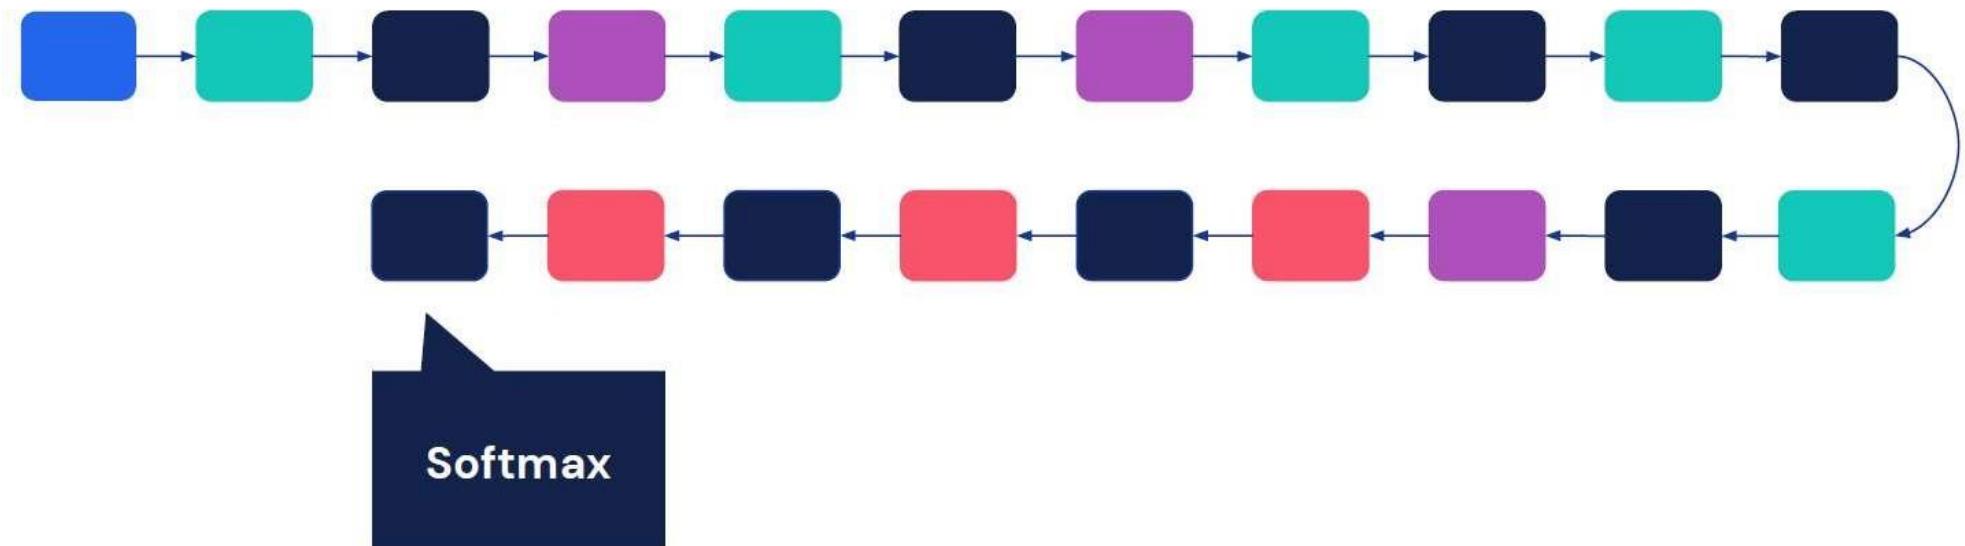
Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



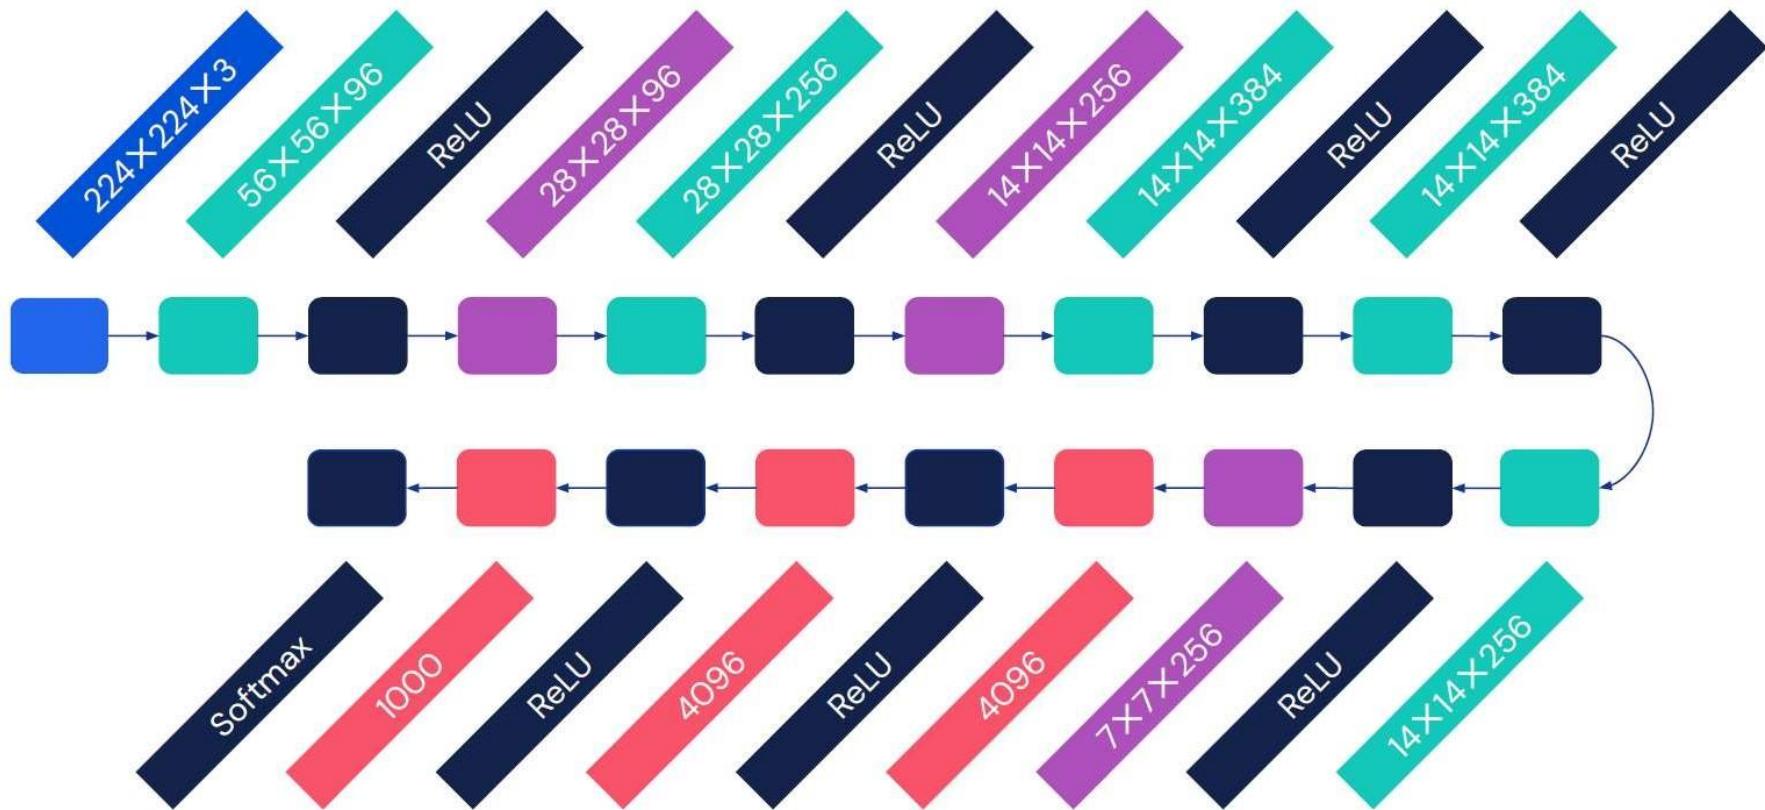
Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



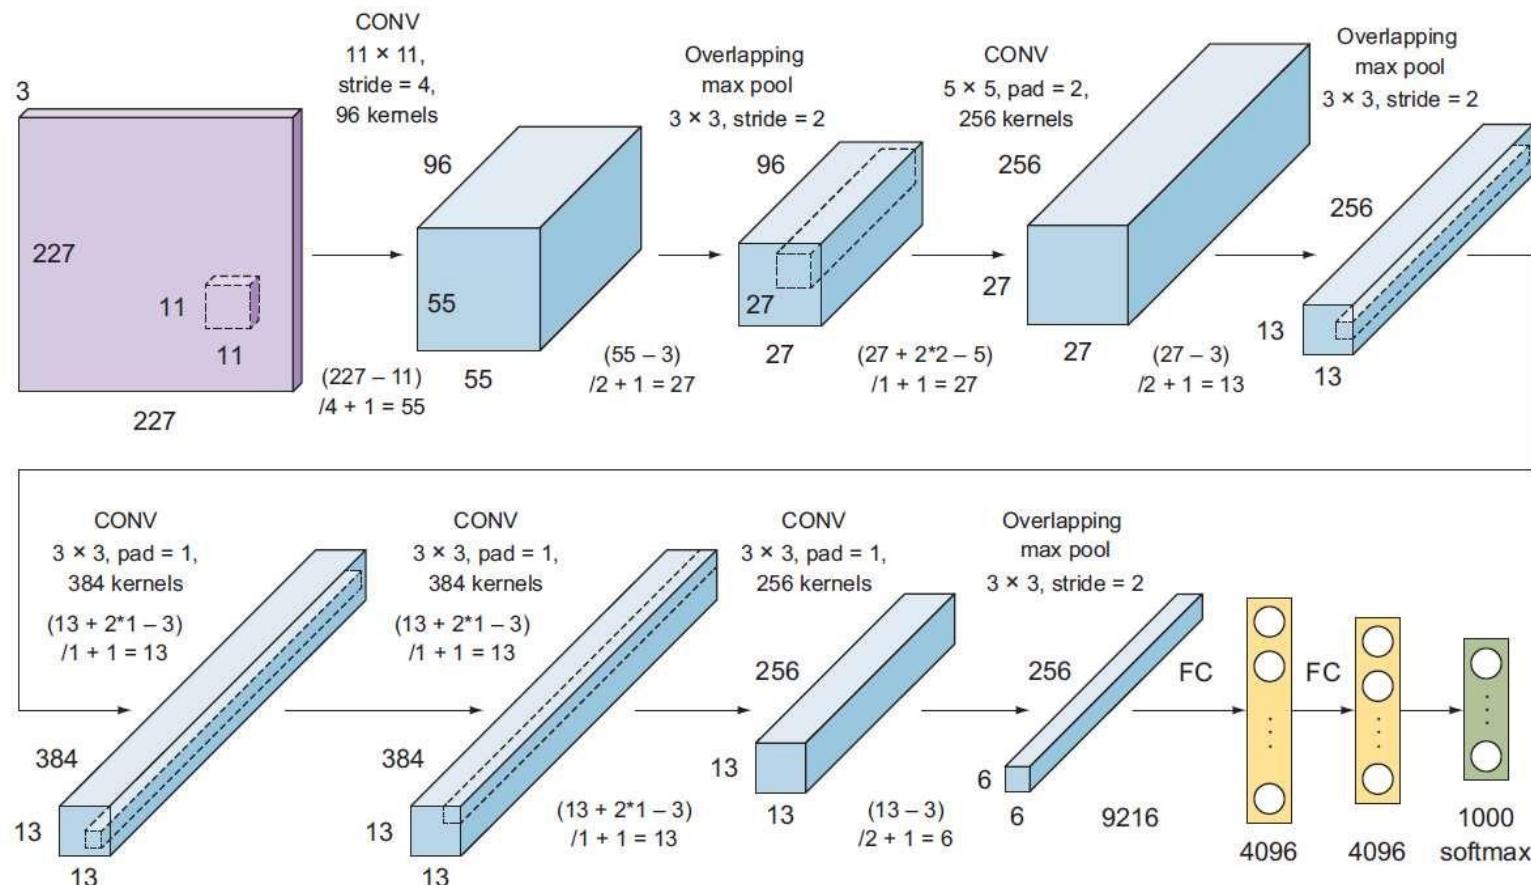
Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



Convolutional Neural Networks

Some Case Studies: AlexNet (2012)



Convolutional Neural Networks

Some Case Studies: AlexNet (2012)

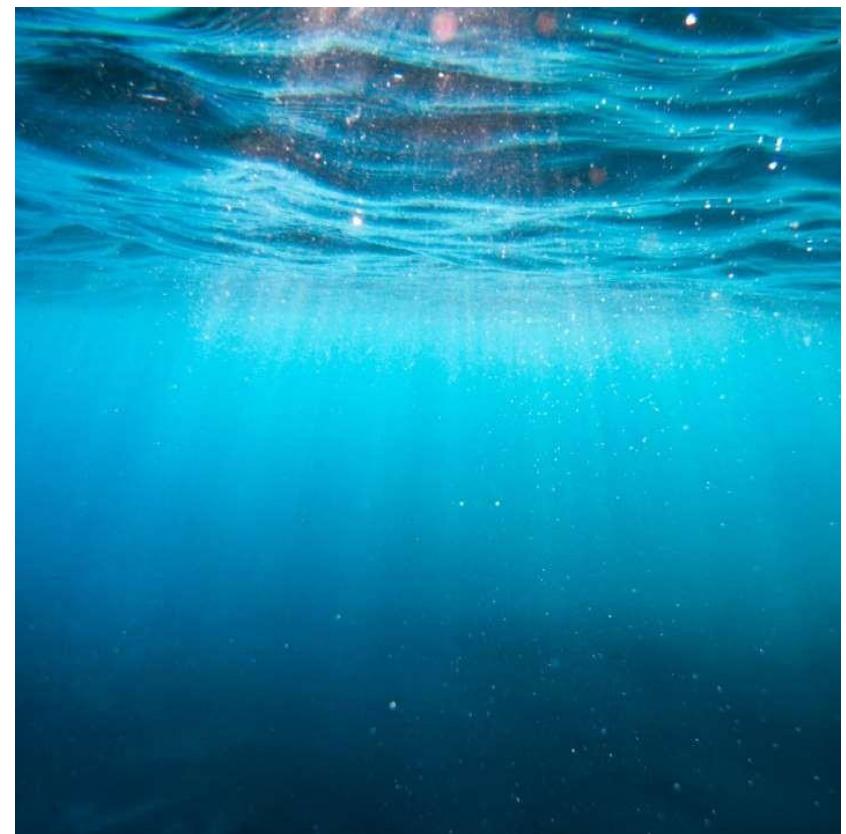
Stacking the building blocks

Each layer is a linear classifier by itself

More layers: more non-linearities

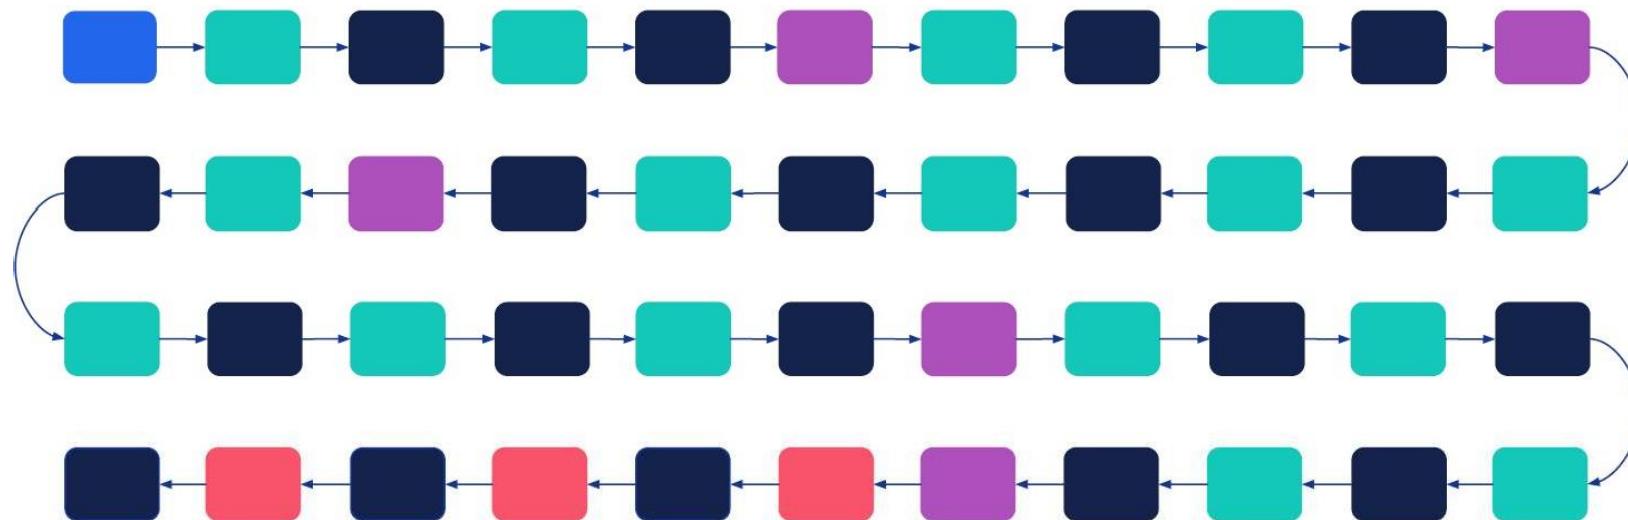
What limits the number of layers of ConvNets?

Computation, Optimization, Dissection



Convolutional Neural Networks

Some Case Studies: VGGNet (2014)



Want to learn more?

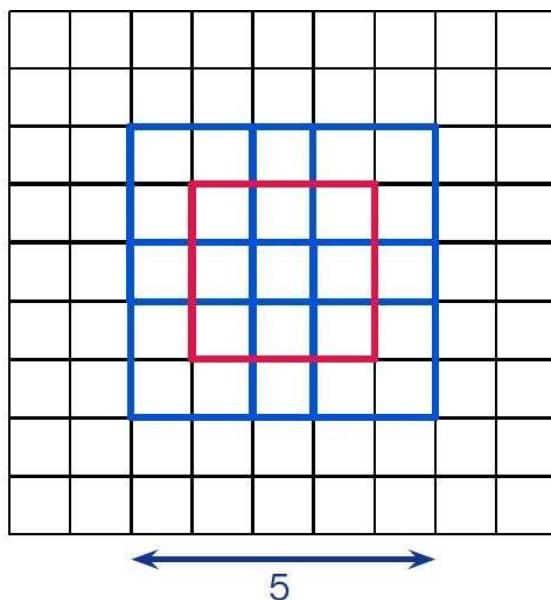


Simonyan, K.; Zisserman, A.
*Very deep convolutional networks for
large-scale image recognition* International
Conference on Learning Representations (2015)

**Stack many convolutional layers before pooling
Use “same” convolutions to avoid resolution reduction**

Convolutional Neural Networks

Some Case Studies: VGGNet (2014)



Stacking 3×3 kernels



1st 3×3 conv. layer



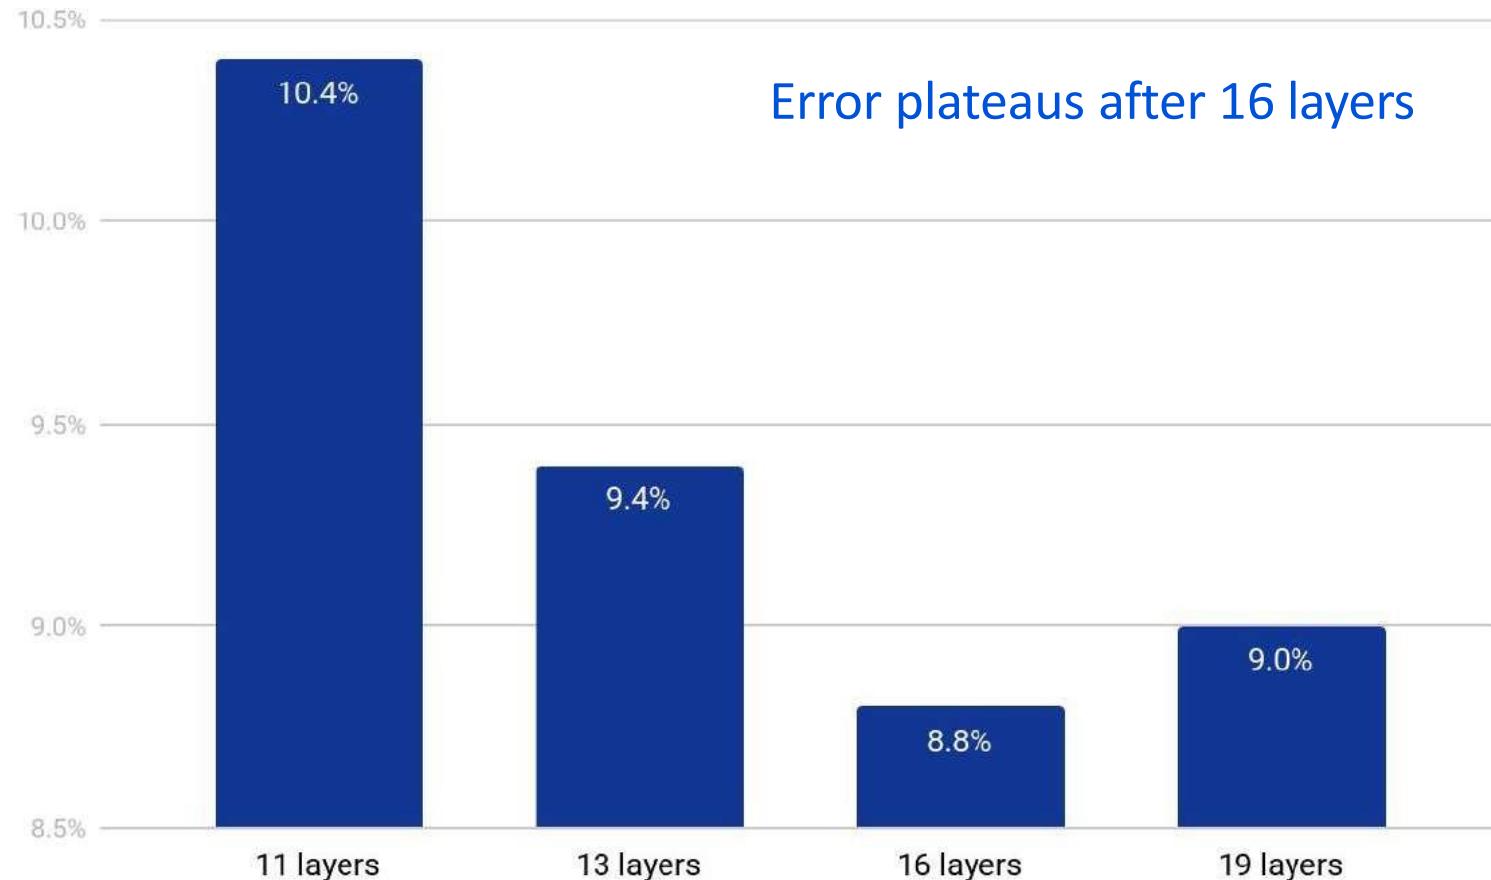
2nd 3×3 conv. layer

Architecture: up to 19 layers, 3×3 kernels only, “same” convolutions

Infrastructure: trained for 2–3 weeks on 4 GPUs (data parallelism)

Convolutional Neural Networks

Some Case Studies: VGGNet (2014)



Convolutional Neural Networks

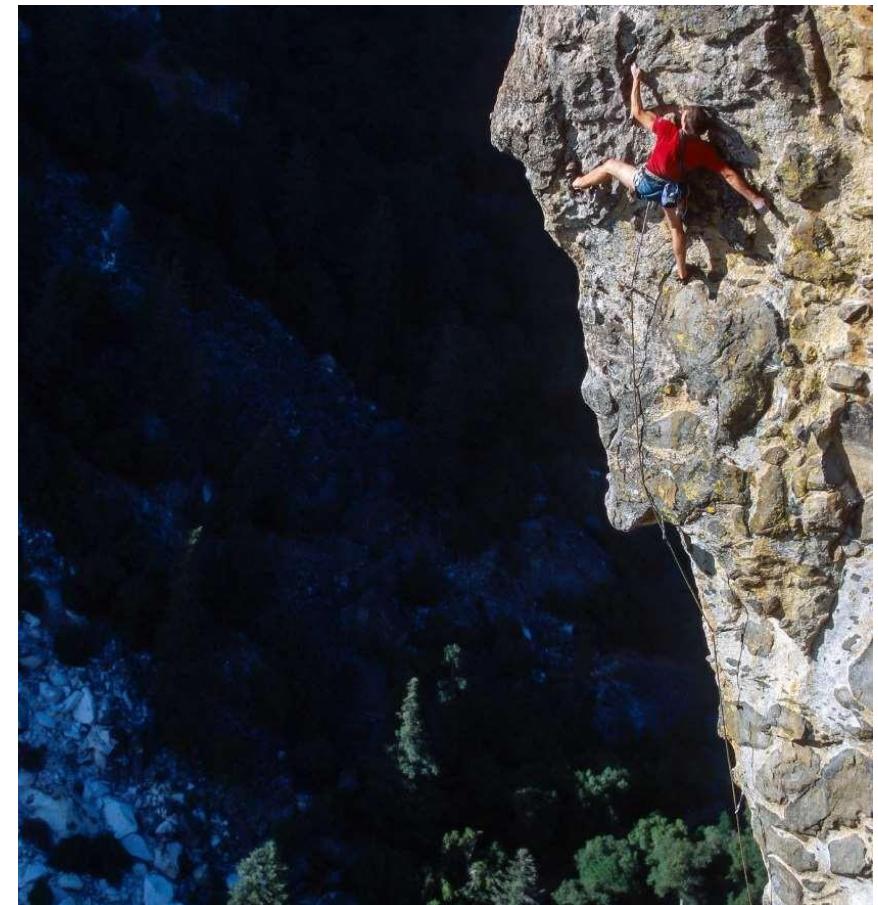
Some Case Studies:VGGNet (2014)

Challenges of depth

Computational Complexity
Optimisation difficulties

Improving optimisation

Careful Initialisation
Sophisticated Optimizers
Normalization Layers
Network Design



Convolutional Neural Networks

Some Case Studies: GoogleLeNet (2014)

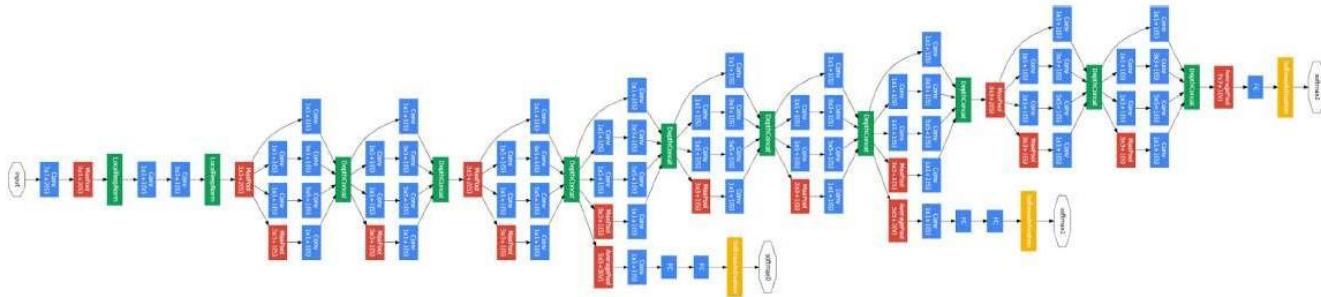
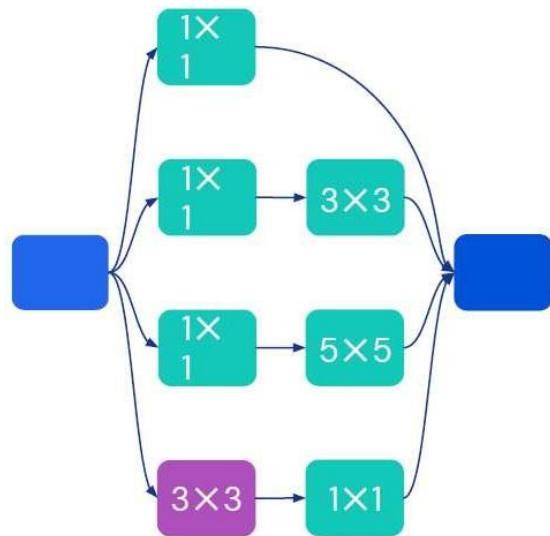


Figure from Szegedy et al. (2015)



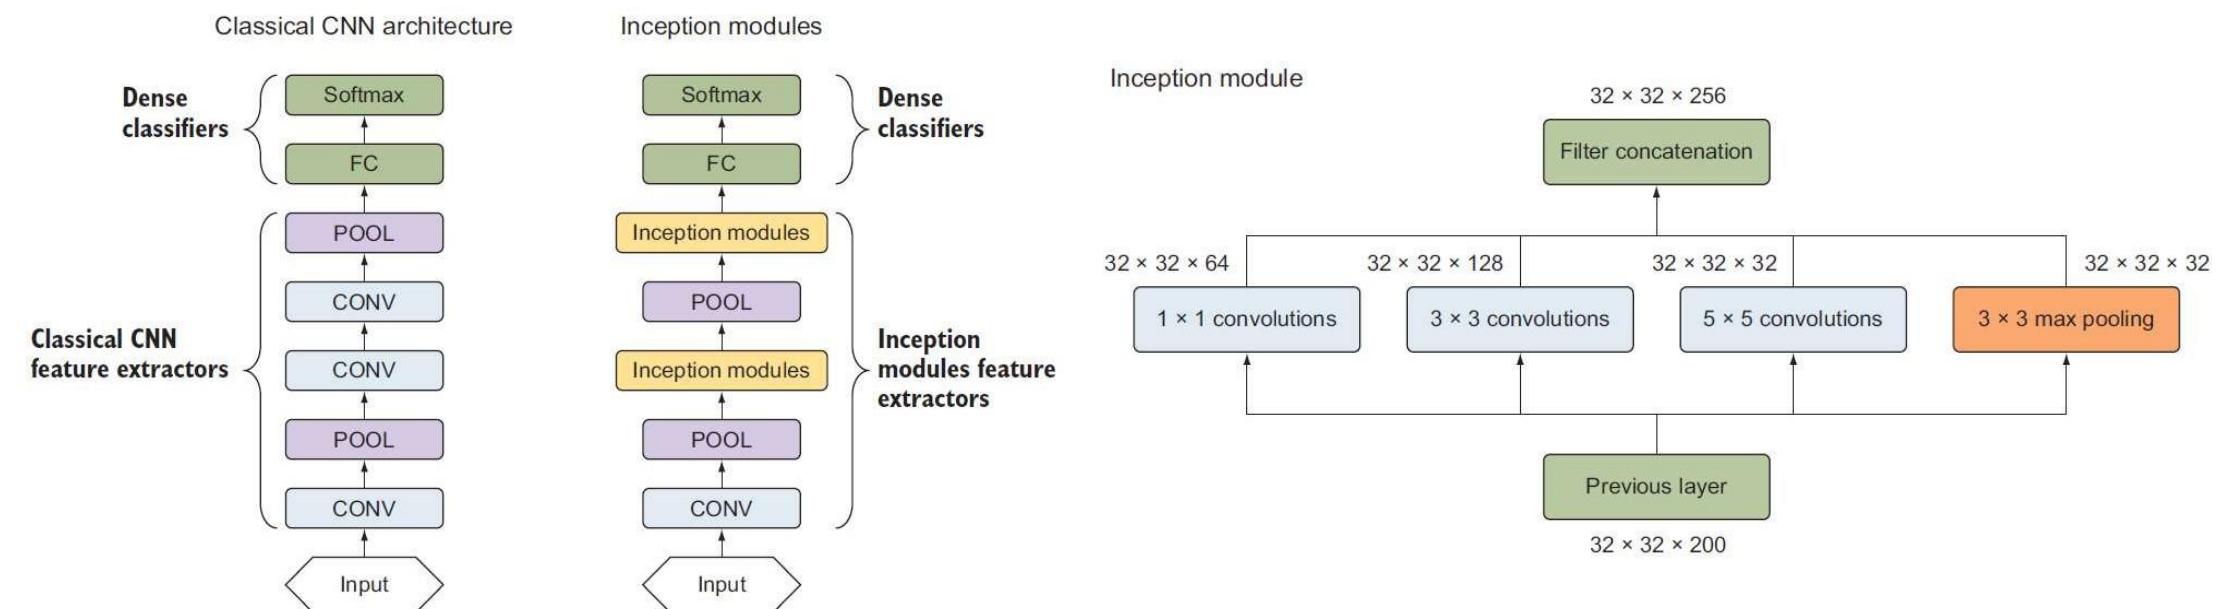
Want to learn more?



Szegedy, C. et al.
Going deeper with convolutions IEEE conference on computer vision and pattern recognition (2015)

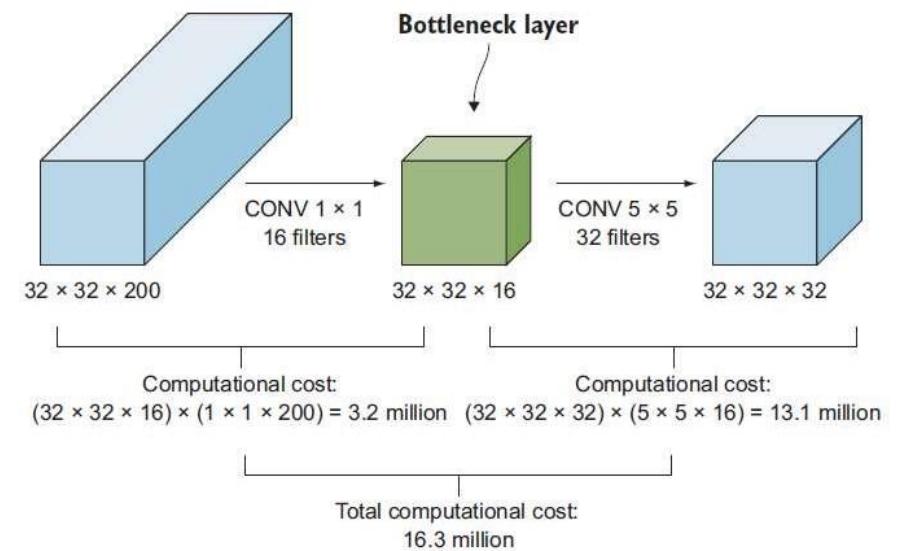
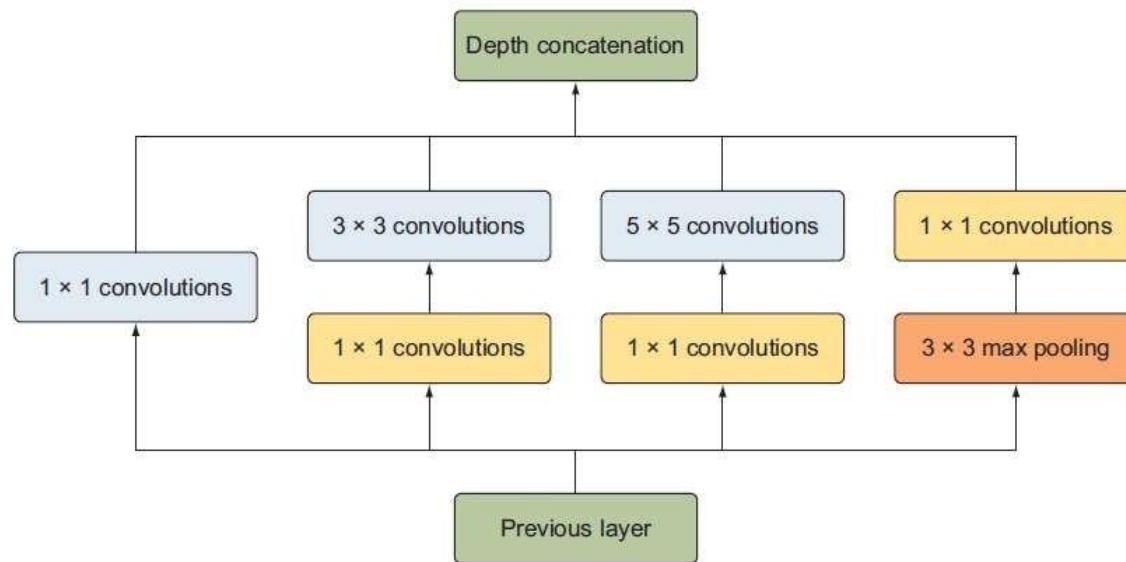
Convolutional Neural Networks

Some Case Studies: GoogleLeNet (2014)



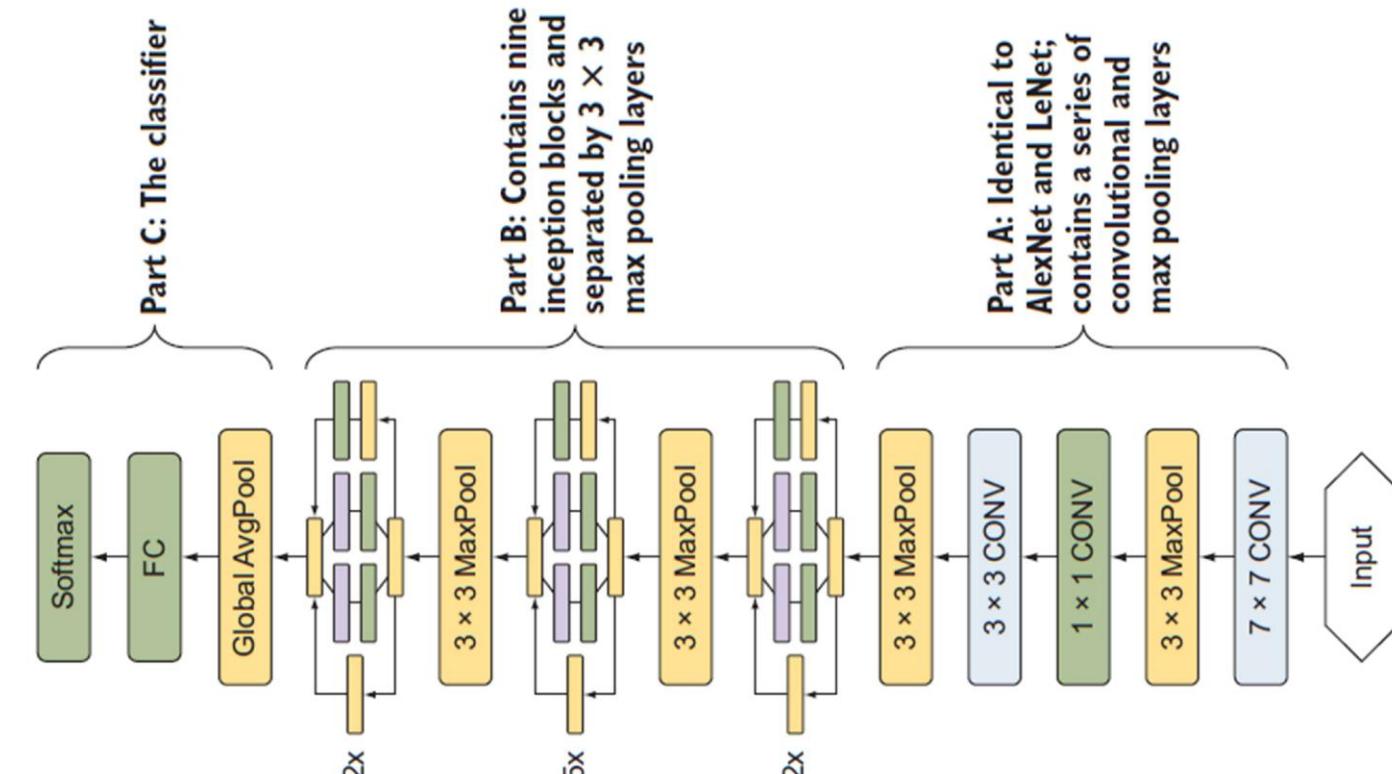
Convolutional Neural Networks

Some Case Studies: GoogleLeNet (2014)



Convolutional Neural Networks

Some Case Studies: GoogleLeNet (2014)



Convolutional Neural Networks

Some Case Studies: GoogleLeNet (2014)

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;
Parameters to be learned: γ, β
Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\begin{aligned}\mu_{\mathcal{B}} &\leftarrow \frac{1}{m} \sum_{i=1}^m x_i && // \text{mini-batch mean} \\ \sigma_{\mathcal{B}}^2 &\leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 && // \text{mini-batch variance} \\ \hat{x}_i &\leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} && // \text{normalize} \\ y_i &\leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) && // \text{scale and shift}\end{aligned}$$

Figure from Ioffe et al. (2015)

Want to learn more?



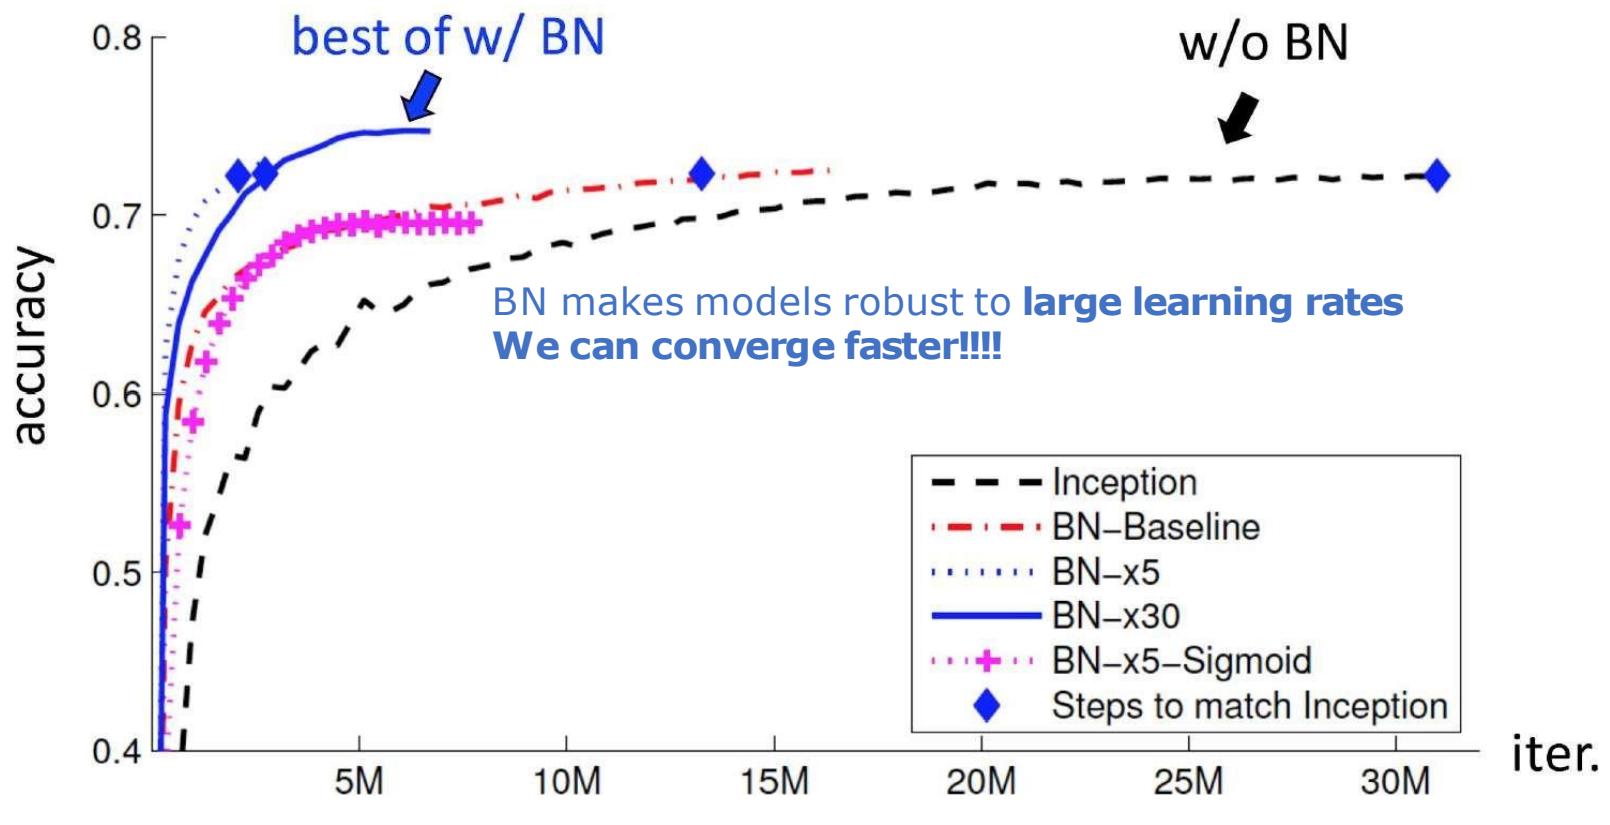
Ioffe, S.; Szegedy, C.
Batch normalization: Accelerating deep network training by reducing internal covariate shift. International conference on machine learning (2015)

Reduces sensitivity to **initialisation**
Introduces **Noise** and makes robust to **large learning rates**
Introduces stochasticity and acts as a **regulariser**

Convolutional Neural Networks

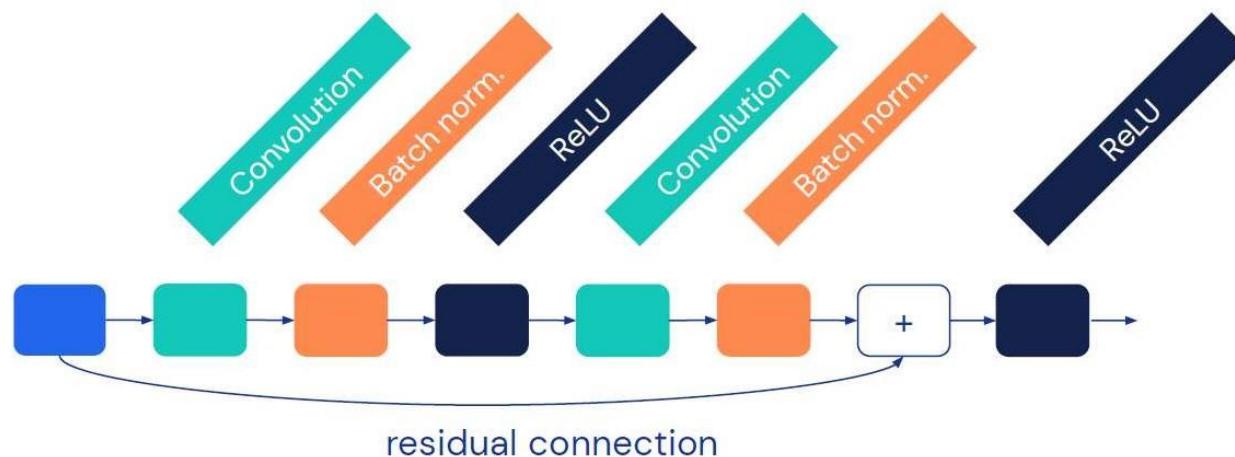
Some Case Studies: GoogleLeNet (2014)

Batch Normalization



Convolutional Neural Networks

Some Case Studies:ResNet (2015)



Want to learn more?

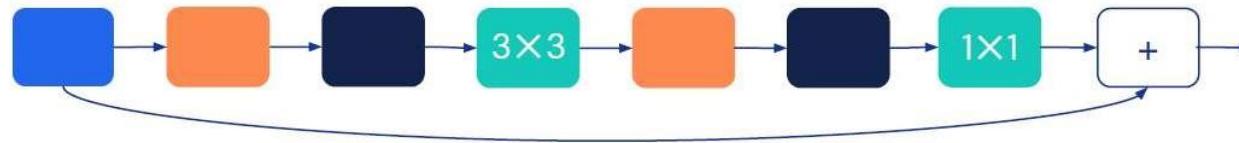
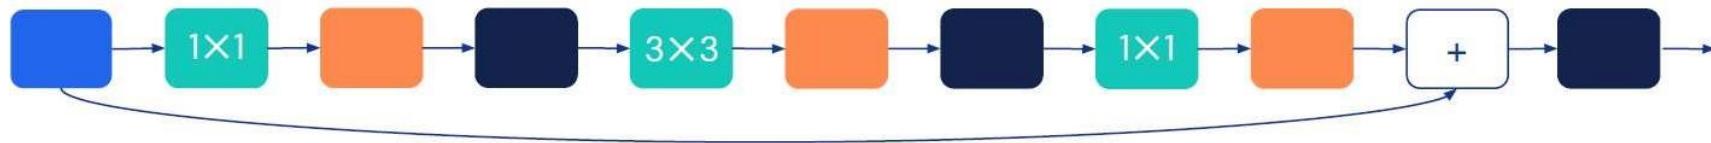
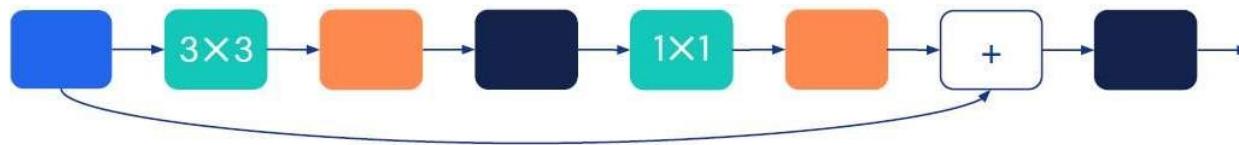


He, K. et al.
Deep residual learning for image recognition
IEEE conference on computer vision and pattern recognition (2016)

Residual connections facilitate training deeper networks

Convolutional Neural Networks

Some Case Studies:ResNet (2015)



Want to learn more?



He, K. et al.
Identity mappings in deep residual networks
European conference on computer vision
(2016)

ResNet V2 (bottom) avoids all
nonlinearities in the residual pathway

Convolutional Neural Networks

Some Case Studies:ResNet (2015)

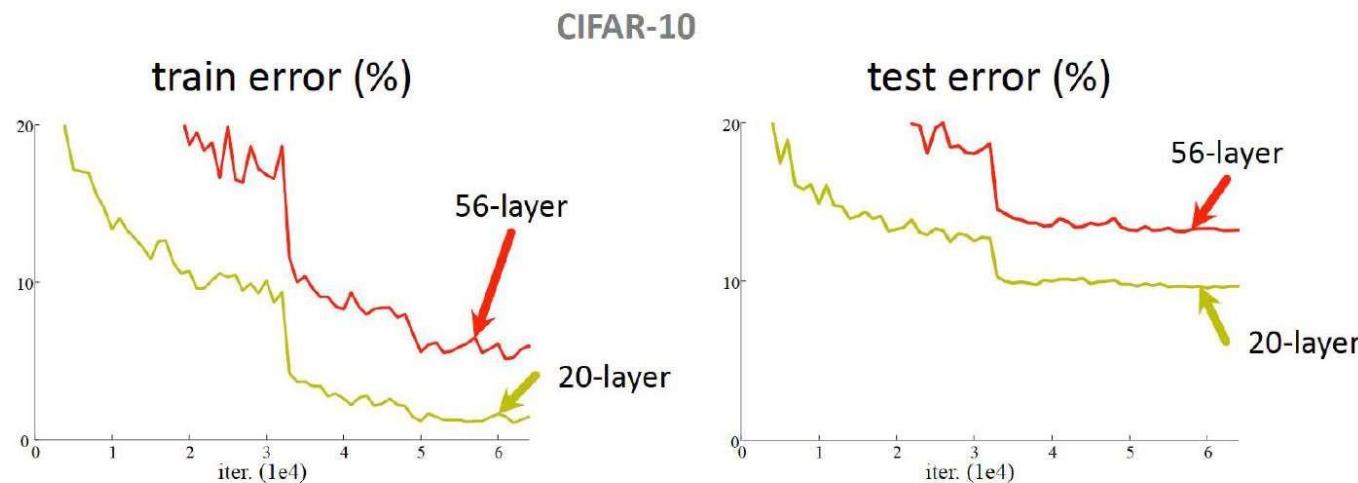
| layer name | output size | 18-layer | 34-layer | 50-layer | 101-layer | 152-layer |
|------------|-------------|---|---|---|--|--|
| conv1 | 112×112 | | | 7×7, 64, stride 2 | | |
| | | | | 3×3 max pool, stride 2 | | |
| conv2_x | 56×56 | $\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$ |
| conv3_x | 28×28 | $\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$ | $\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$ | $\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$ | $\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$ |
| conv4_x | 14×14 | $\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$ | $\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$ | $\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$ | $\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$ |
| conv5_x | 7×7 | $\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$ |
| | 1×1 | | | average pool, 1000-d fc, softmax | | |
| FLOPs | | 1.8×10^9 | 3.6×10^9 | 3.8×10^9 | 7.6×10^9 | 11.3×10^9 |

Table from He et al. (2015)

Convolutional Neural Networks

Some Case Studies:ResNet (2015)

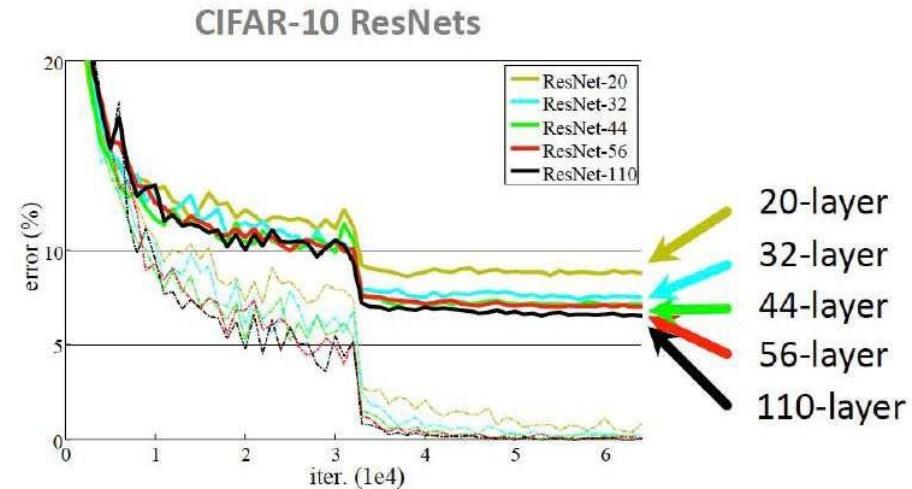
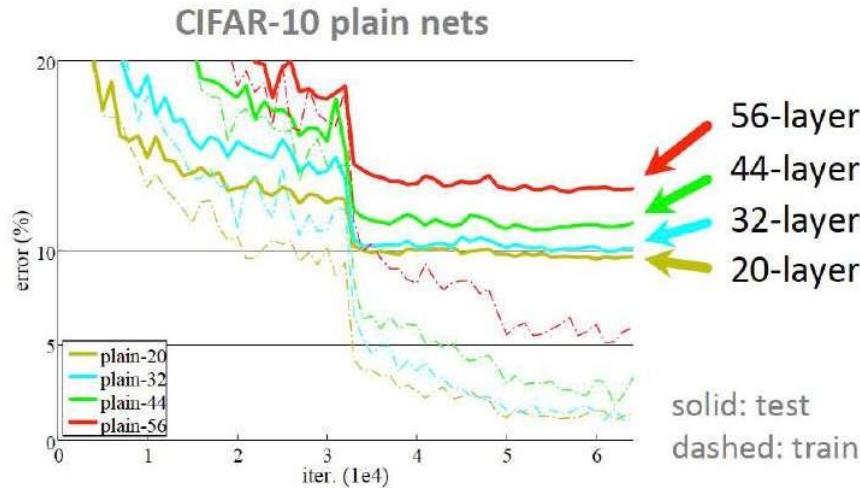
Simply stacking layers?



- Plain nets: stacking 3x3 conv layers
- 56-layer net has **higher training error and test error** than 20-layer net
- A deeper model should not have higher training error

Convolutional Neural Networks

Some Case Studies:ResNet (2015)



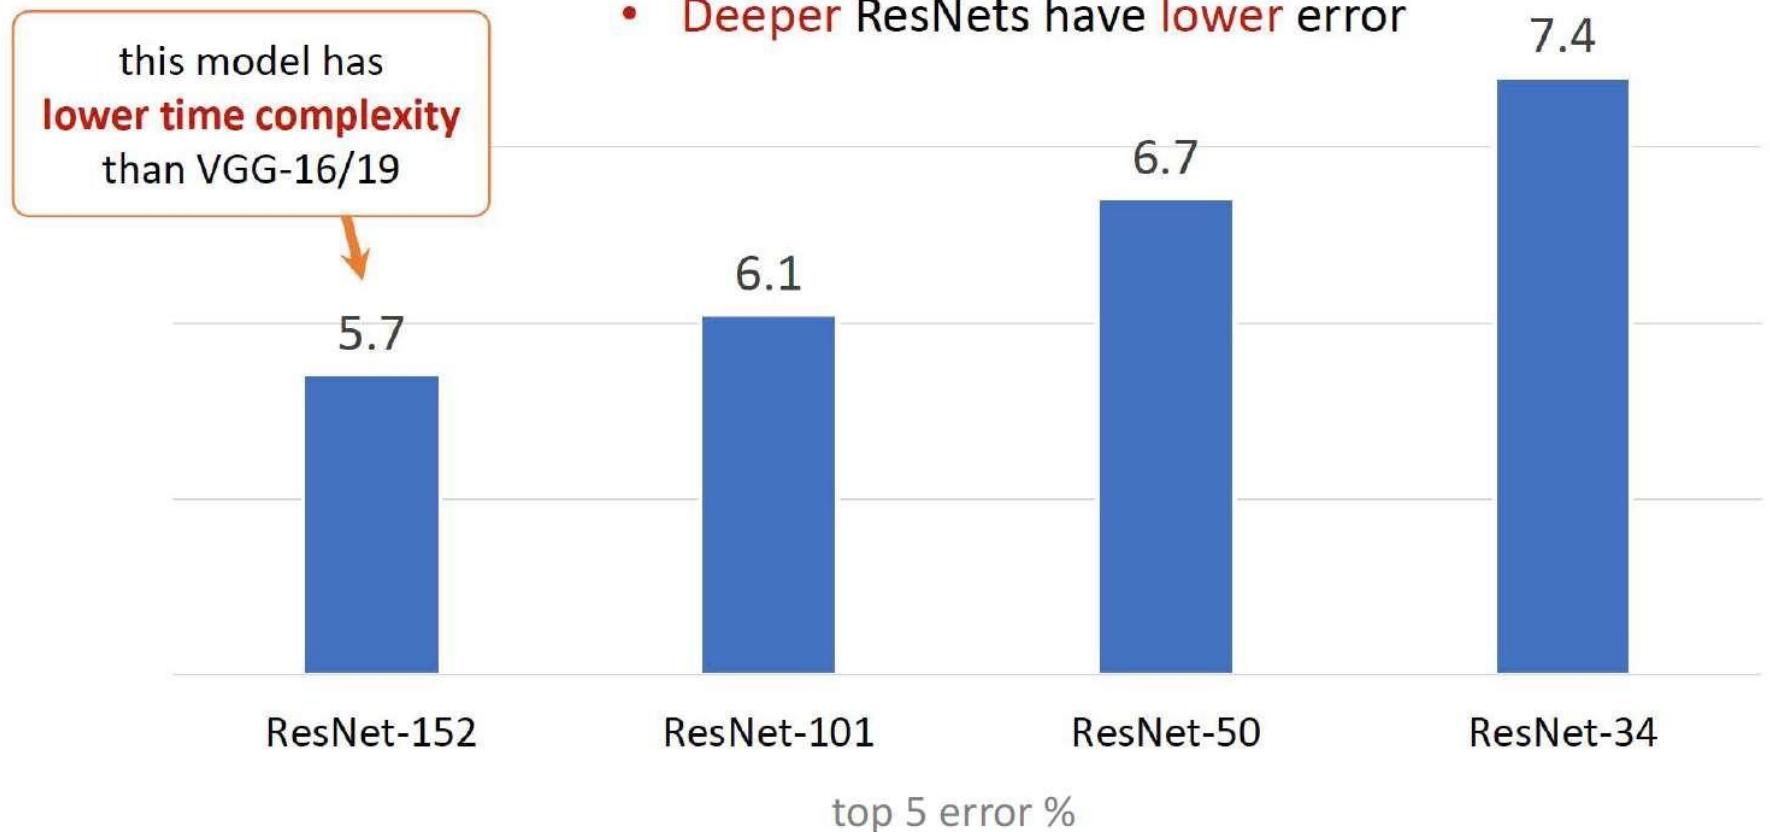
- Deep ResNets can be trained easier
- Deeper ResNets have lower training error, and also lower test error

Convolutional Neural Networks

Some Case Studies:ResNet (2015)

this model has
lower time complexity
than VGG-16/19

- Deeper ResNets have **lower** error



8

7

6

5

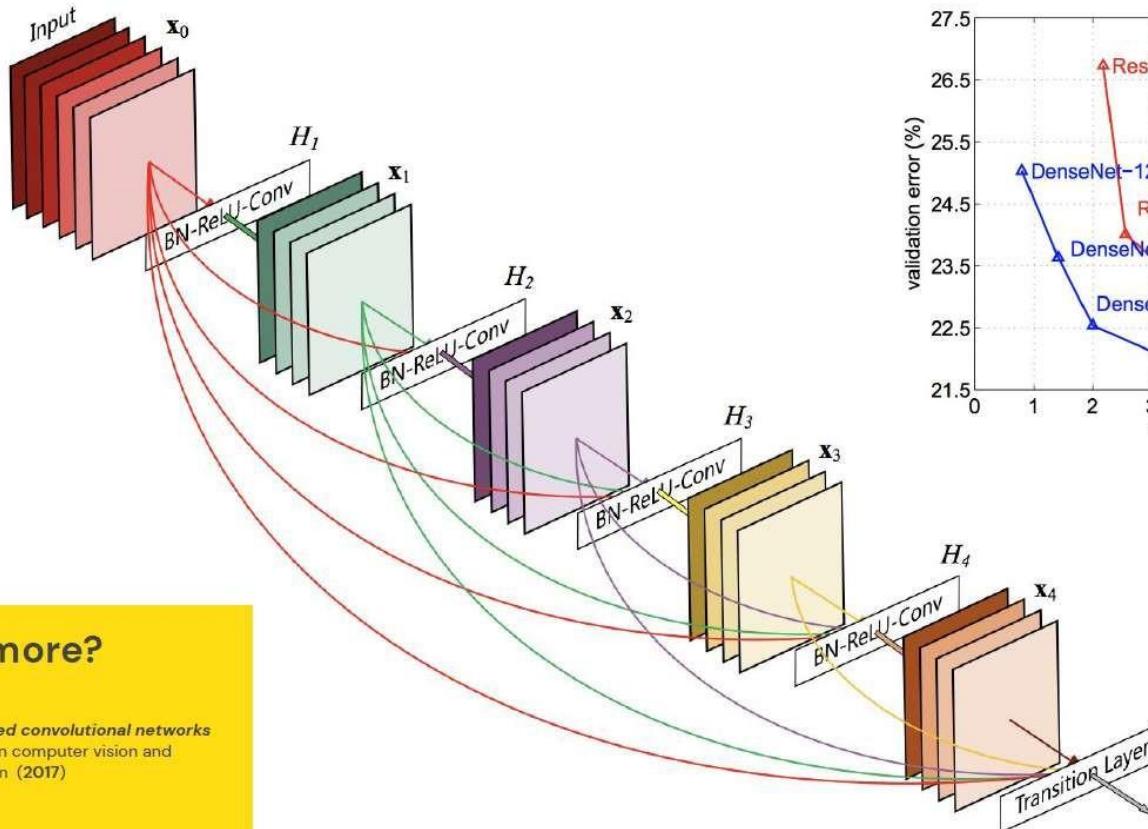
4

13

3

Convolutional Neural Networks

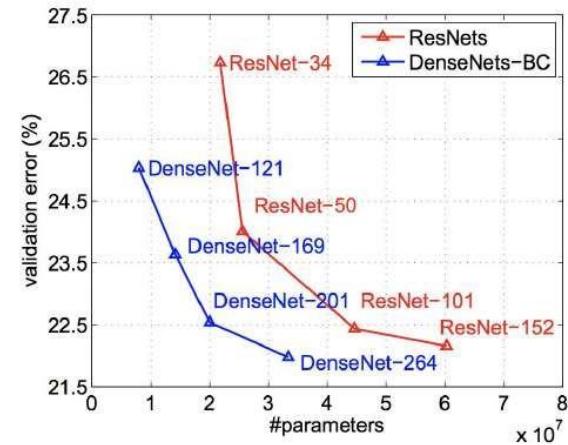
Some Case Studies:DenseNet (2016)



Want to learn more?



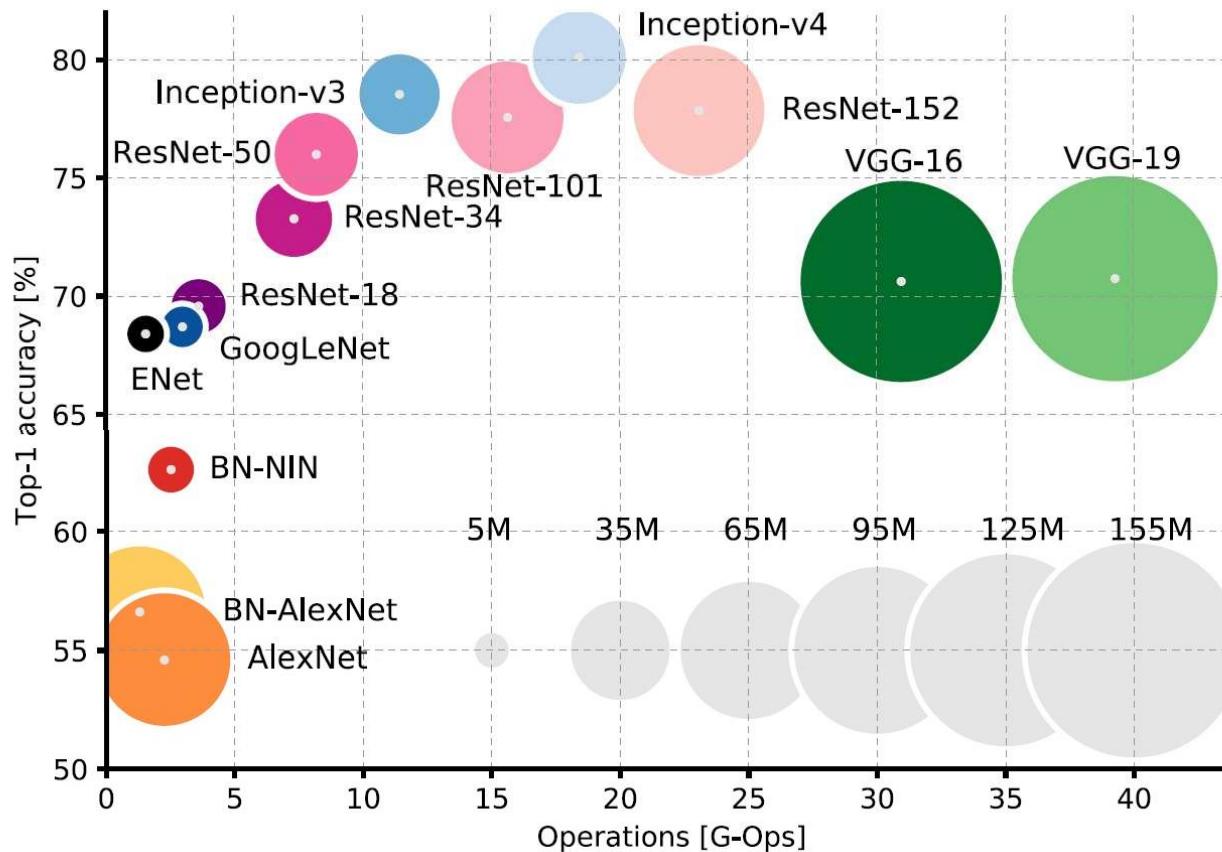
Huang, G. et al.
Densely connected convolutional networks
IEEE conference on computer vision and pattern recognition (2017)



Figures from Huang et al. (2015)

Convolutional Neural Networks

Summary



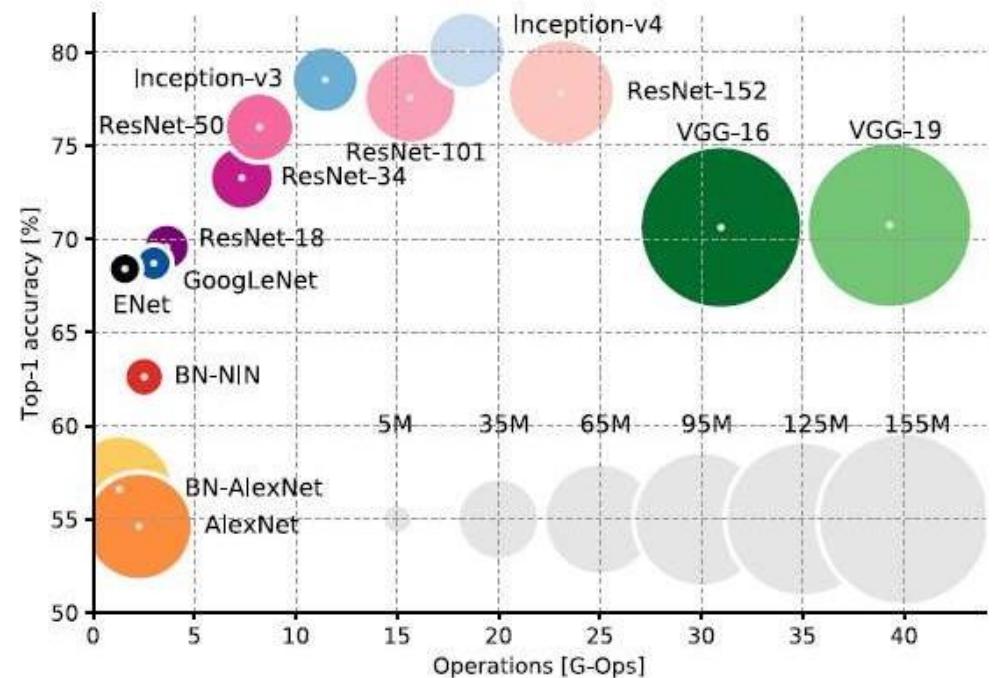
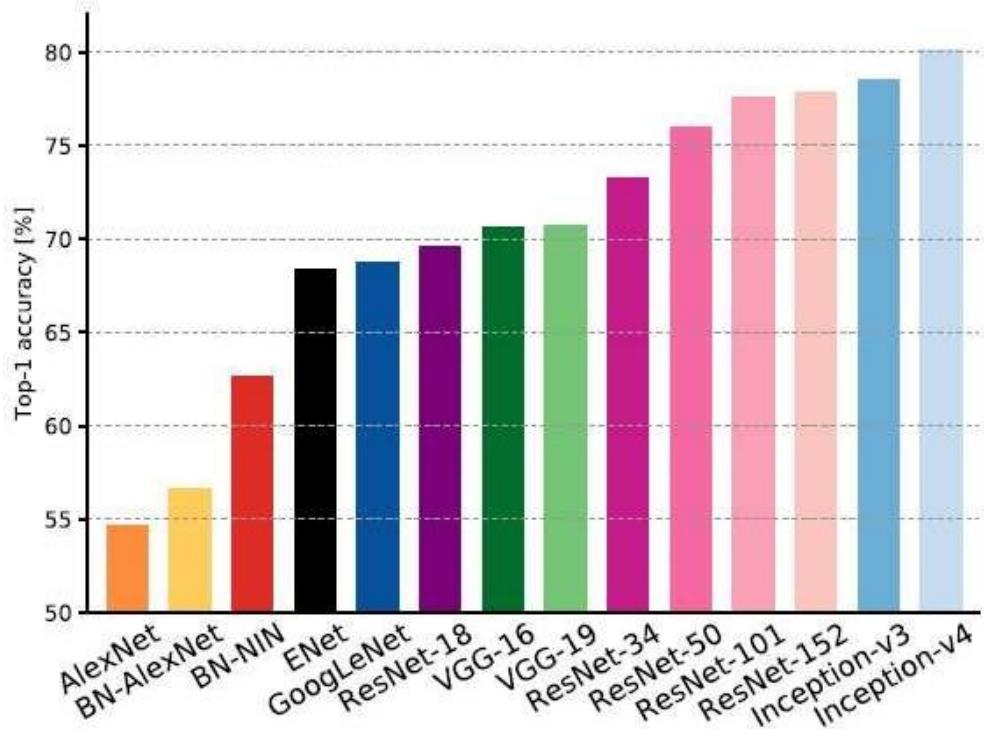
Canziani, A., Paszke, A., & Culurciello, E. (2016).

An analysis of deep neural network models for practical applications.

arXiv preprint arXiv:1605.07678.

Convolutional Neural Networks

Summary



Convolutional Neural Networks

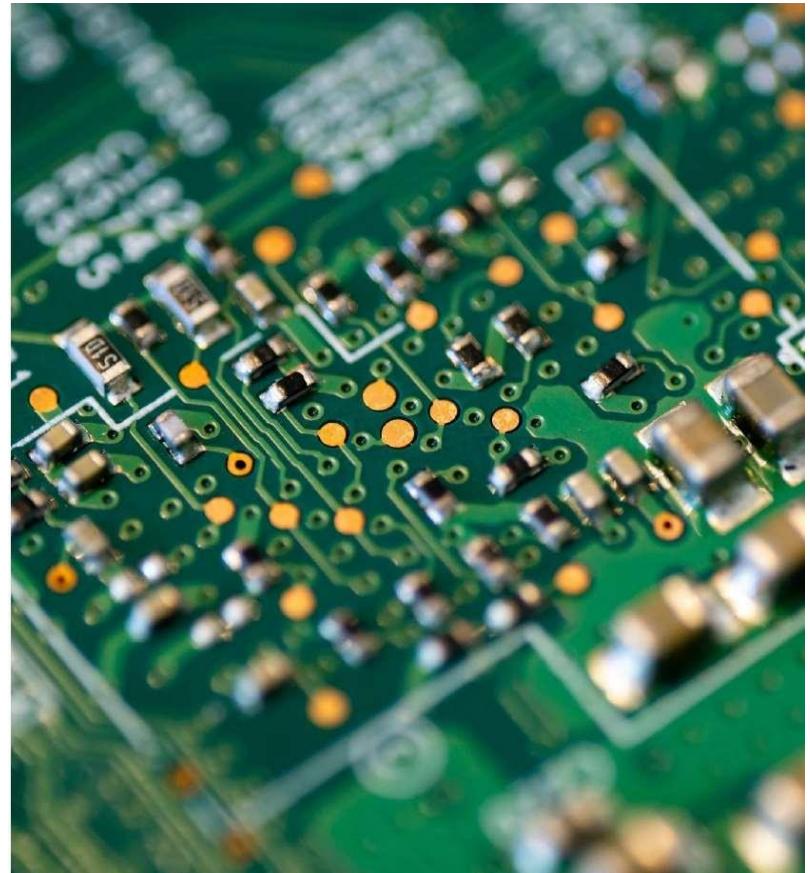
Reducing Complexity

Reducing Complexity

Depthwise Convolutions Separable
Convolutions Inverted BottleNecks
(MobileNet, MNASNet, EfficientNet)

Improving resource usage

Quantized Neural Networks
Binarized Neural Networks



Convolutional Neural Networks

Data Augmentation



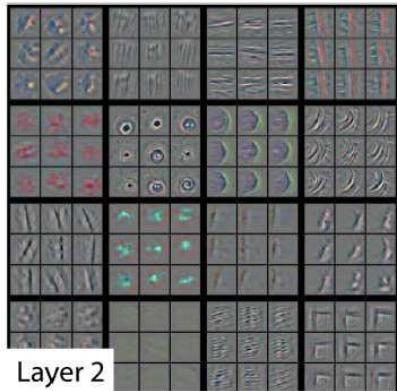
By design, convnets are only robust against **translation**

Data augmentation makes them robust against other transformations: rotation, scaling, shearing, warping, ...

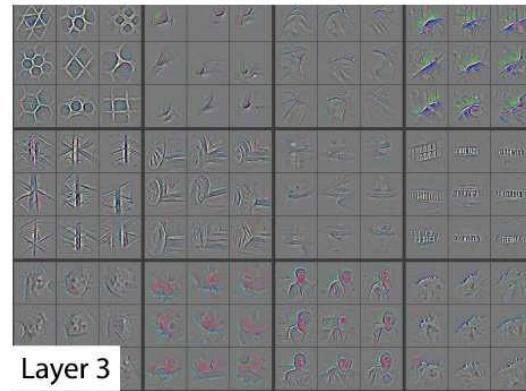


Convolutional Neural Networks

Visualizing what a convnet learns



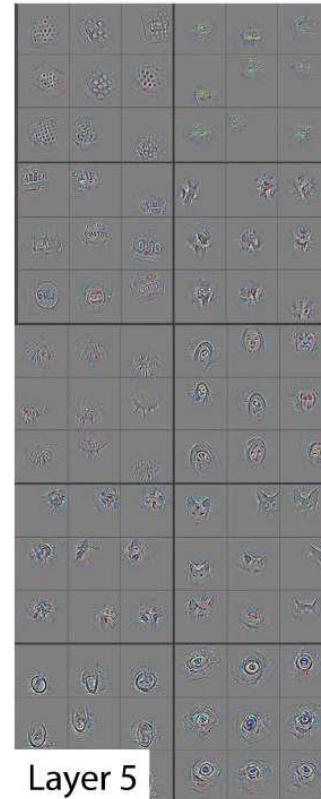
Layer 2



Layer 3



Layer 4



Layer 5

Want to learn more?



Zeiler, M.D.; Fergus, R.
*Visualizing and understanding convolutional
networks* European conference on computer
vision (2014)

Convolutional Neural Networks

Visualizing what a convnet learns

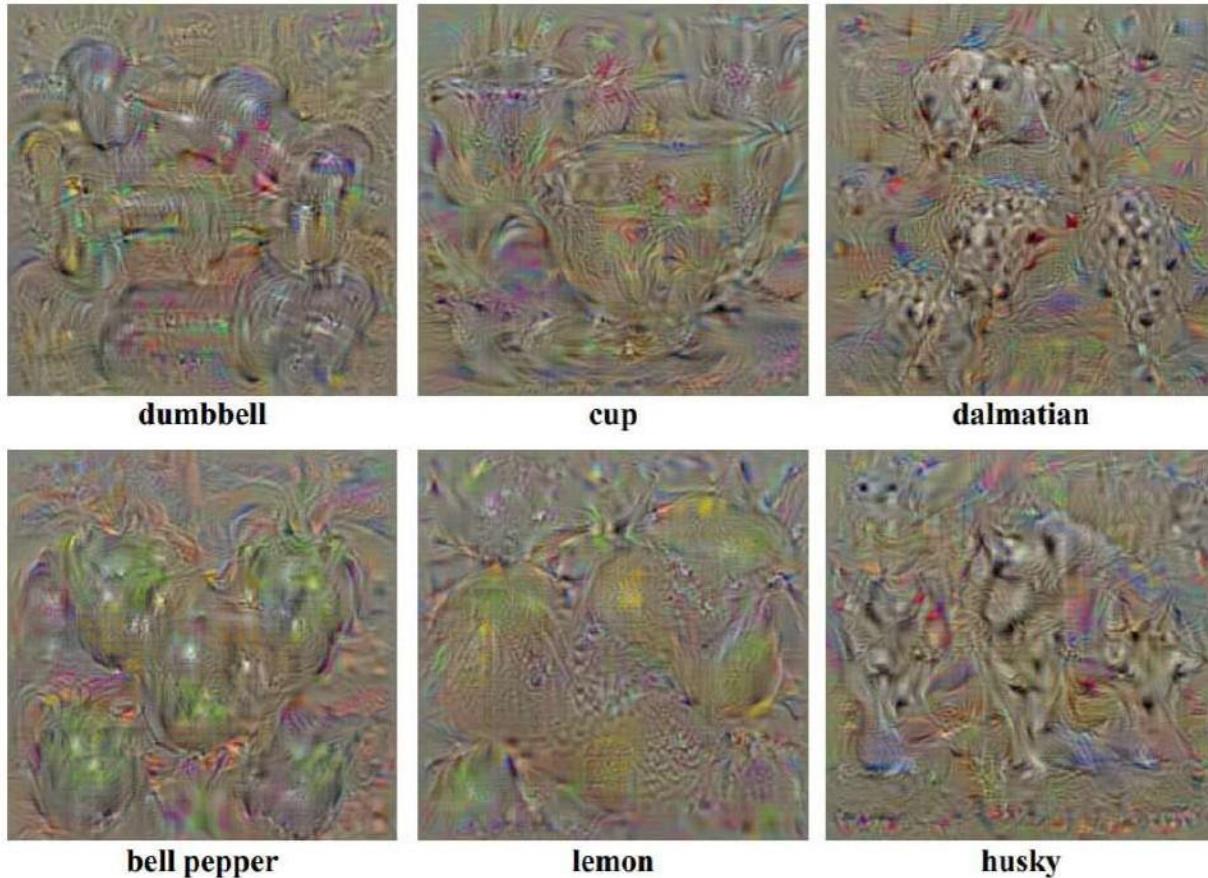


Figure from Simonyan et al. (2013)

Convolutional Neural Networks

Visualizing what a convnet learns

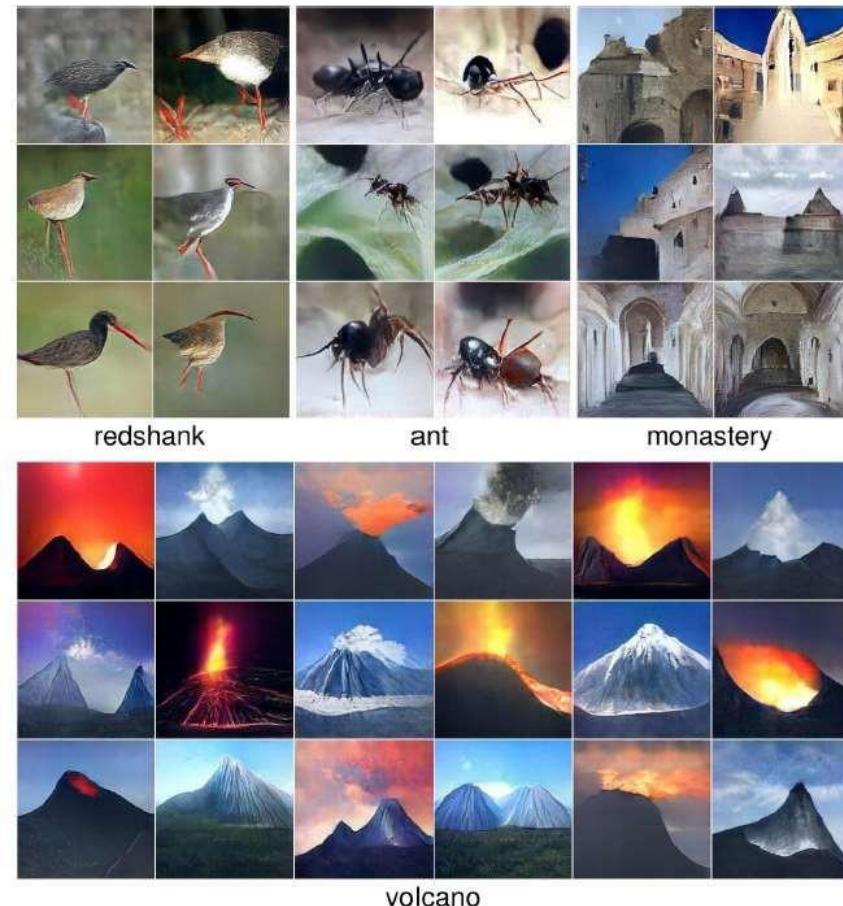


Figure from Nguyen et al. (2016)

Convolutional Neural Networks

Visualizing what a convnet learns

<https://distill.pub/2017/feature-visualization/> by Chris Olah, Alexander Mordvintsev and Ludwig Schubert

Feature Visualization

How neural networks build up their understanding of images



Feature visualization allows us to see how GoogLeNet^[1], trained on the ImageNet^[2] dataset, builds up its understanding of images over many layers. Visualizations of all channels are available in the [appendix](#).

AUTHORS

Chris Olah

Alexander Mordvintsev

Ludwig Schubert

AFFILIATIONS

Google Brain Team

Google Research

Google Brain Team

PUBLISHED

Nov. 7, 2017

DOI

10.23915/distill.00007